

Introduction to Kubernetes and OpenShift for System Administrators

Ohio LinuxFest (OLF) '23

Jay Ryan
Senior Solutions
Architect

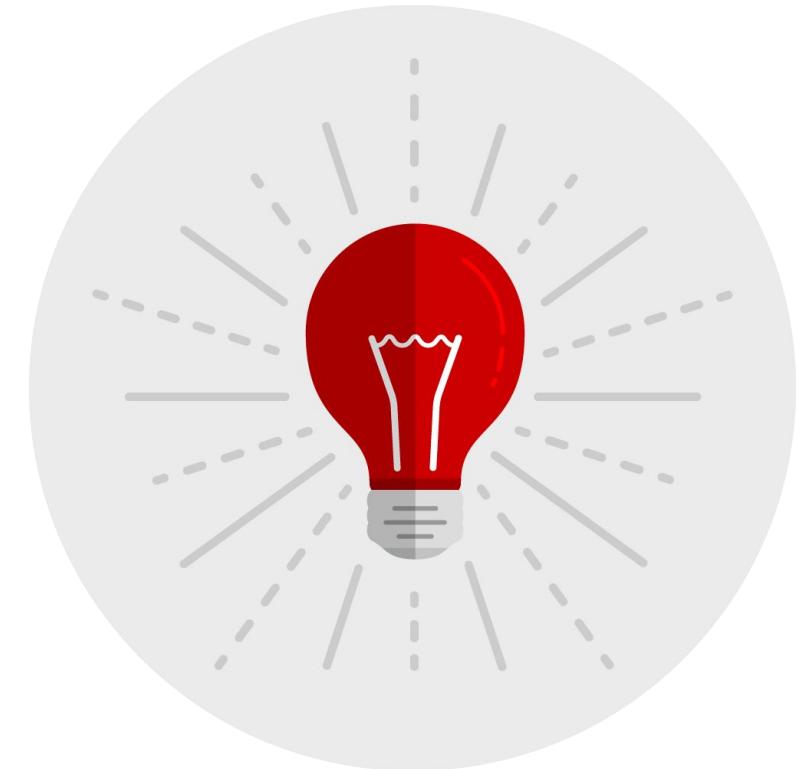
Brad Krumme
Senior Platform Solutions
Architect

AGENDA

Morning Session 09:00→12:00 EST

Lunch Break 12:00→1:00 EST

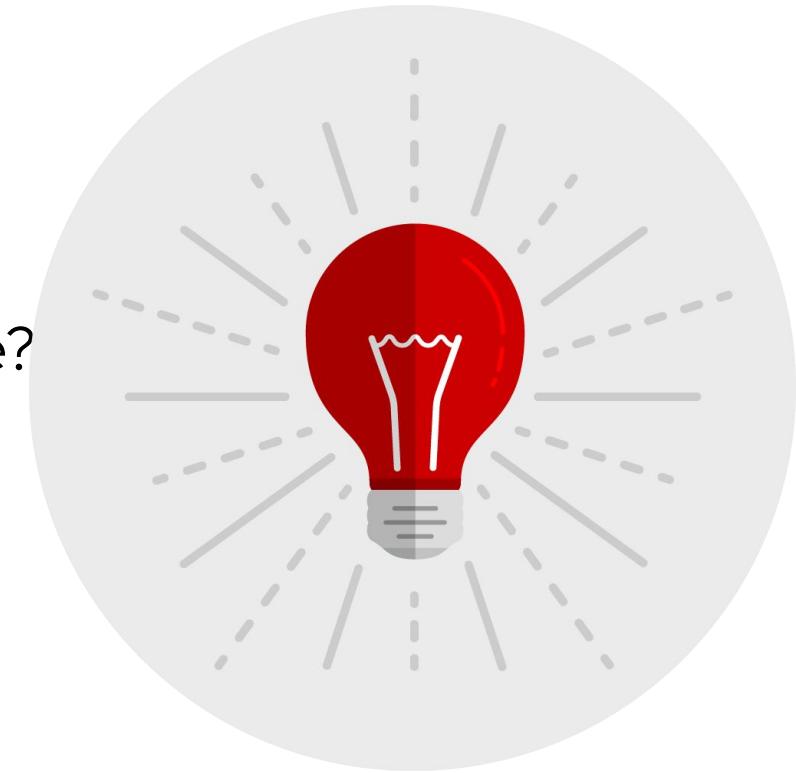
Afternoon Session 1:00→5:00 EST



AGENDA

Morning Session

- Kubernetes and Containers - Why are we here?
- Container Technology Introduction
- What is Kubernetes
- Kubernetes Cluster Fundamentals
- 10 minute break (~10:00-10:30?)
- Workshop Introduction
- LAB: Module 1 & 2
- LAB: Installing / Managing Applications



AGENDA

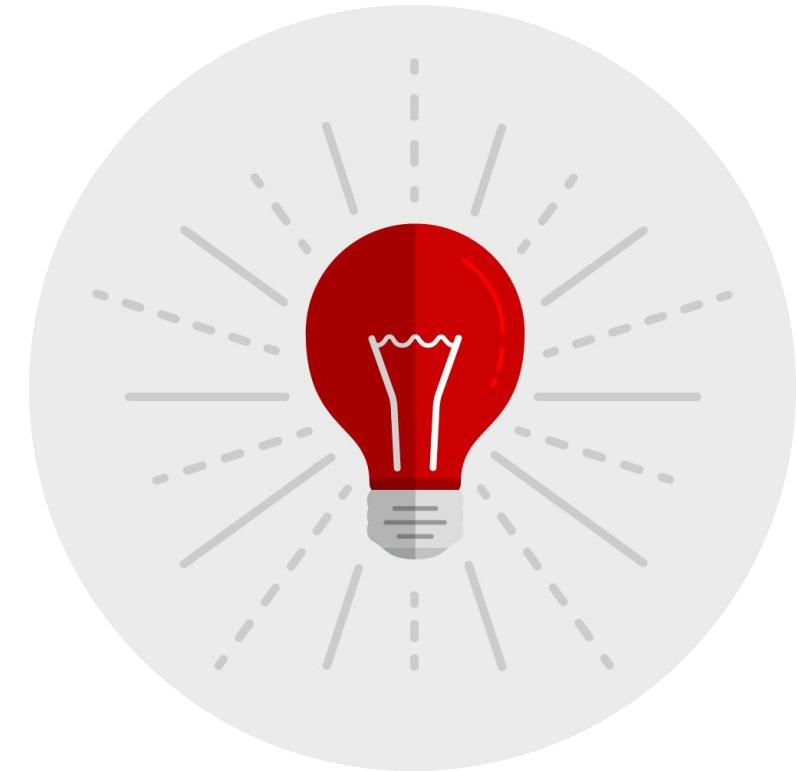
Afternoon Session

- Security fundamentals
- LAB: Managing and Scaling Nodes
- Day 2 Operations:
 - LAB: Logging / Authentication /Monitoring
 - LAB: Templates / Quotas / Limits - Oh My!
 - LAB: Networking / NetworkPolicy
- LAB: GitOps
- Bonus LAB: Windows Containers!!!



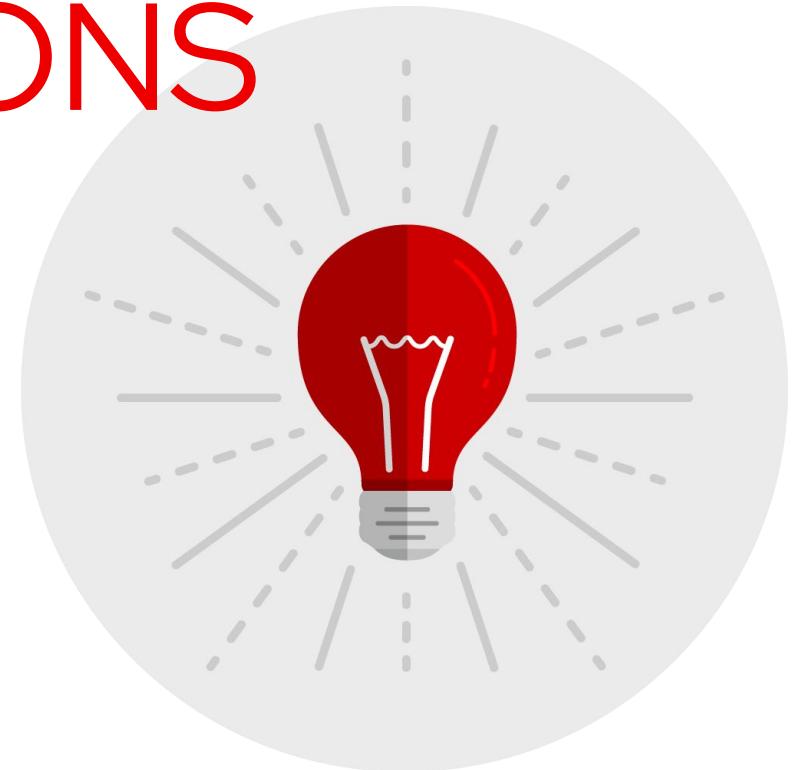
LOGISTICS

- Requirements - A laptop with internet access ~~and ssh access to a public address~~
- No food or drinks on the tables
- There are no stupid questions
- I don't know / I'll get back to you
- Small group / Let's make it Interactive
- You will get out what you put in
- Building Logistics / Questions?



INTRODUCTIONS

- What Company or Organization are you from?
- What's your role (not your title!)
- What are you looking to get out of this Lab today?



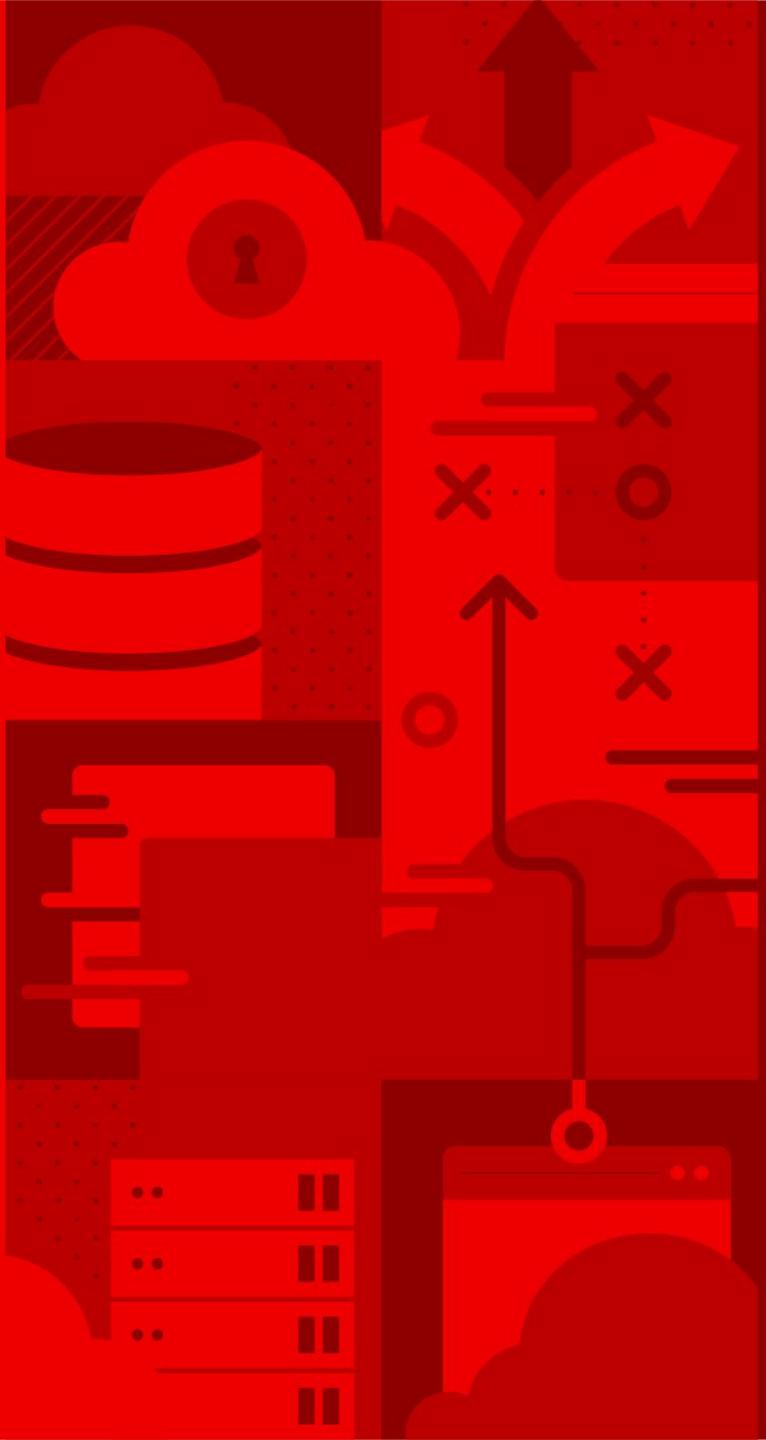
About Me!

Jay Ryan

Senior Account Solutions Architect
@Red Hat

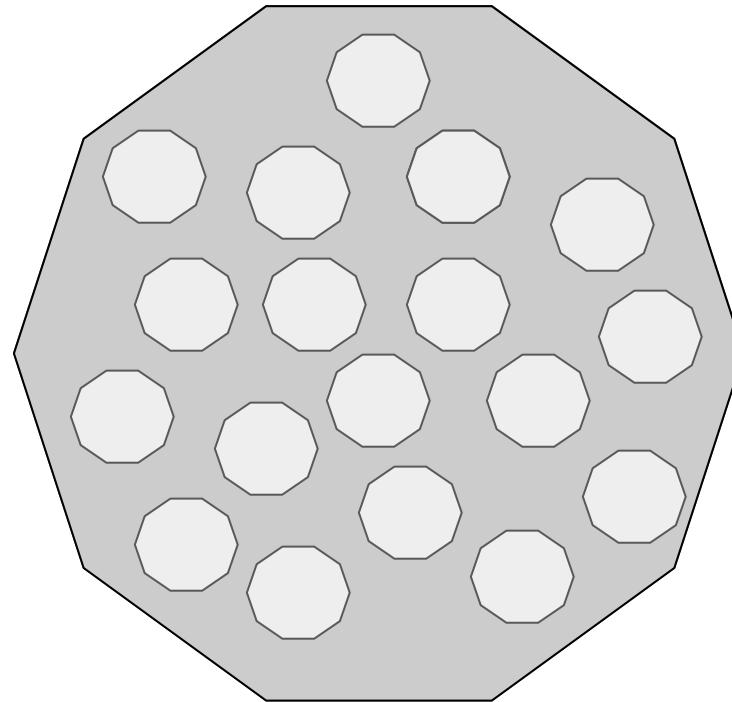
@jaywryan



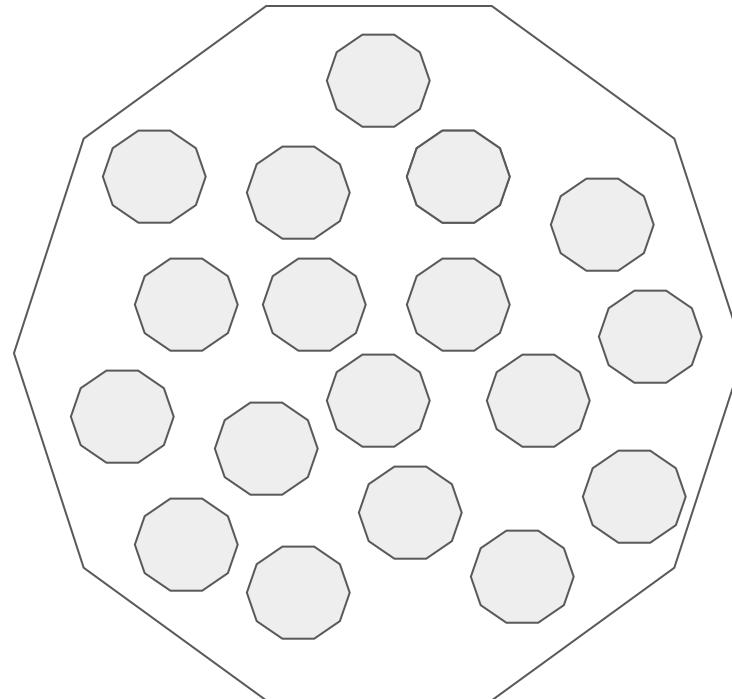


Kubernetes and Containers: Why are we here?

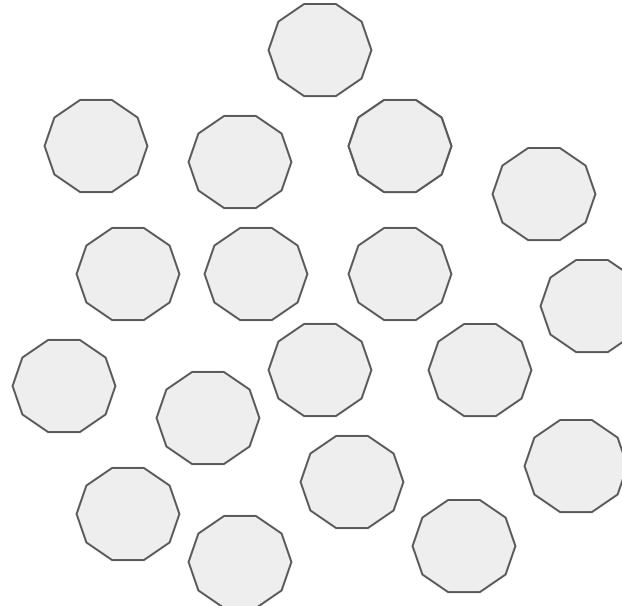
The Application



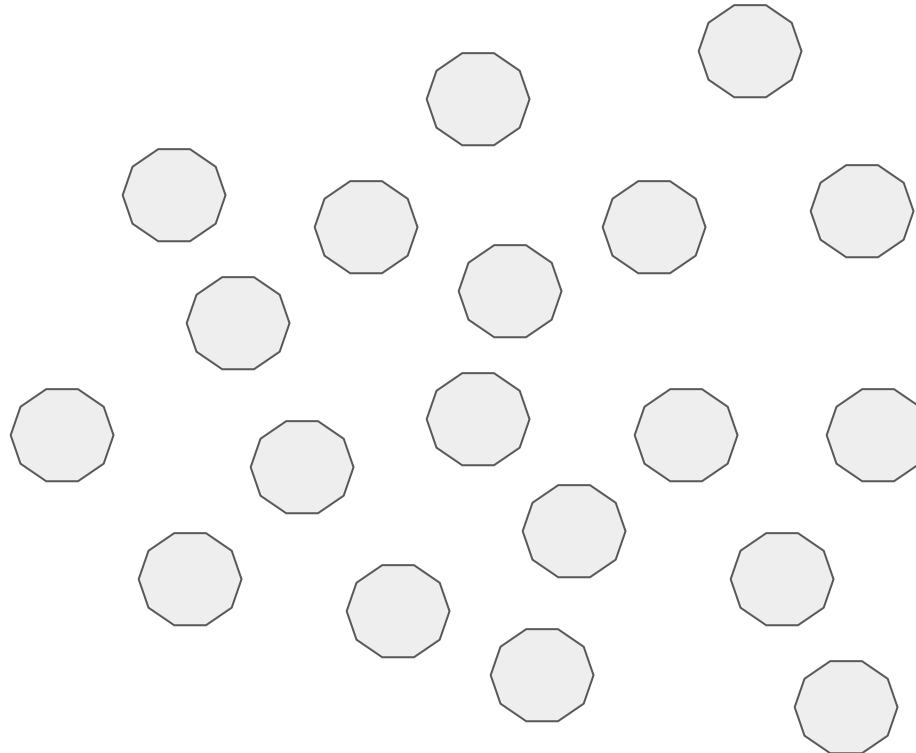
Modules



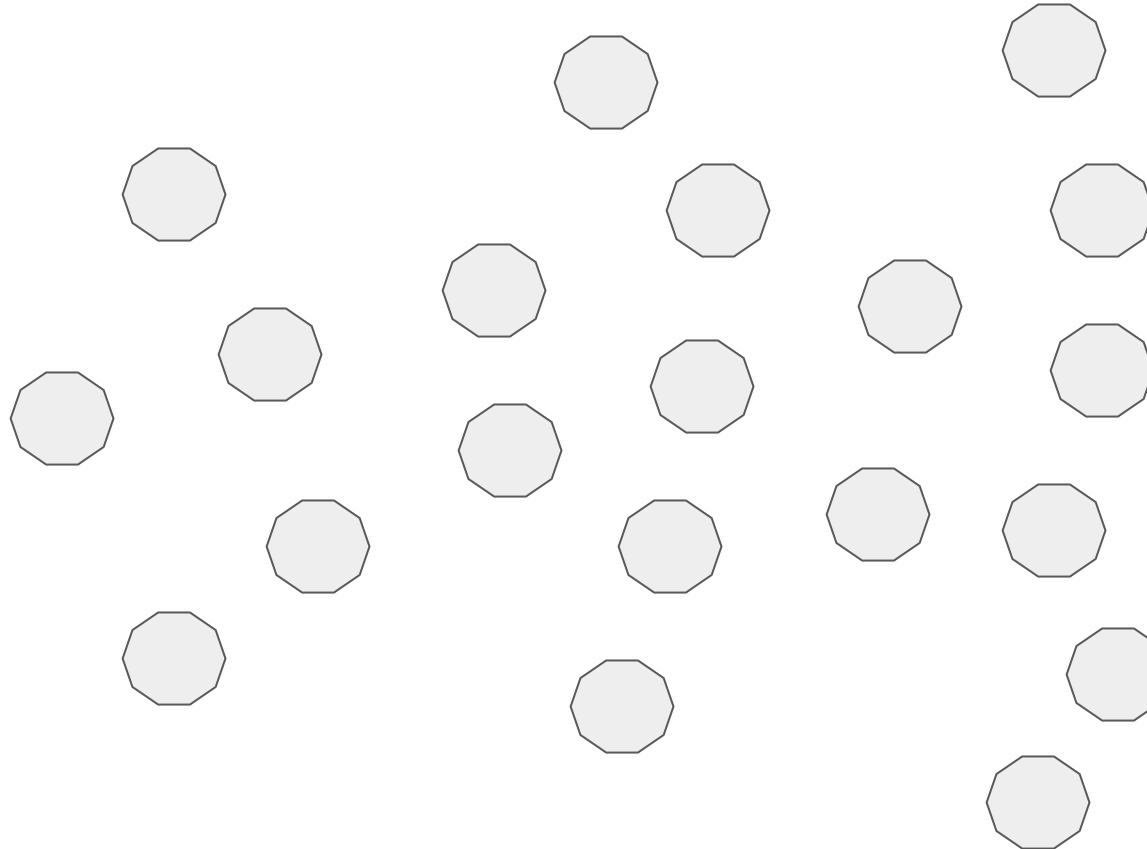
Microservices



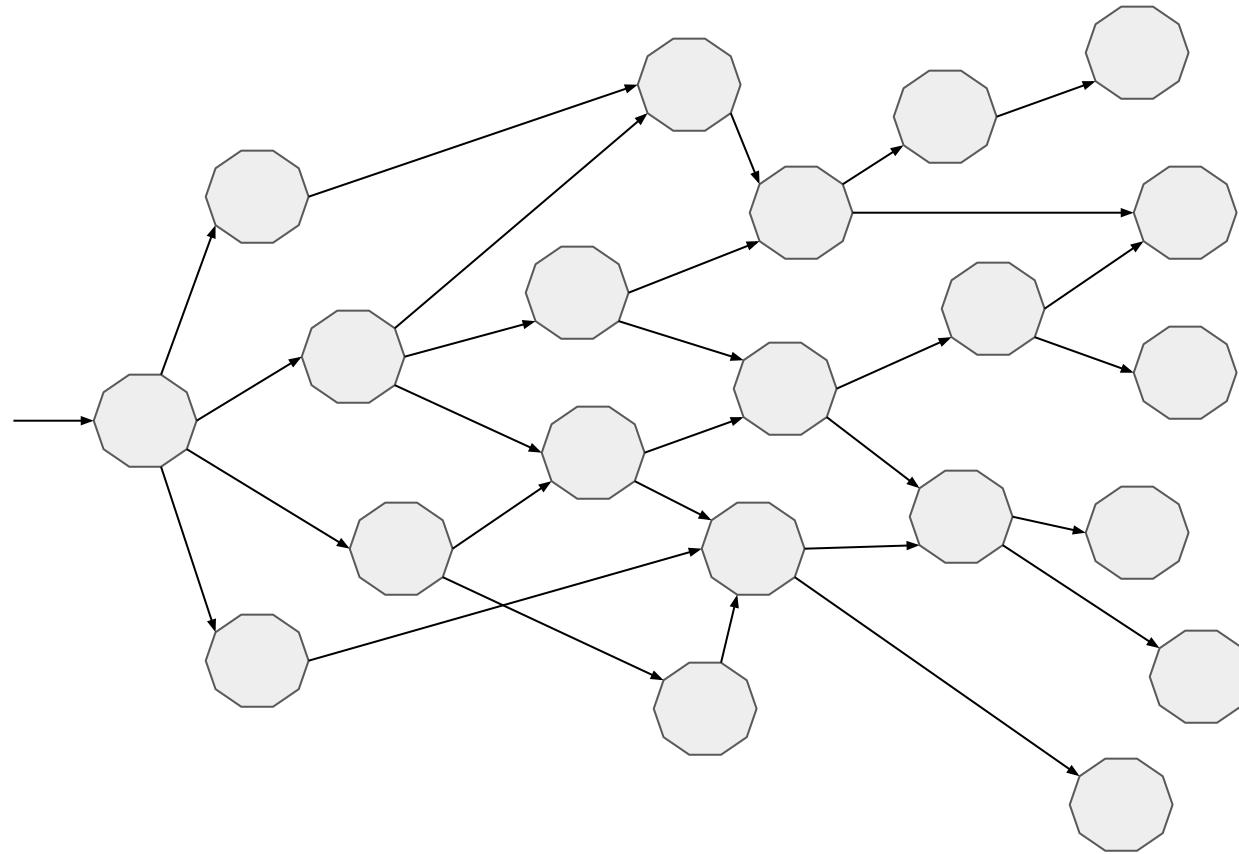
Microservices



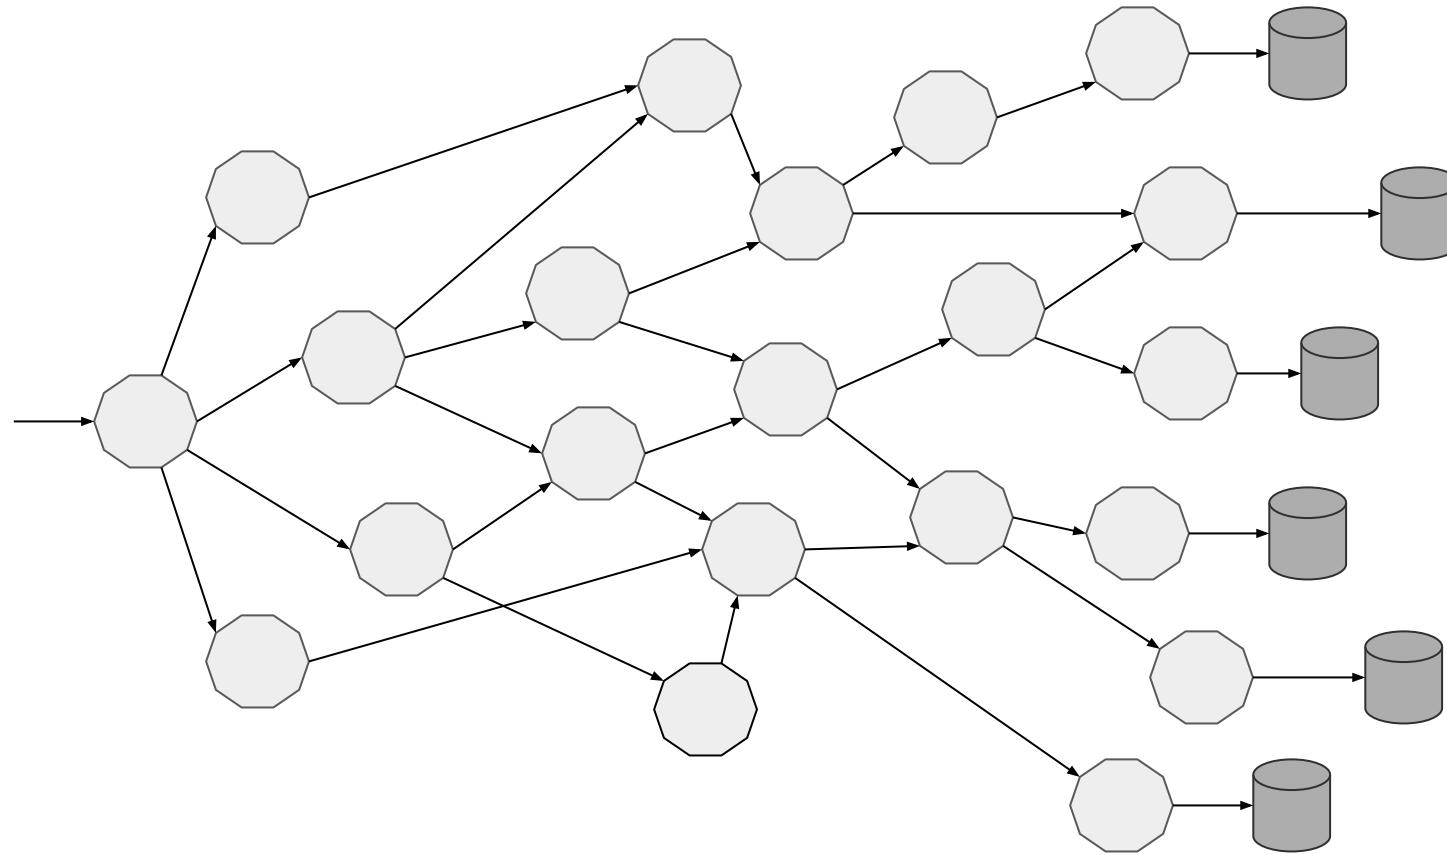
Microservices



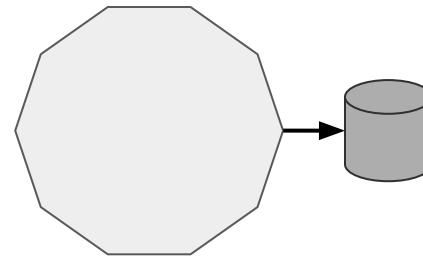
Network of Services



Microservices own their Data

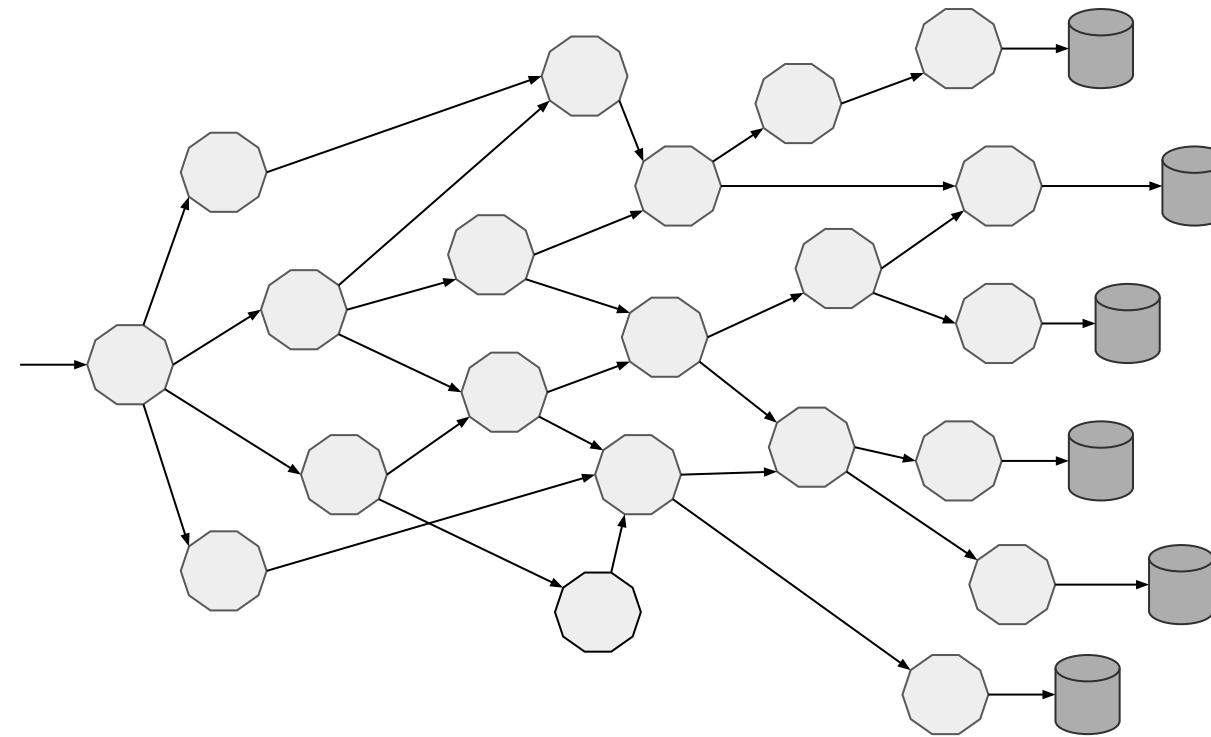


Old School

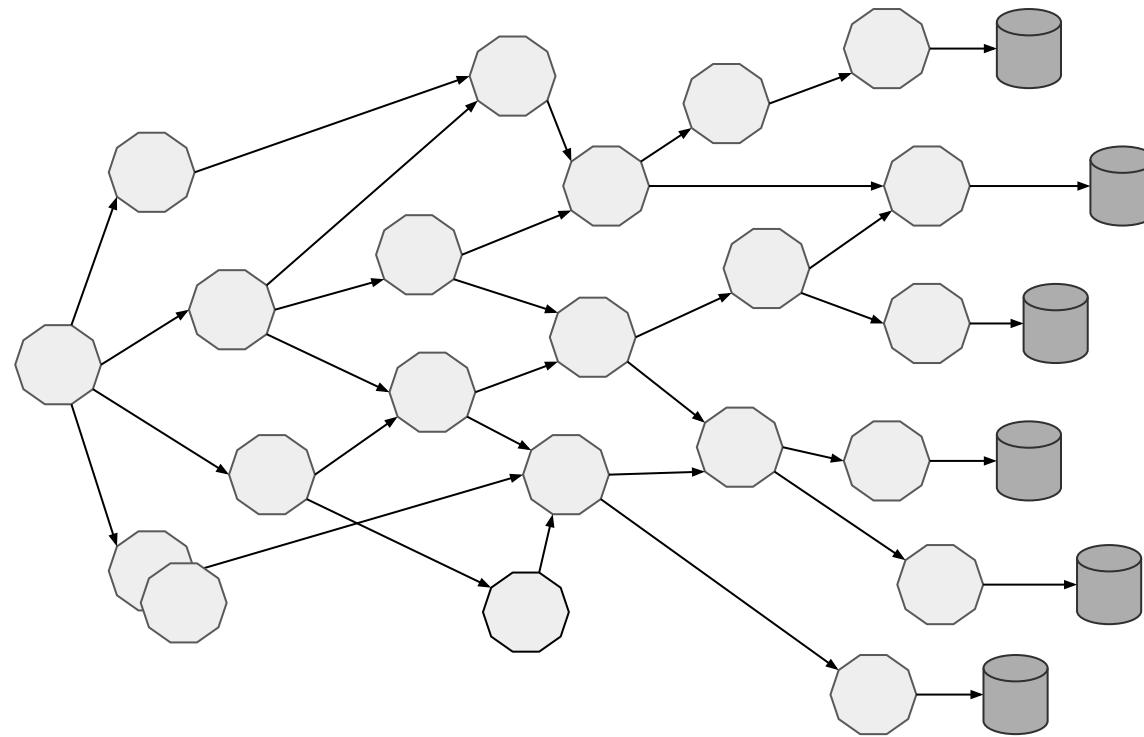
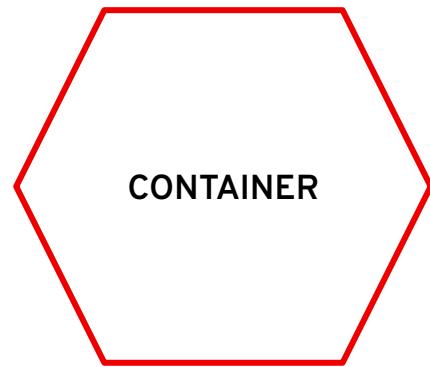


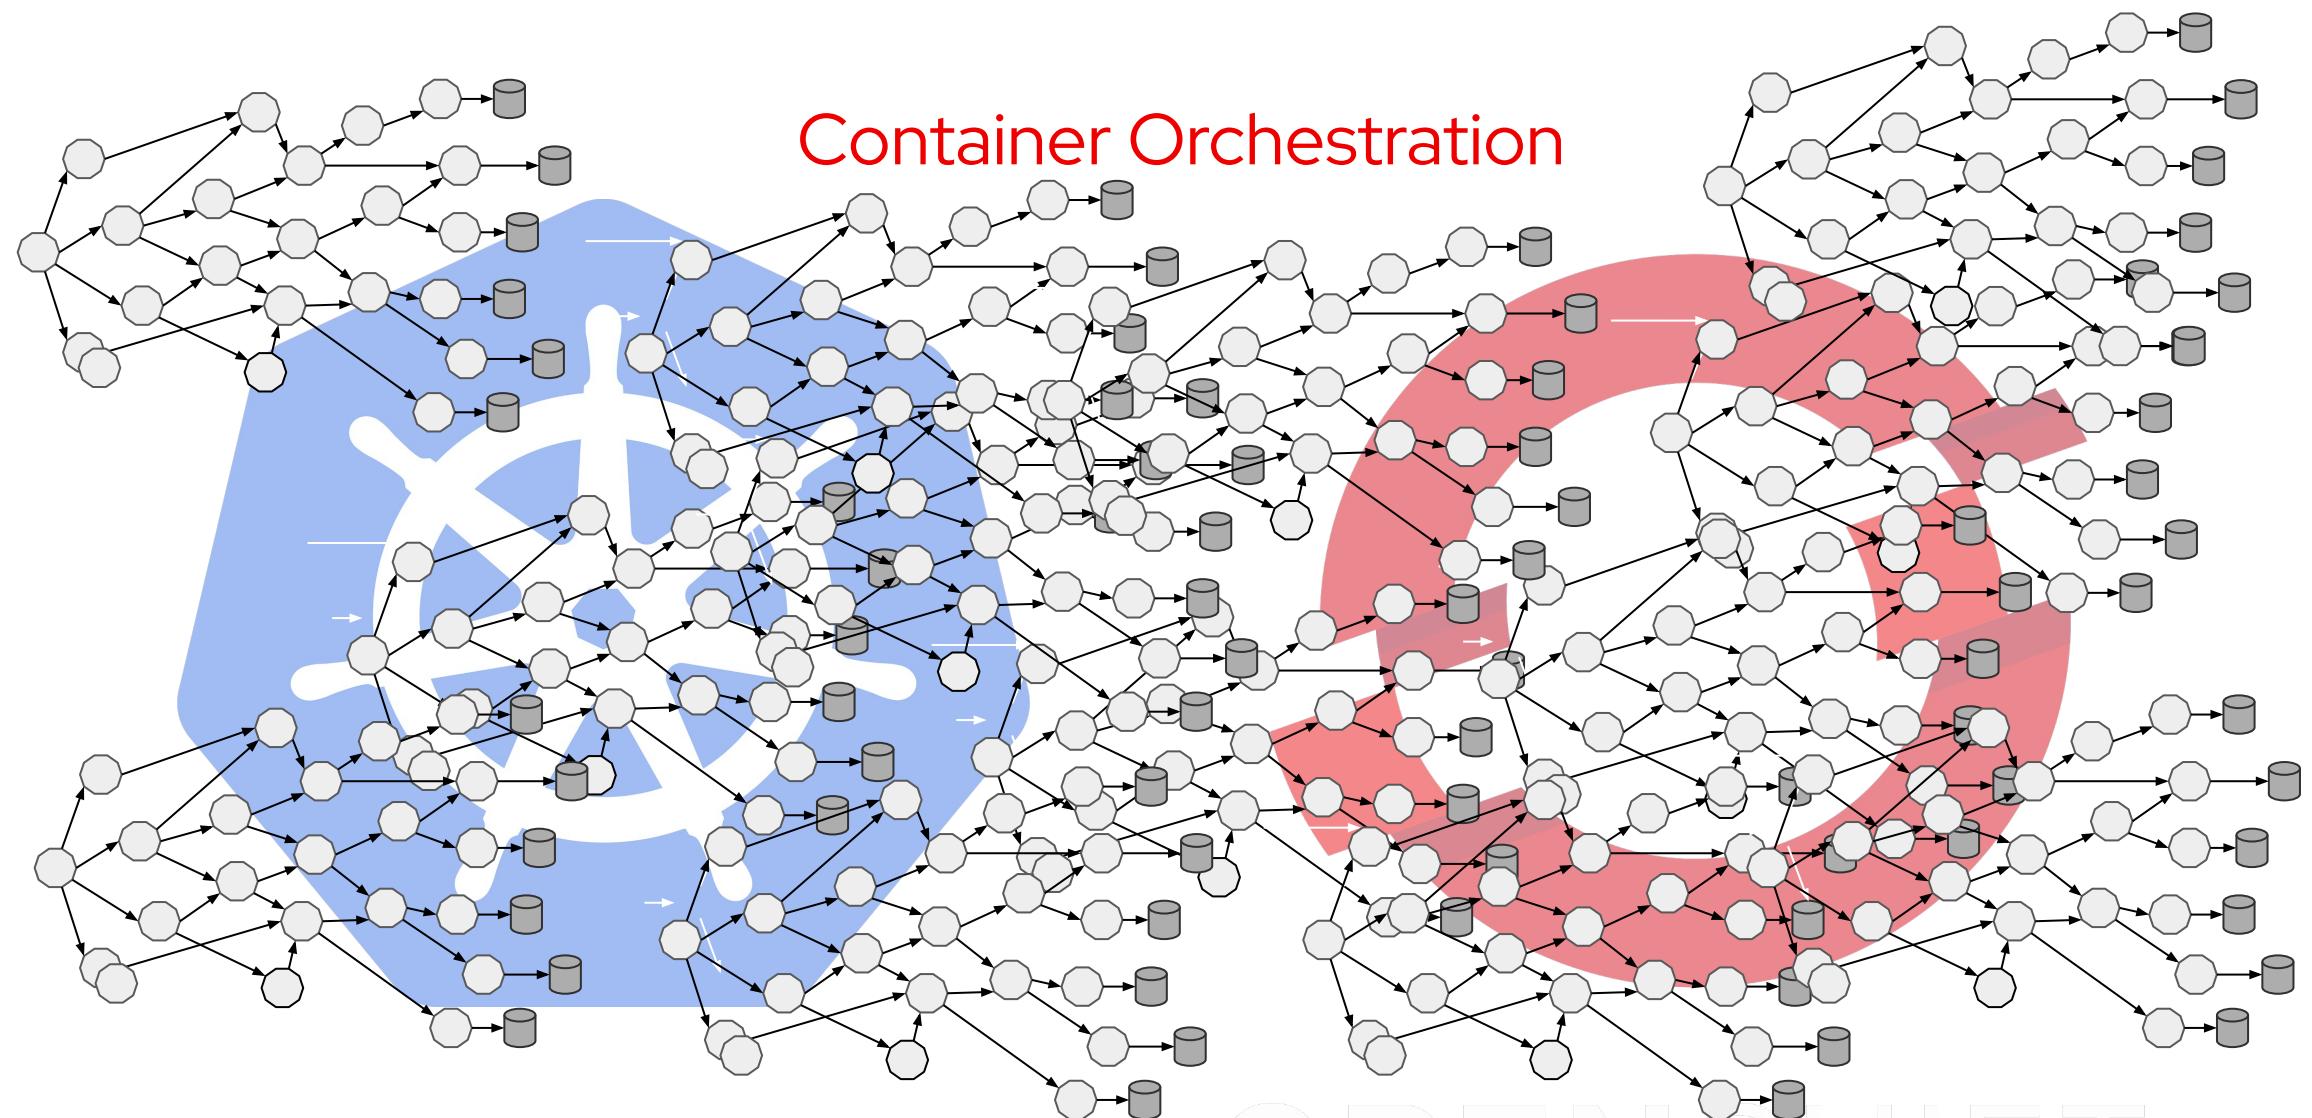
Love Thy Mono

New School



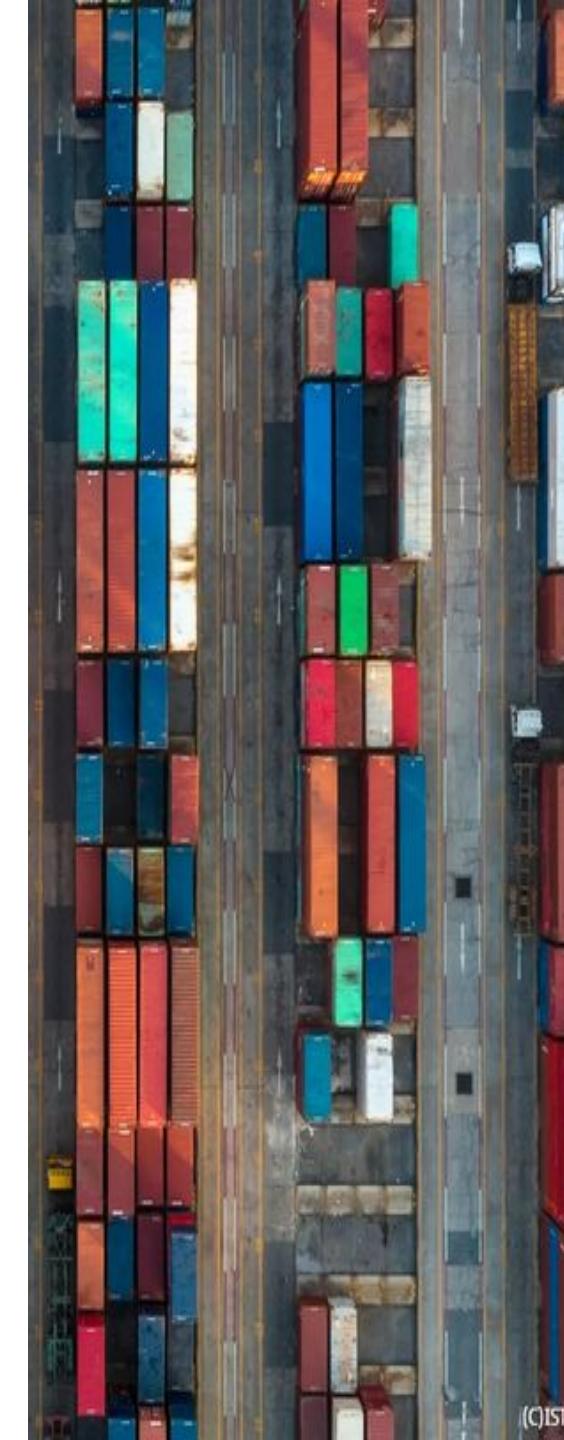
Containers





Container Technology

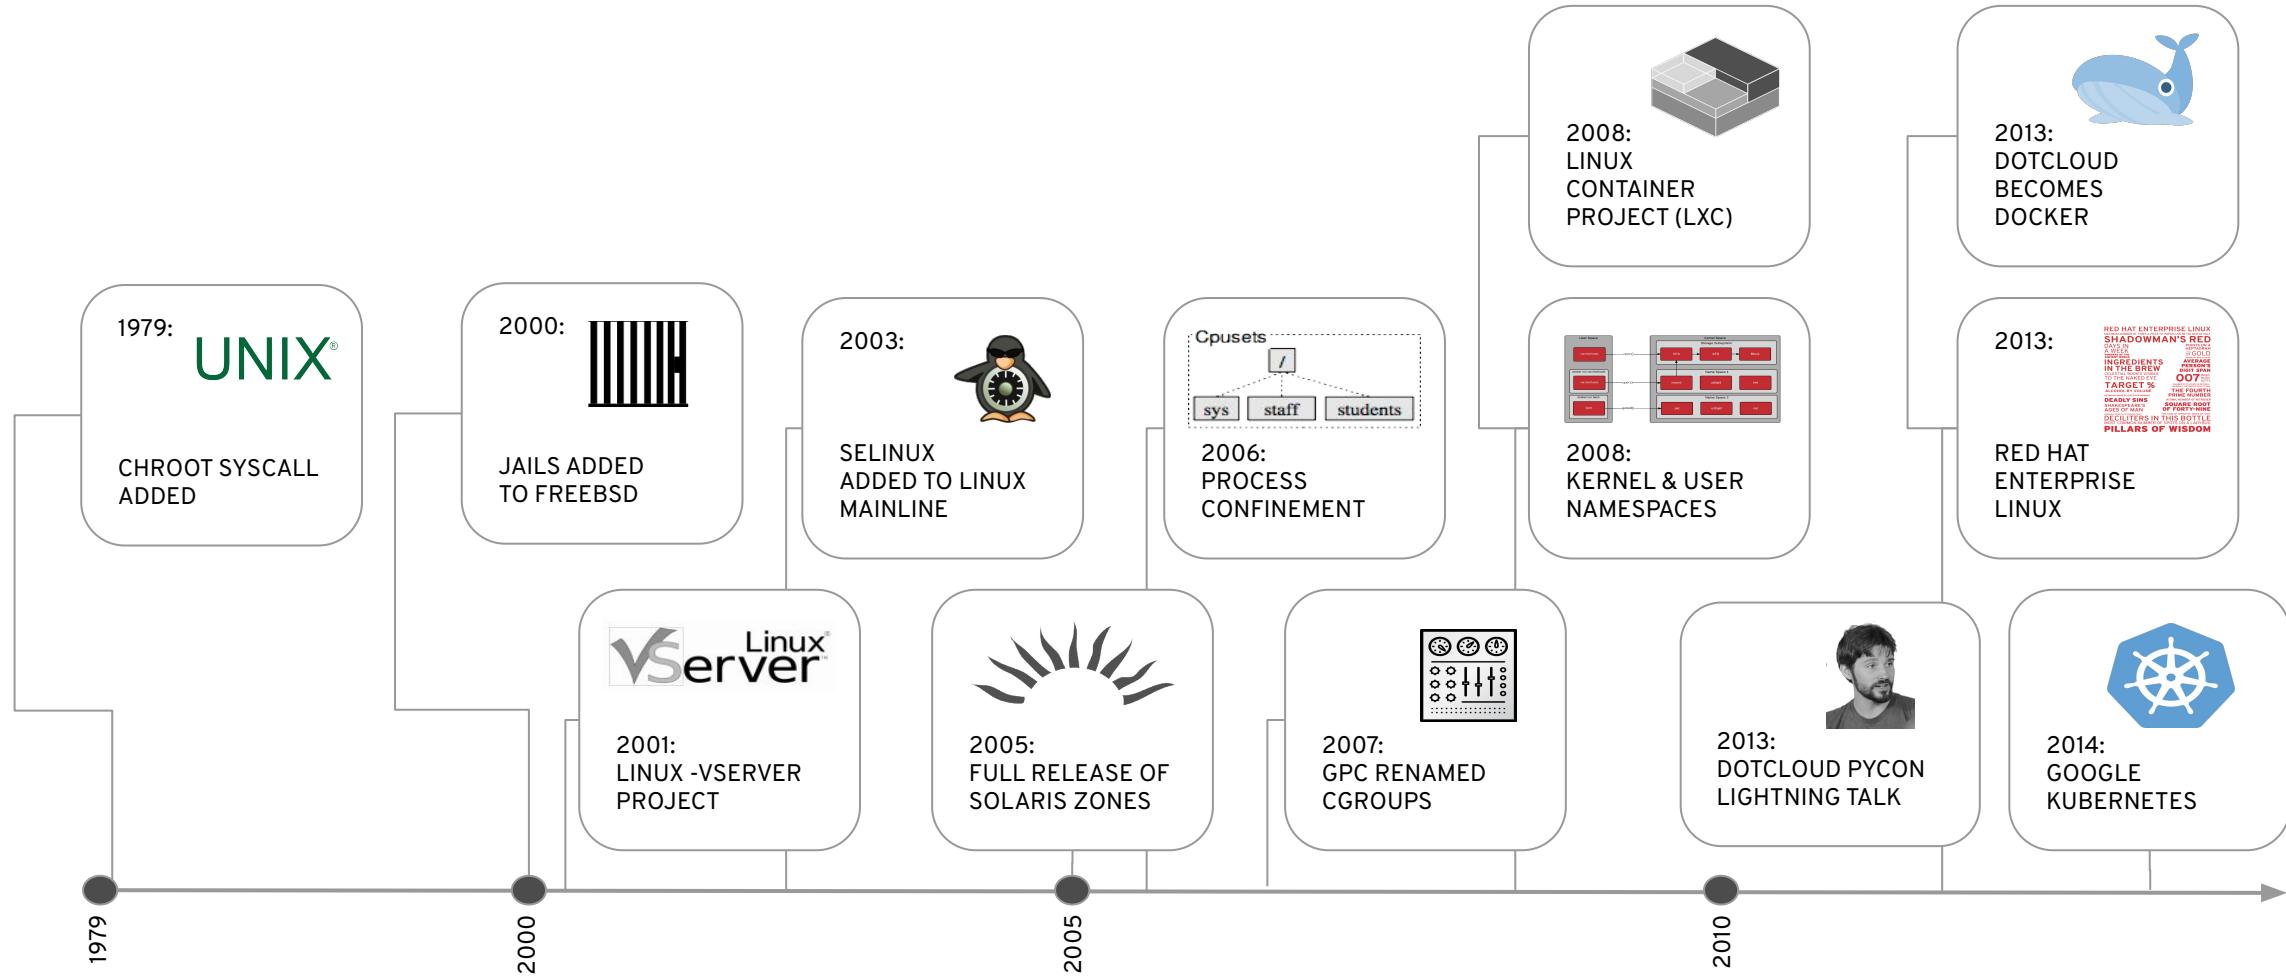
Introduction



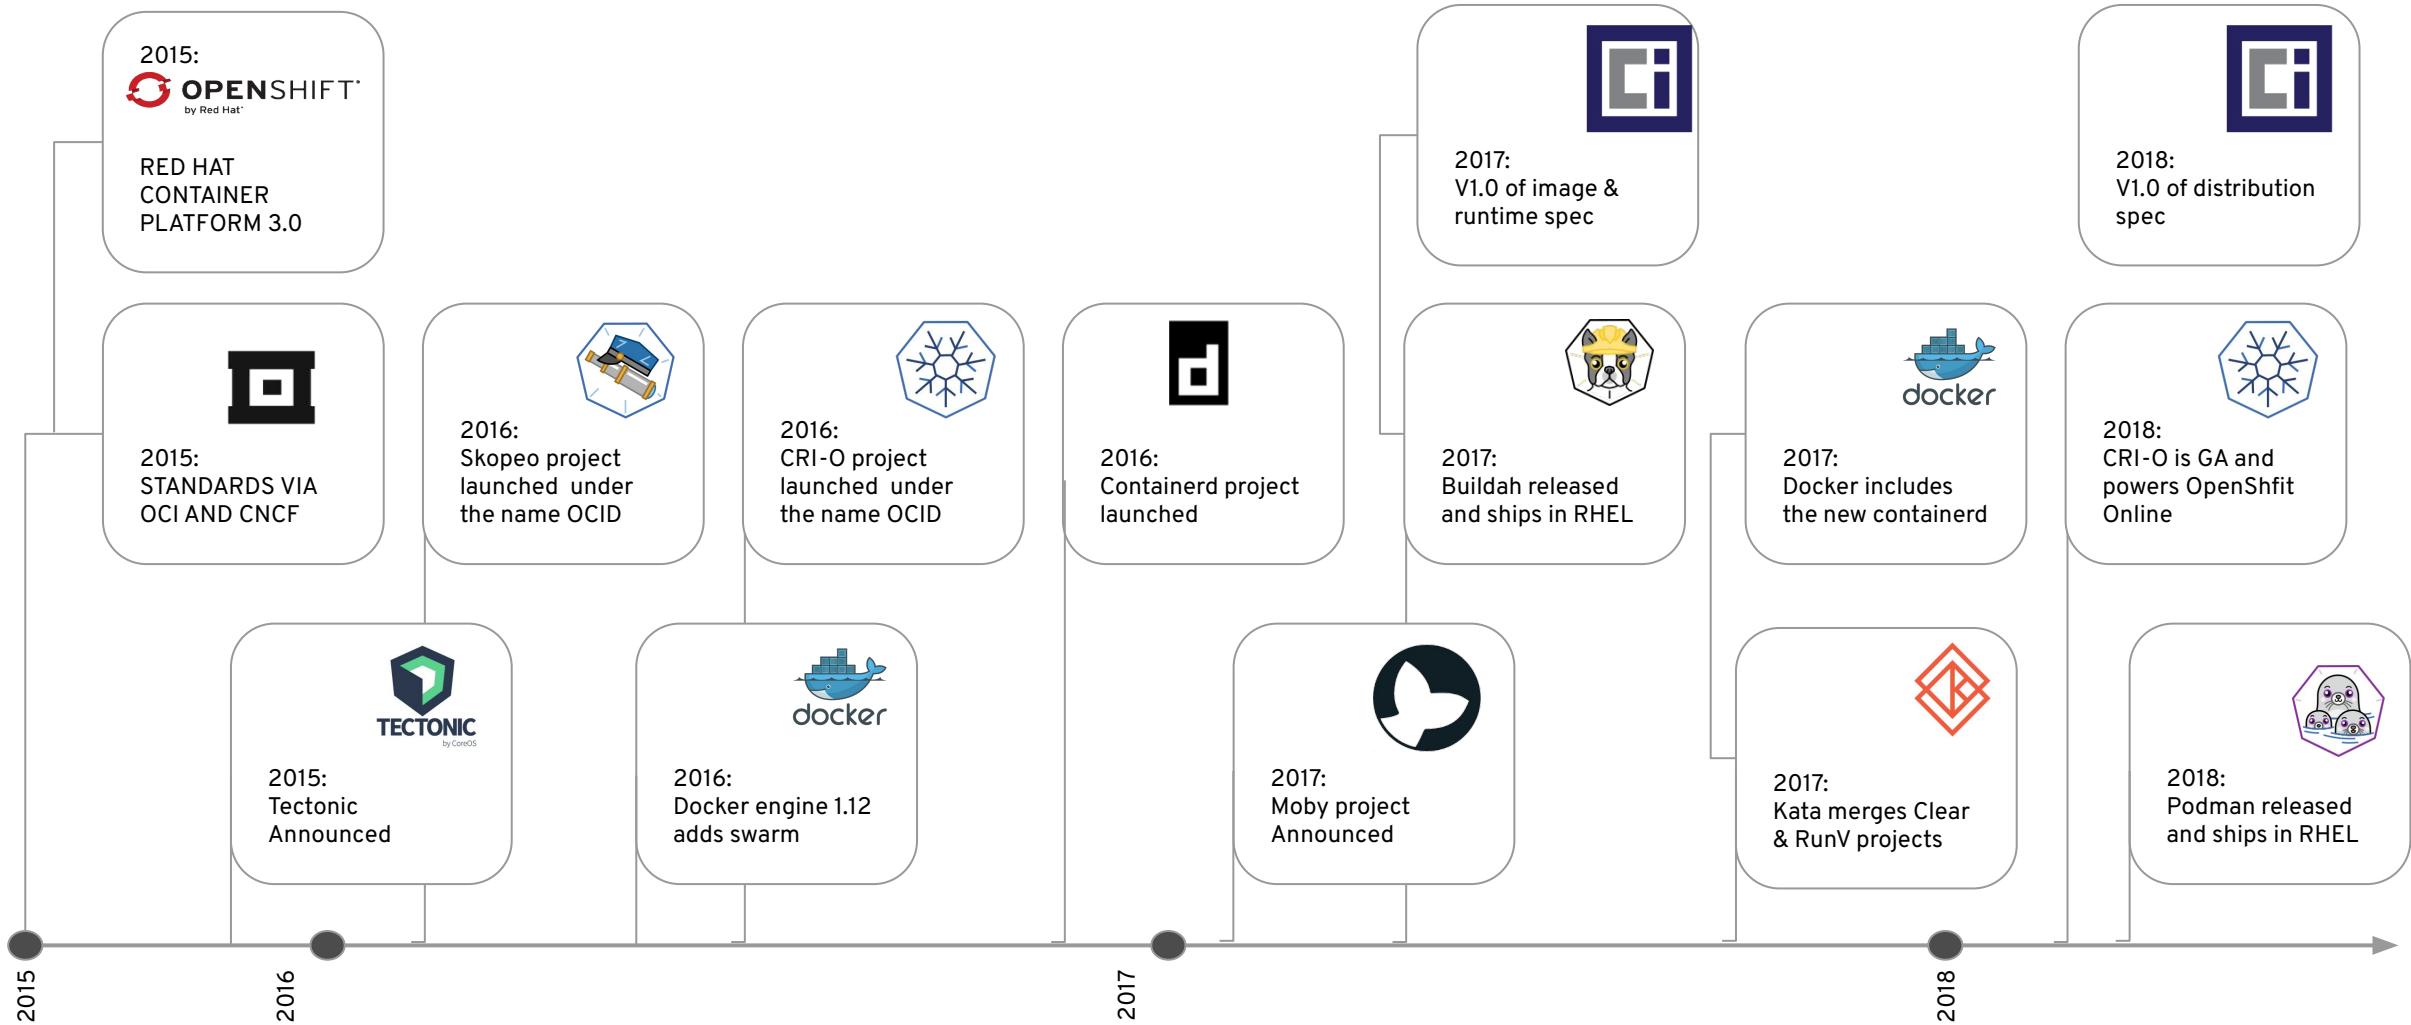
What is a container?

- A unit of software that contains code and all its dependencies so that the application can be run in the same way in different computational environments.
- **Operating system** - A process in a sandbox
- **System Administrator** - Portability for an app and its runtime dependencies
- **Developer** - Universal Application Packager

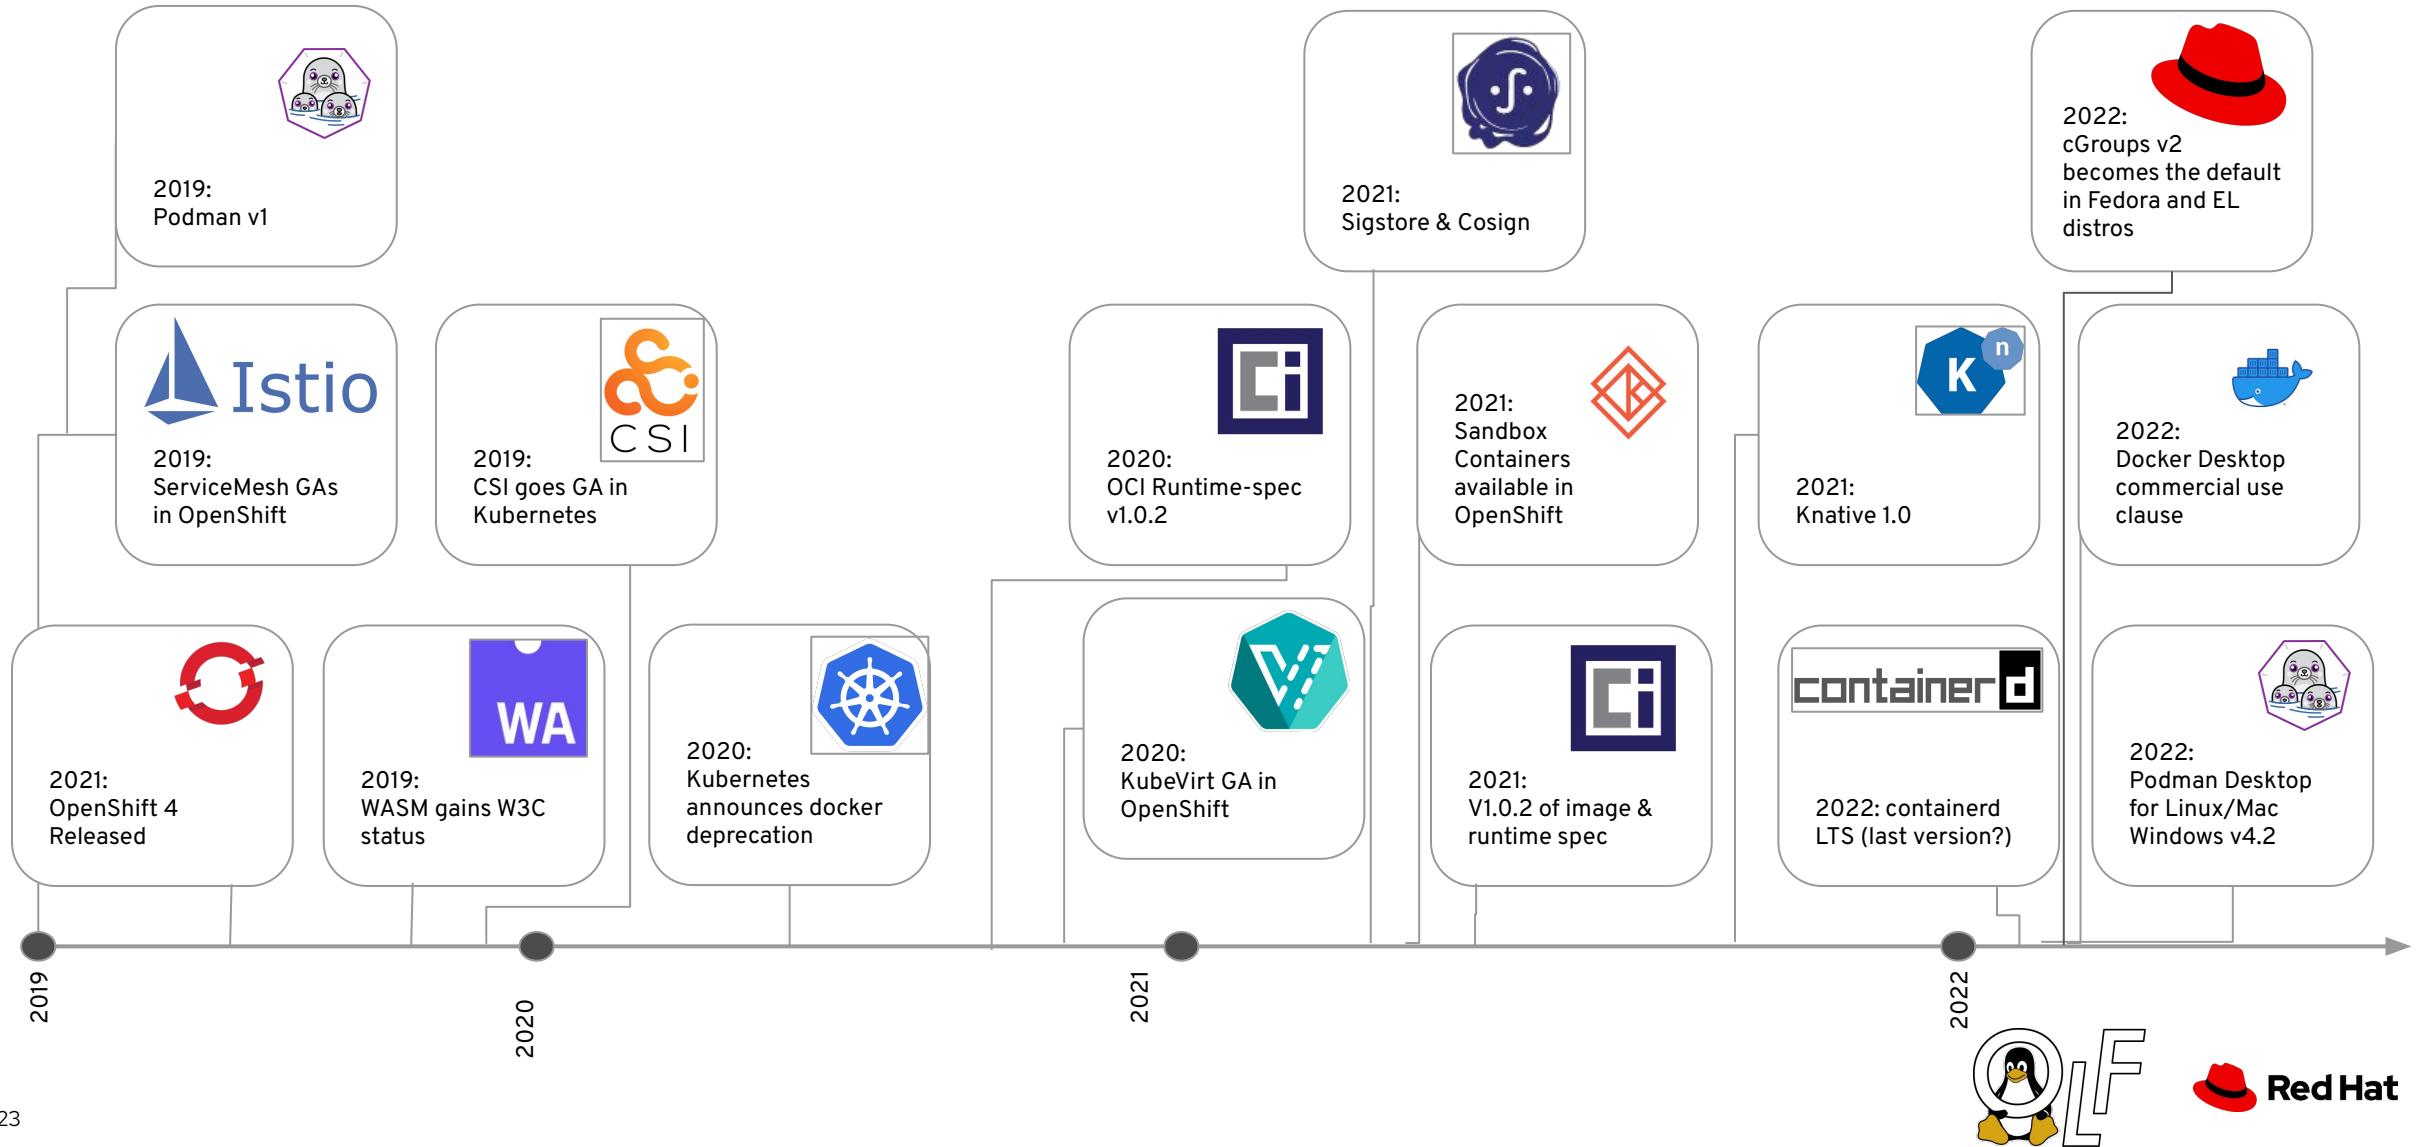
The History of Containers



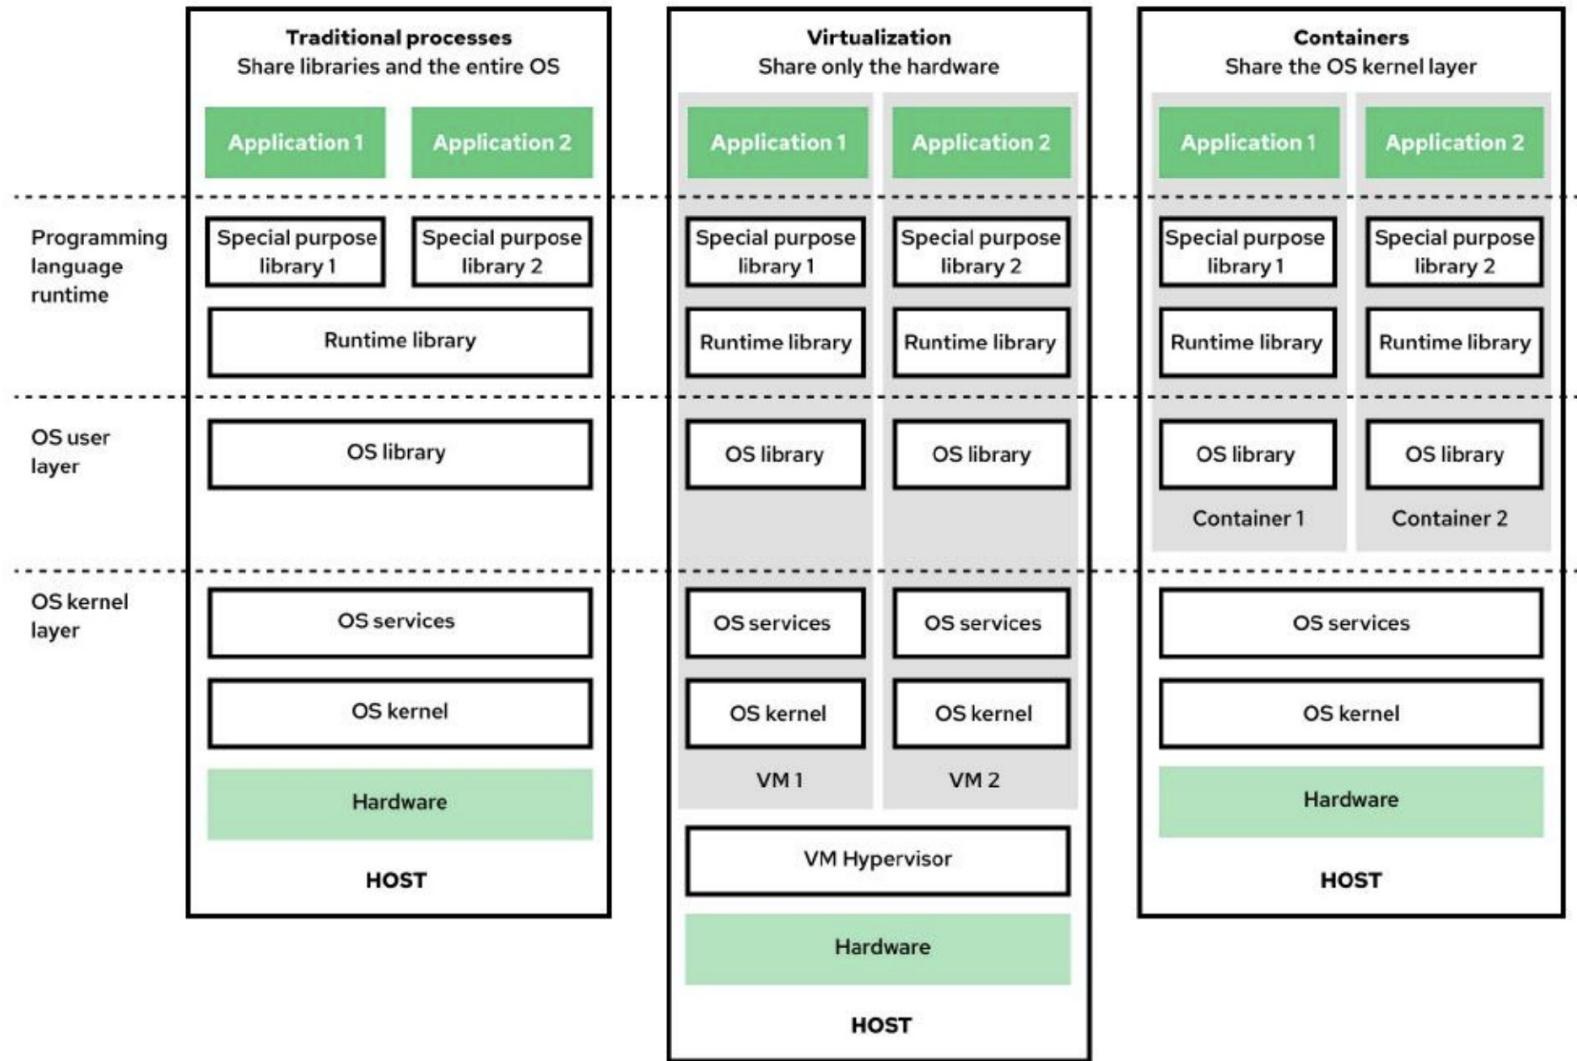
CONTAINER INNOVATION ...



CONTAINER INNOVATION IS NOT FINISHED



What is a container?



How do we make one?

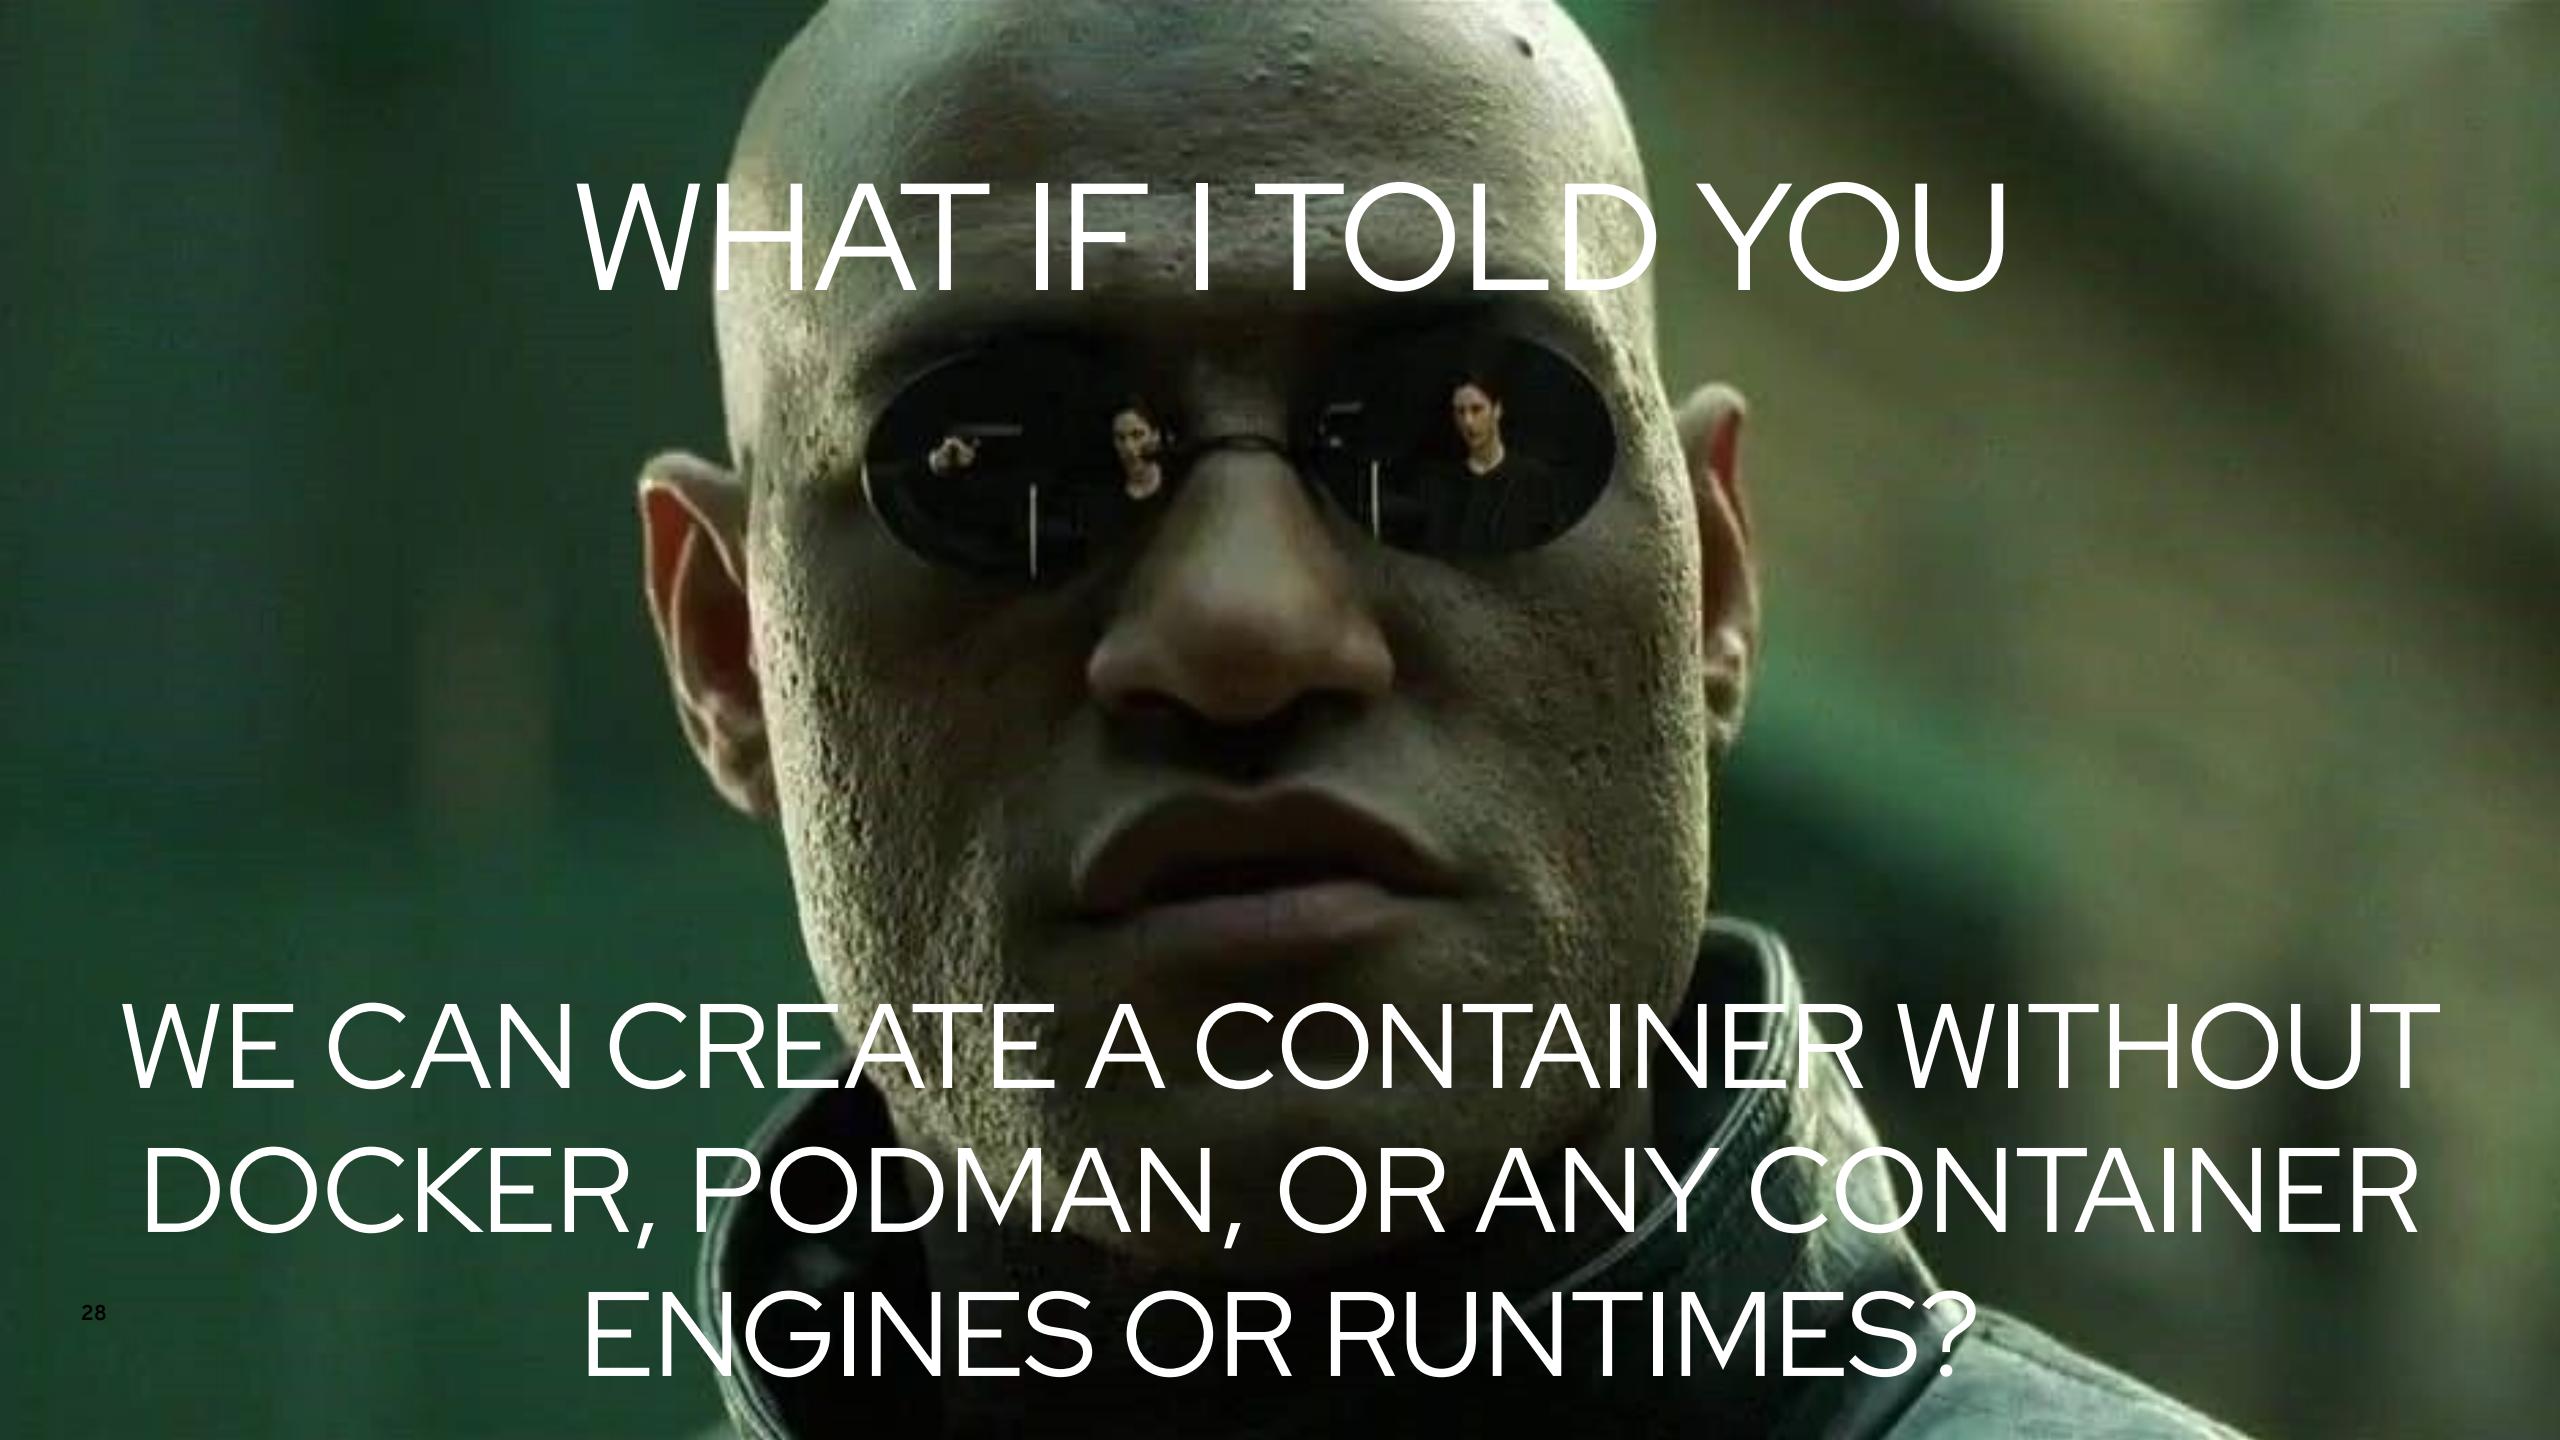
- It all starts with a process
- Then we wrap some boundaries around it that....
 - Partition kernel resources (filesystem, IPC, network, PID, users, more)
 - Limit resource usage
 - Prioritize or de-prioritize CPU share and Disk I/O
 - Control freezing, stopping, starting and checkpointing
 - Account for usage of resources managed above

Namespaces

- It all starts with a process
- Then we wrap some boundaries around it that....
 - Partition kernel resources (filesystem, IPC, network, PID, users, more)
 - Limit resource usage
 - Prioritize or de-prioritize CPU share and Disk I/O
 - Control freezing, stopping, starting and checkpointing
 - Account for usage of resources managed above

CGroups

- It all starts with a process
- Then we wrap some boundaries around it that....
 - Partition kernel resources (filesystem, IPC, network, PID, users, more)
 - **Limit** resource usage
 - **Prioritize** or de-prioritize CPU share and Disk I/O
 - **Control** freezing, stopping, starting and checkpointing
 - **Account** for usage of resources managed above



WHAT IF I TOLD YOU

WE CAN CREATE A CONTAINER WITHOUT
DOCKER, PODMAN, OR ANY CONTAINER
ENGINES OR RUNTIMES?

Technically Speaking....

```
# Creates new linux namespaces for a process (shell) and starts it
unshare --mount --uts --ipc --net --pid --fork --user --map-root-user chroot ${CONTAINER_ROOT} /bin/sh

# mount up the kernel stuff
mount -t proc none /proc
mount -t sysfs none /sys
mount -t tmpfs none /tmp

# OUTSIDE THE CONTAINER!!!
sudo ip link add vethlocal type veth peer name vethNS
sudo ip link set vethlocal up
sudo ip link set vethNS up
sudo ps -ef |grep '/bin/sh'
sudo ip link set vethNS netns <pid of /bin/sh>

# Back in the container
ip link

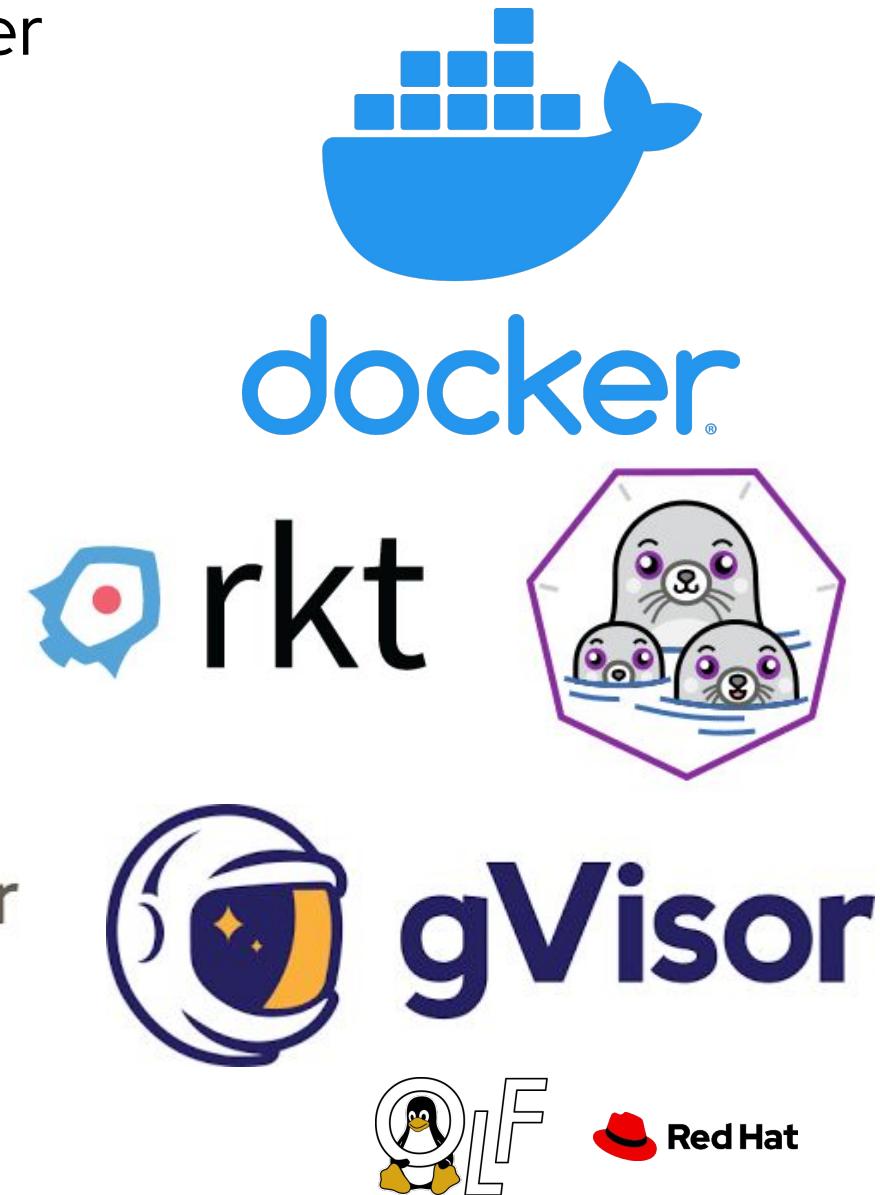
|   0: vethlocal:   
 done
```

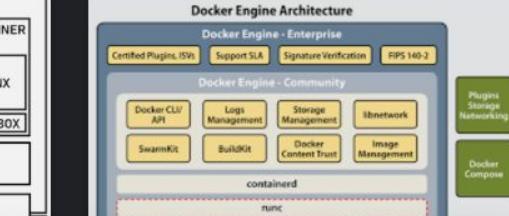
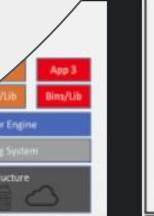
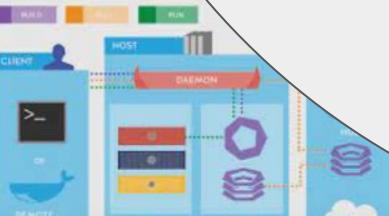
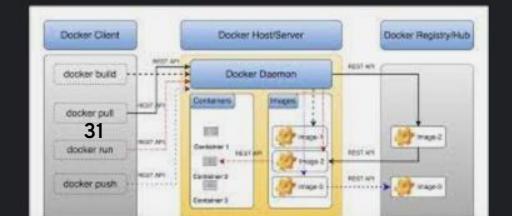
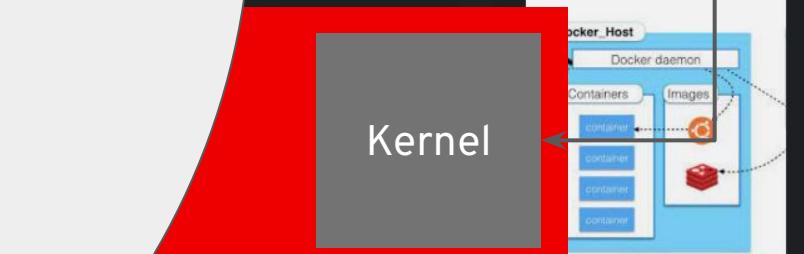
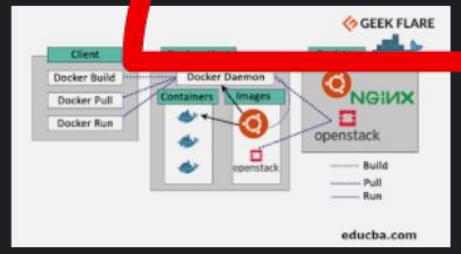
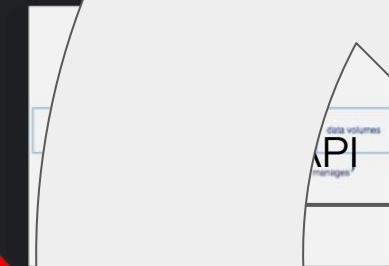
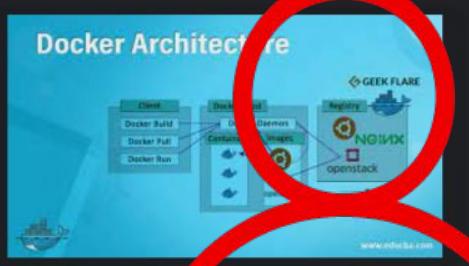
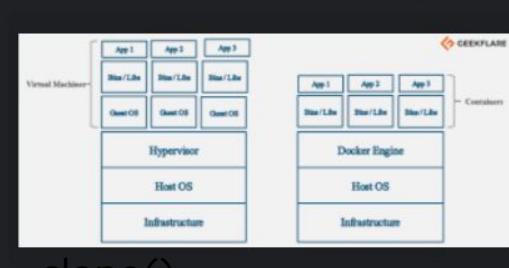
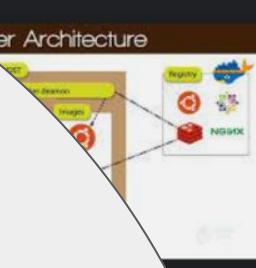
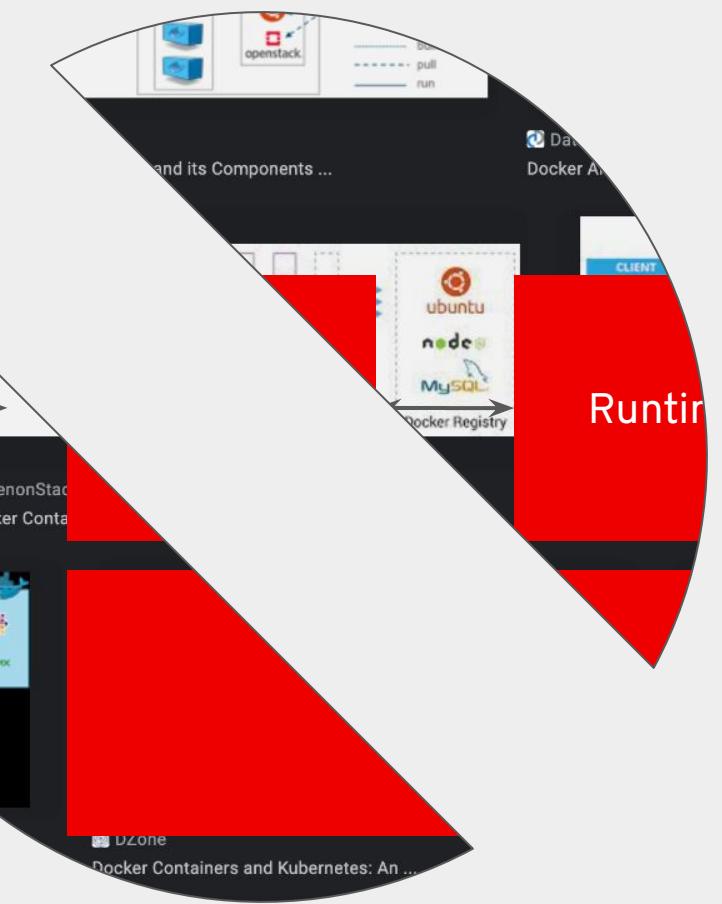
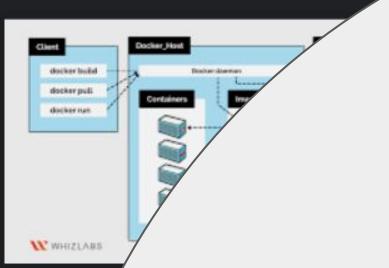
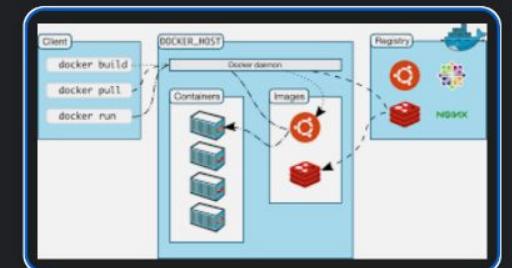
Bonus Lab - If you have time today!

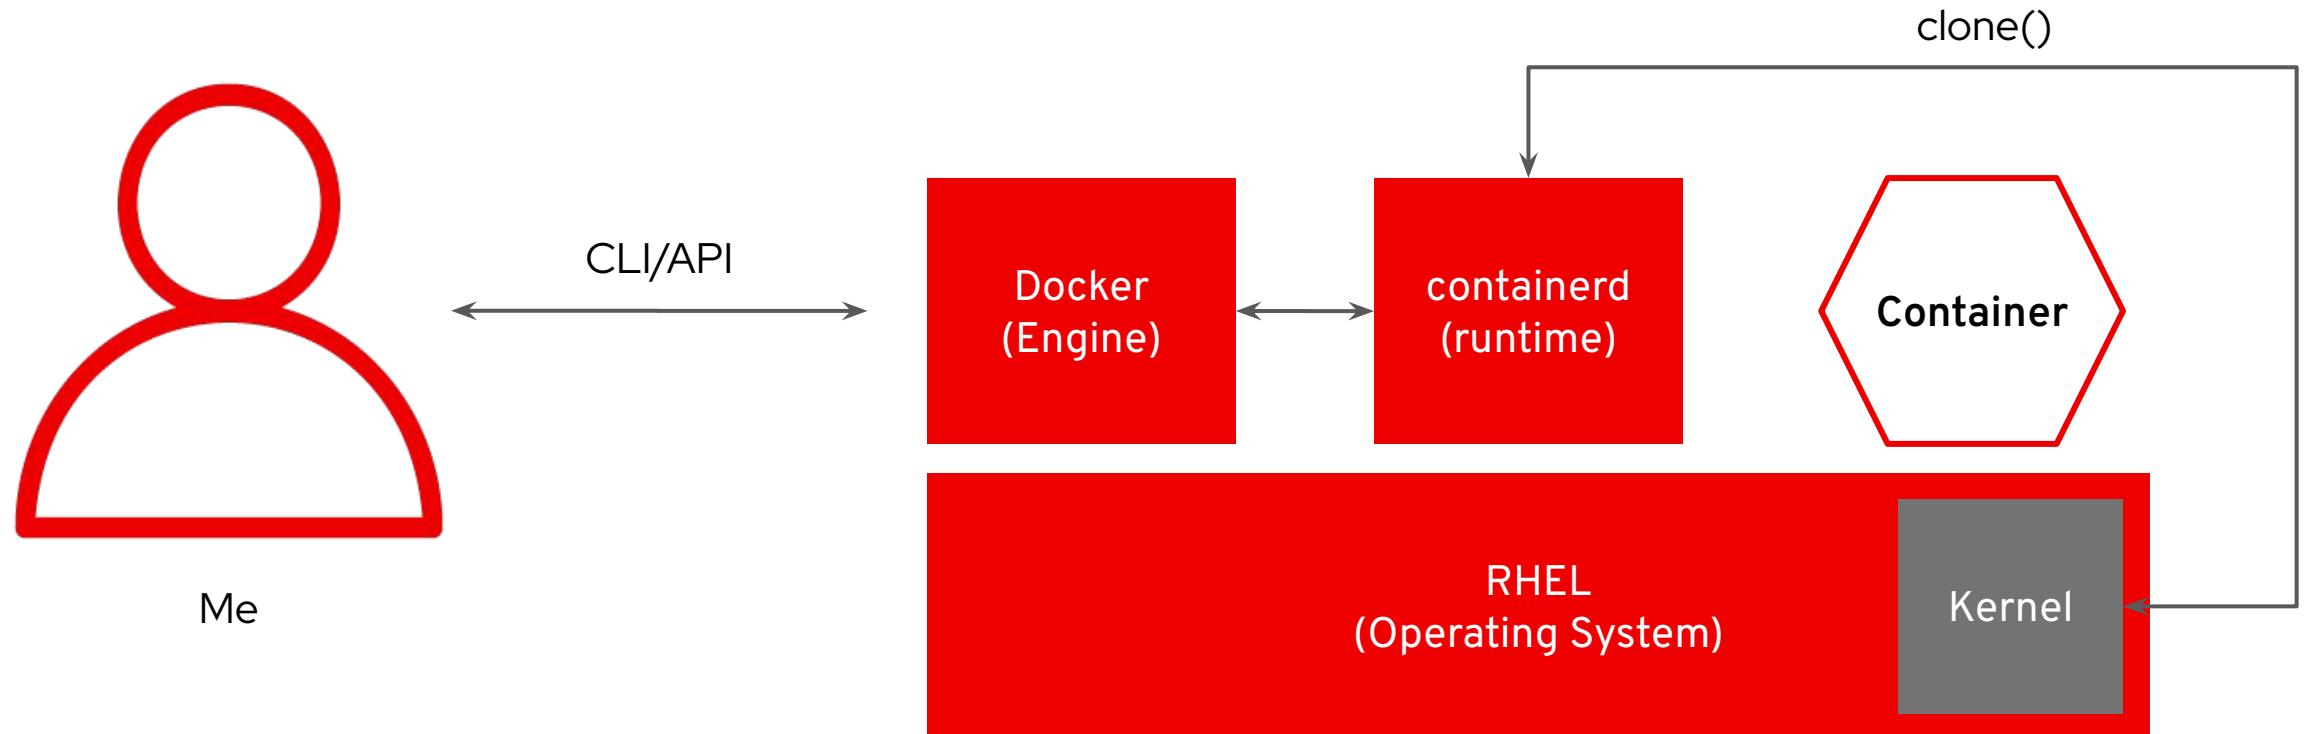


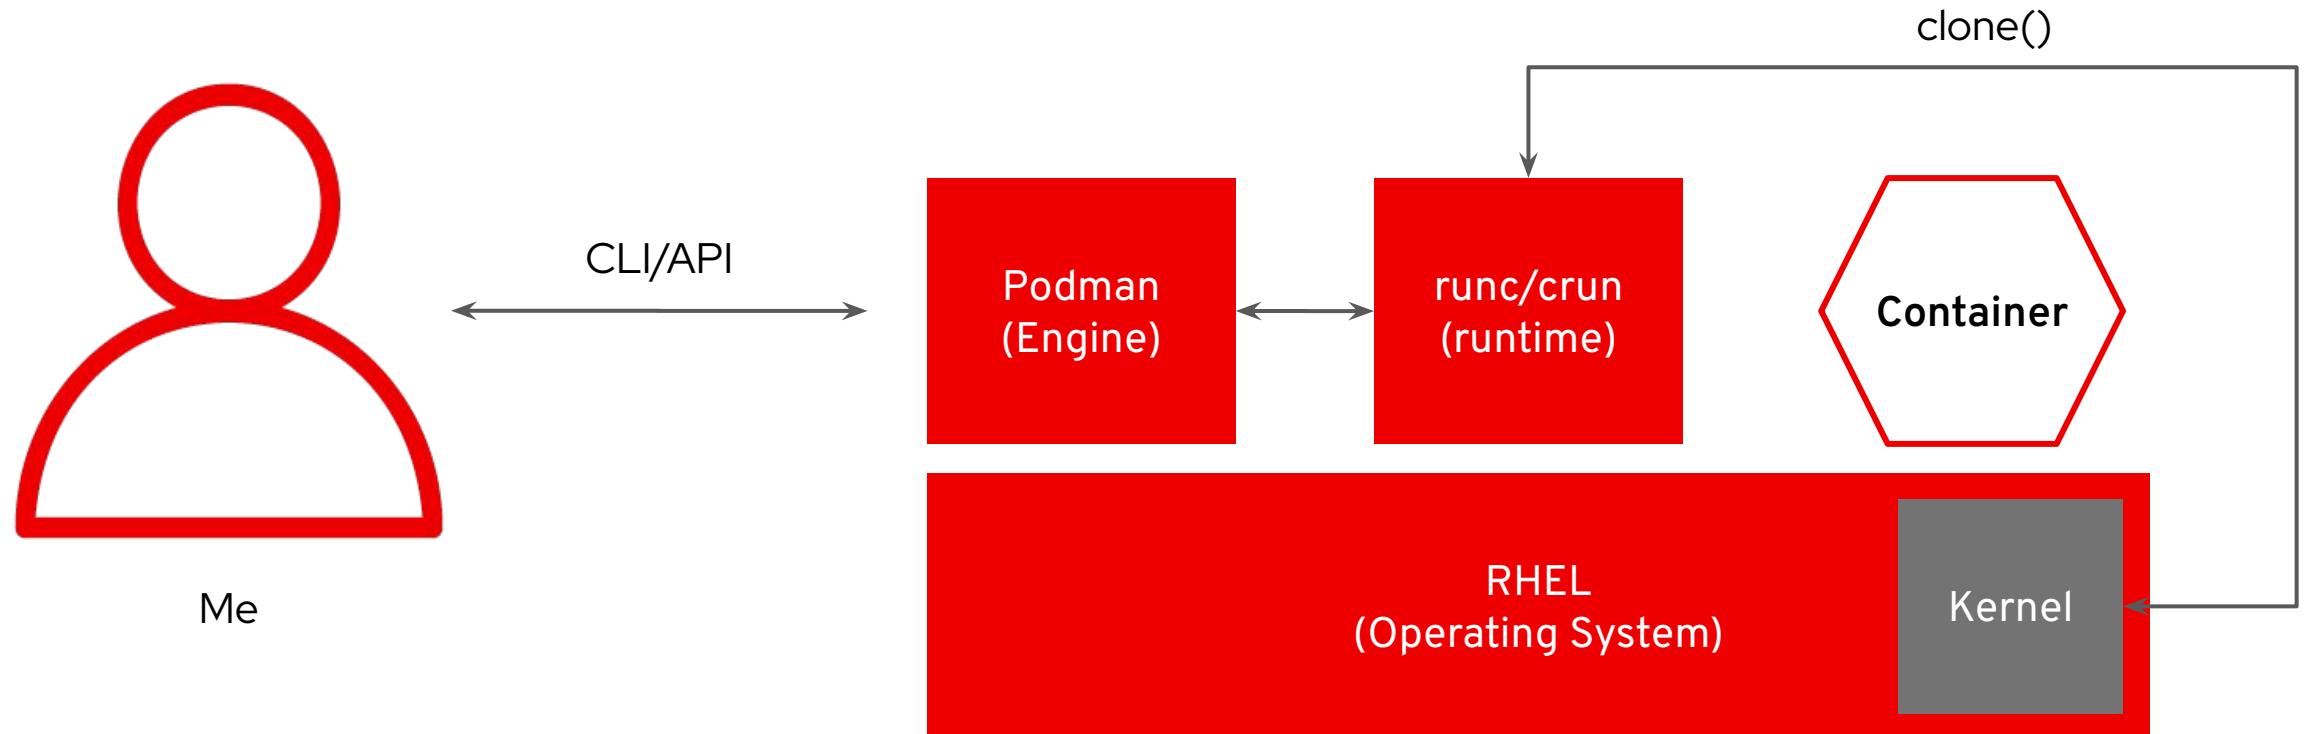
Complexity begets docker

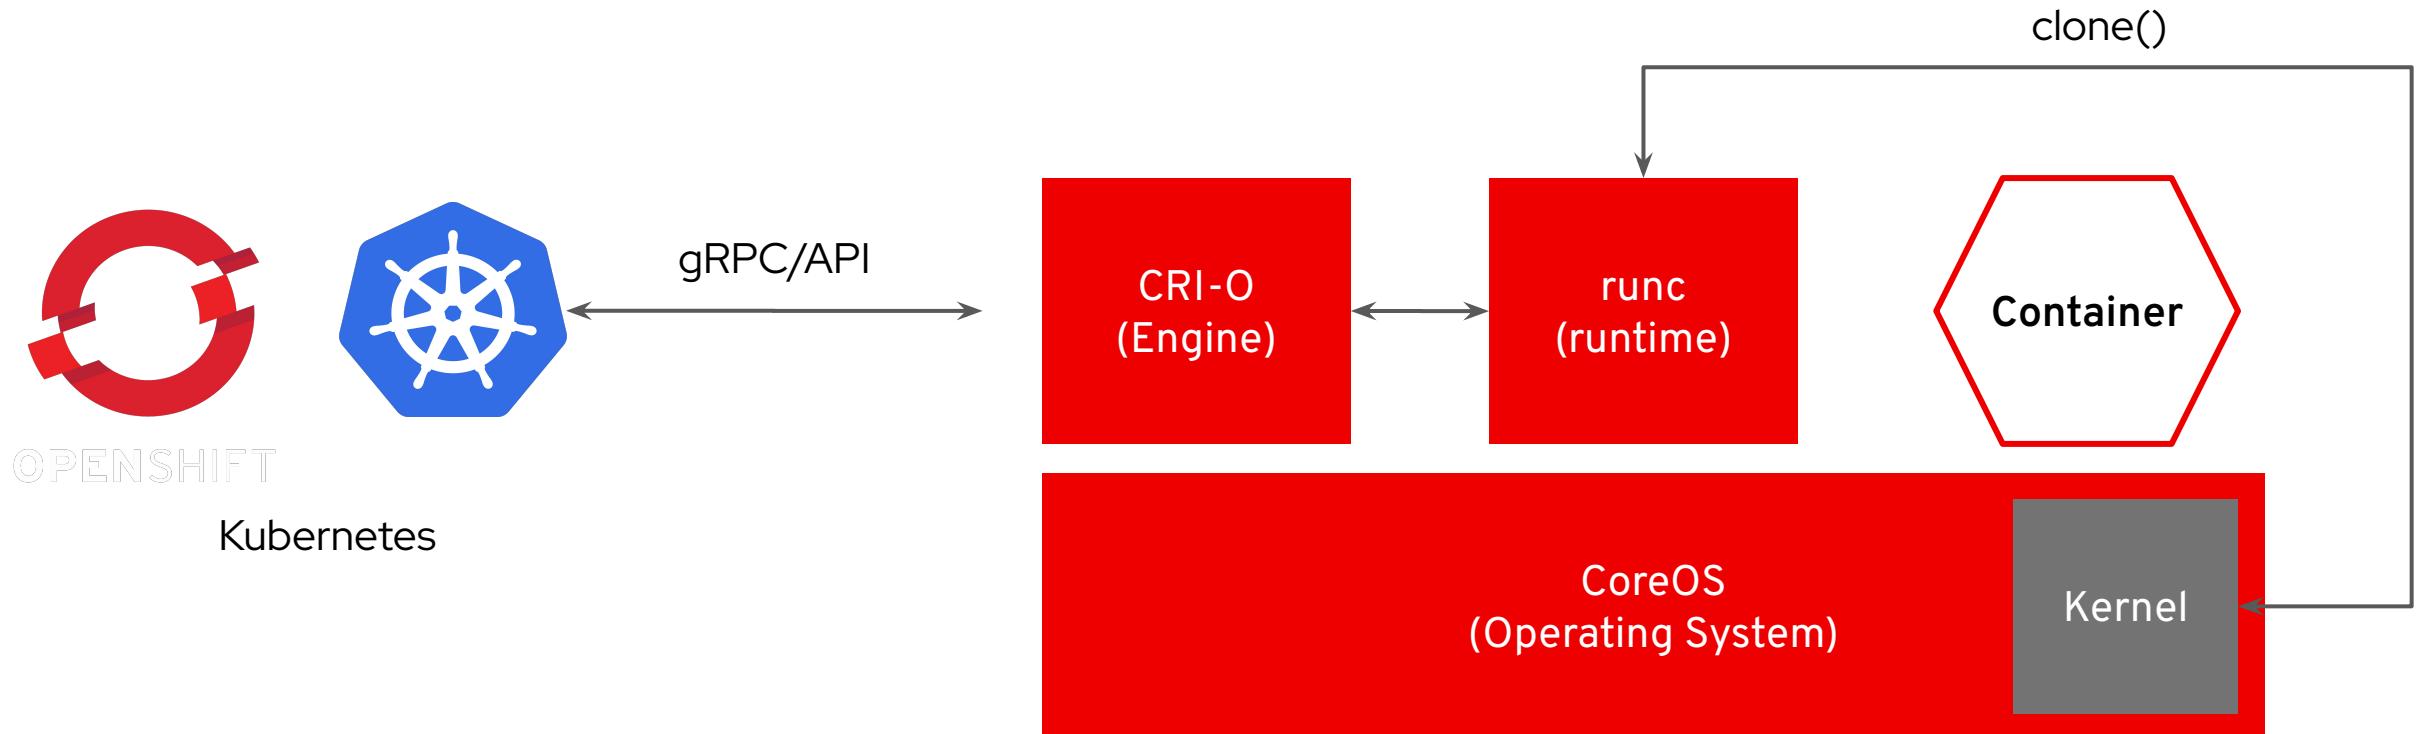
- 2013 - A Star is born!
- Makes all of this hard stuff easy for developers
- New constructs to build, deploy, store, run containers
 - Container Host - brings the kernel to the party
 - Container Runtime - Interface to the kernel
 - Container Engine - Interface to the operator (API/CLI)







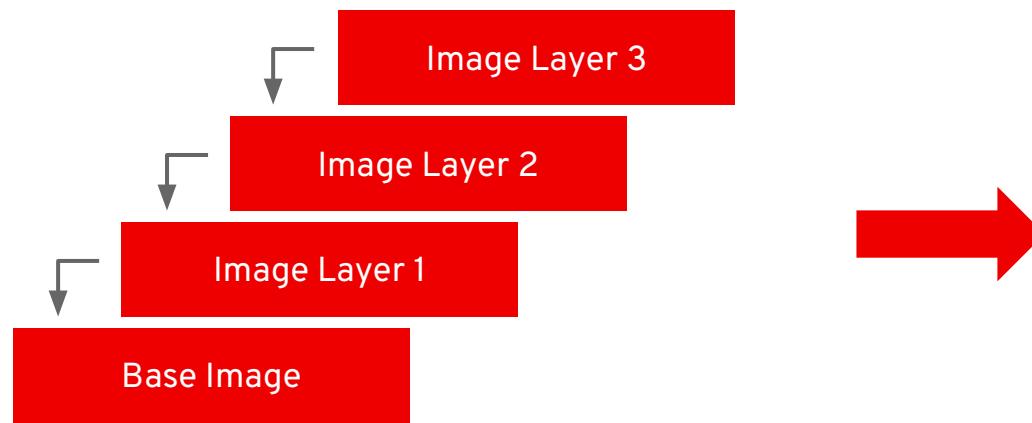




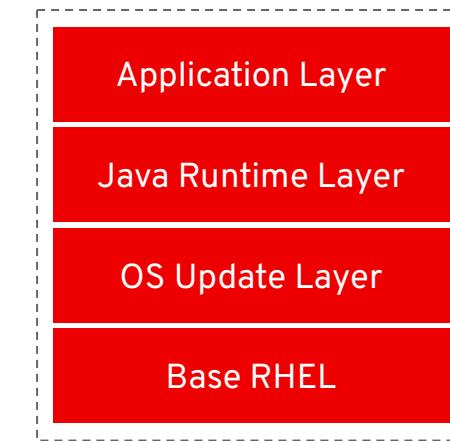
containers are created from
container images



container images are structured in layers



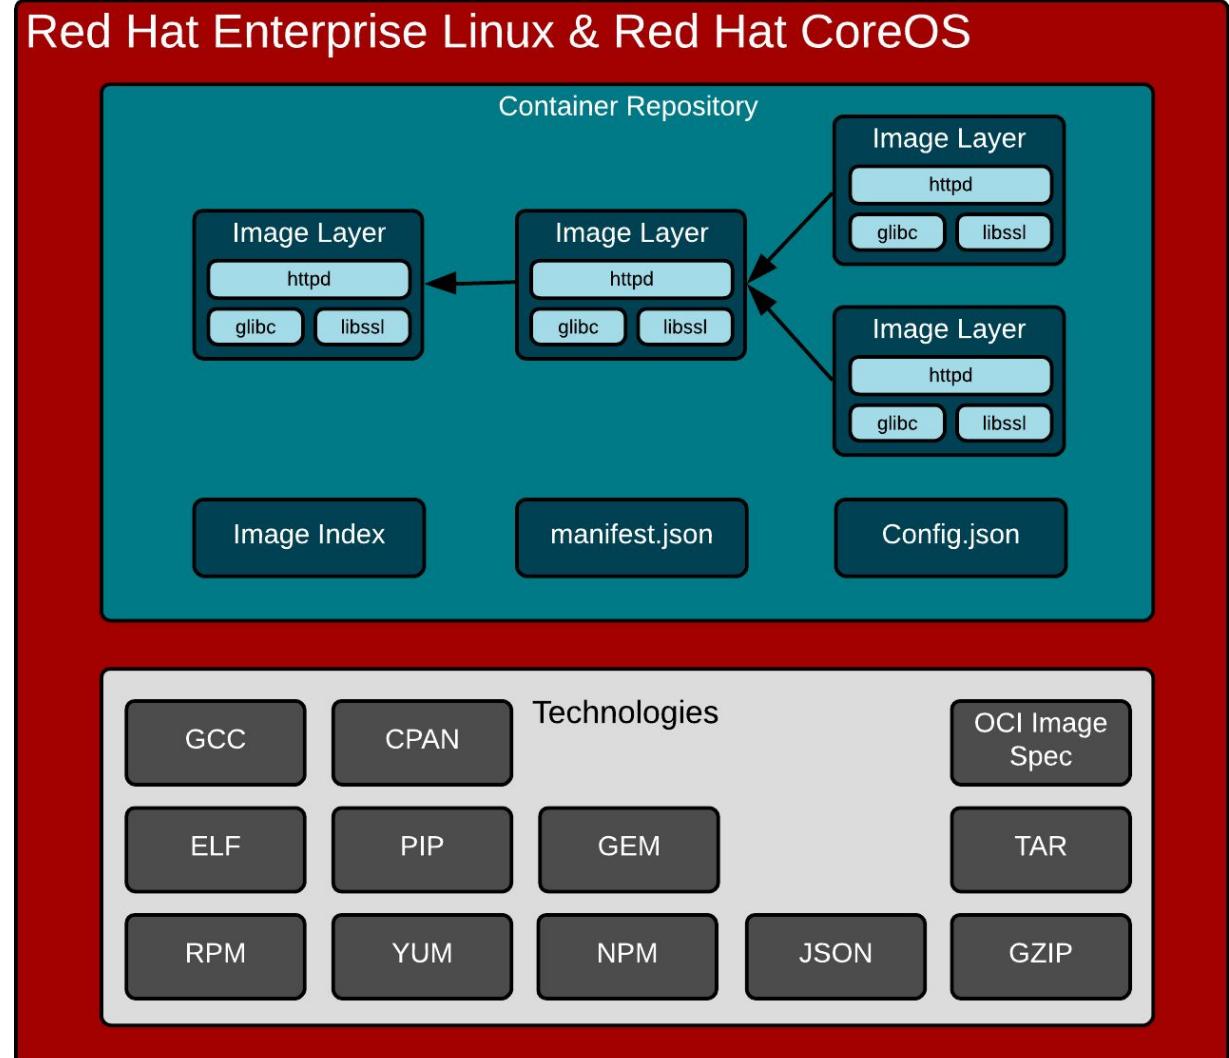
Container Image Layers



Example Container Image

container image parts (OCI Spec)

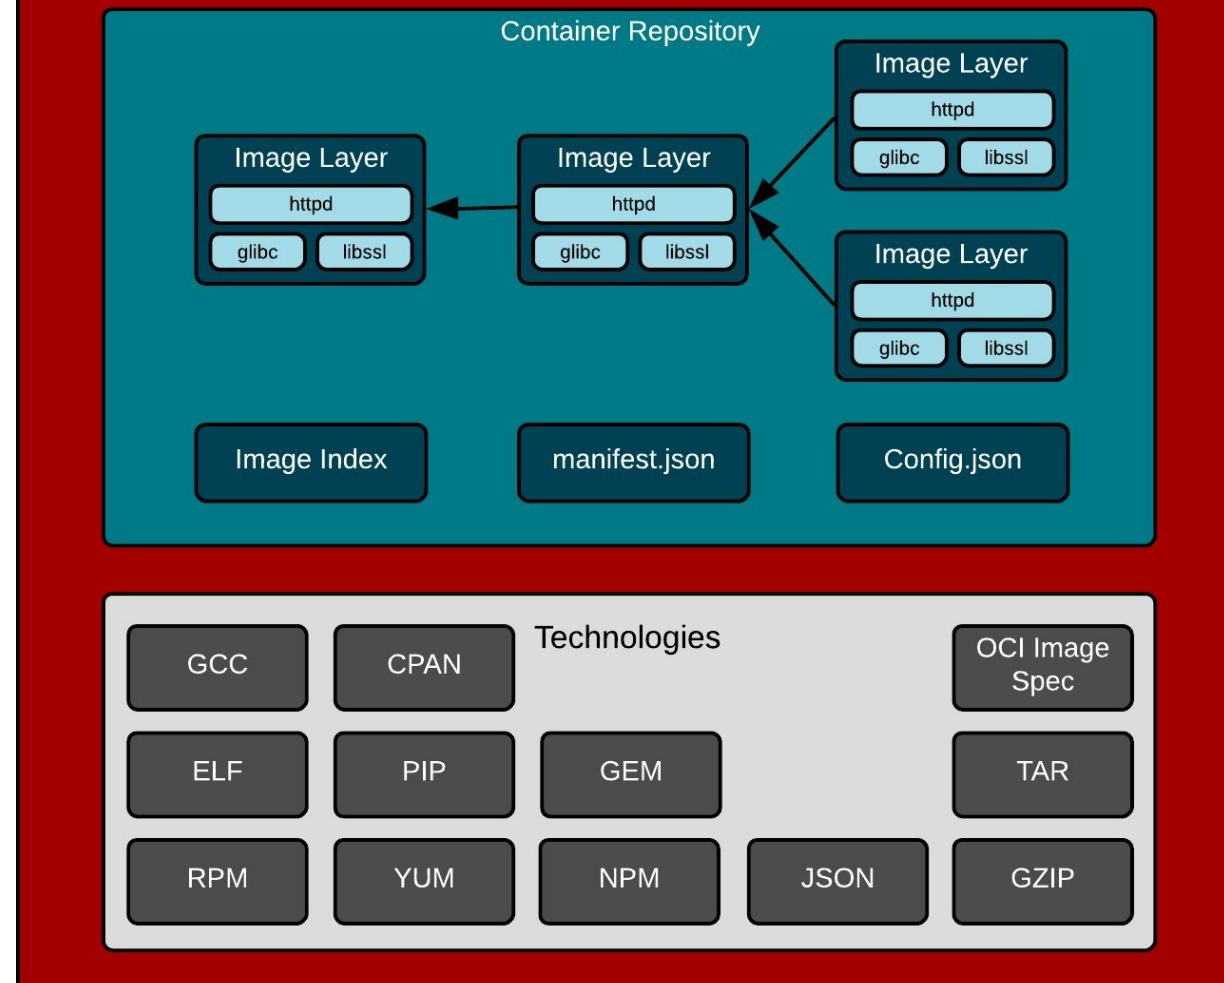
- ▶ Image layers
 - The Bits
- ▶ Image Index
 - Manifest of manifests
 - Multiple architectures
- ▶ Manifest.json
 - Unique hashes of layers
- ▶ Config.json
 - Metadata about the image
- ▶ <https://github.com/opencontainers/runtime-spec/blob/main/spec.md>
- ▶ V1.1.0 was released in July!



container image compatibility

- ▶ OCI Standard guarantees the engine can accept the image
- ▶ Nothing guarantees it is compatible!
- ▶ red.ht/2X9oq0i

Red Hat Enterprise Linux & Red Hat CoreOS



anatomy of a Dockerfile/Containerfile

```
FROM registry.access.redhat.com/ubi8/ubi
```

- 1 Inherit from a base image

```
ENV foo=text
```

- 2 Parameters as environment variables

```
RUN dnf install -y java-11-openjdk
```

- 3 Install dependencies (tooling from base image)

```
ADD my-app.jar /home/my-app.jar
```

- 4 Add your app as a new Layer

```
EXPOSE 8080
```

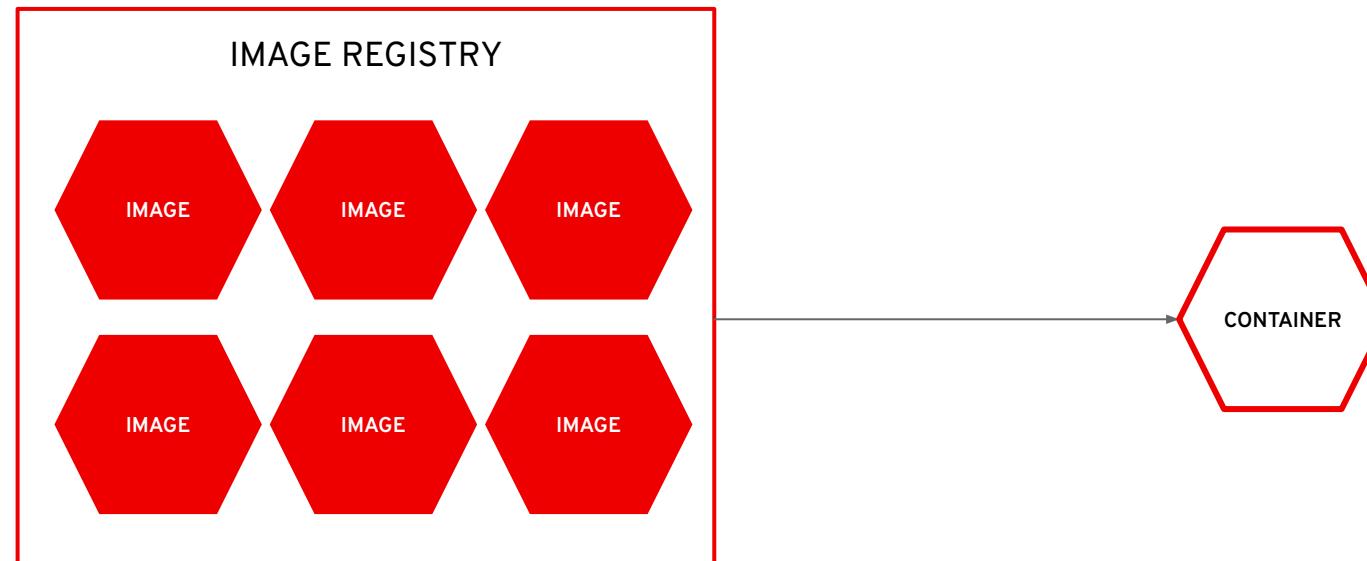
- 5 Expose the port your app will use

```
CMD java -jar /home/my-app.jar
```

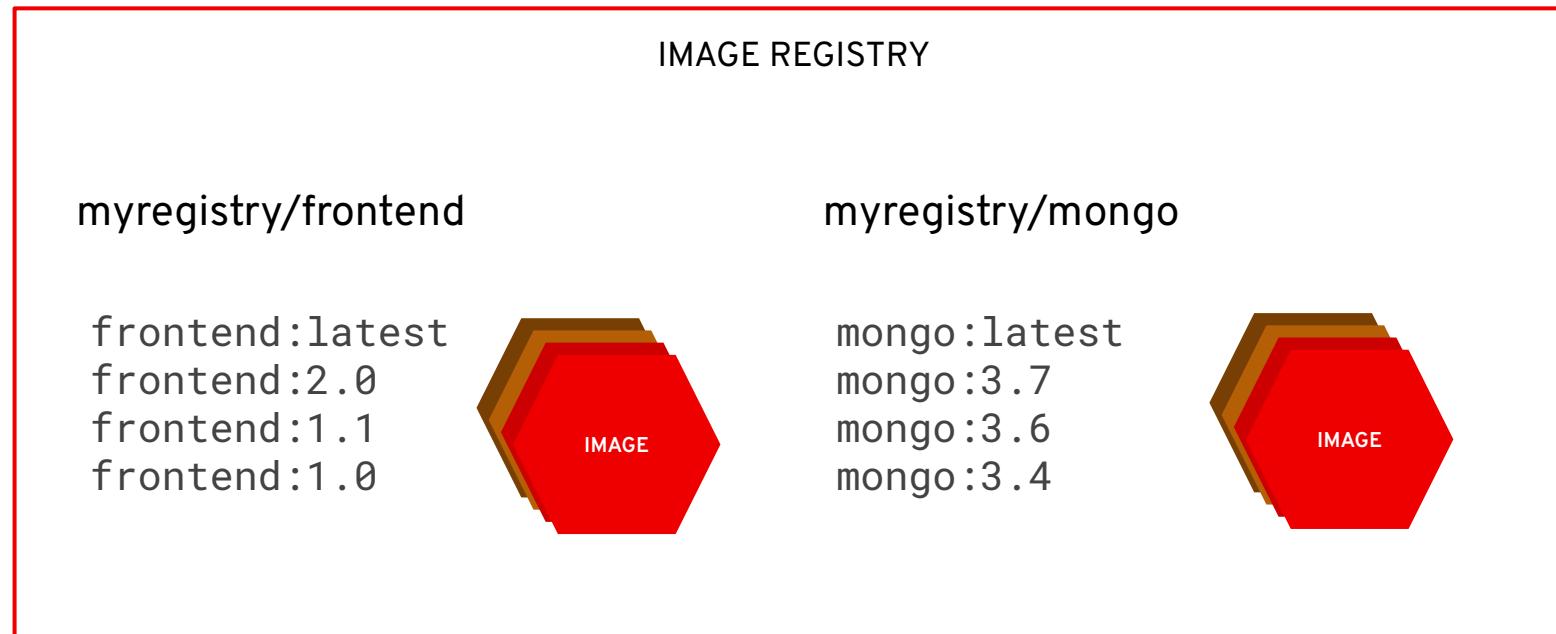
- 6 Run the app

Example for Java app

container images are stored in an image registry



an image repository contains all versions of an image in the image registry



Registry Command Syntax

Command:

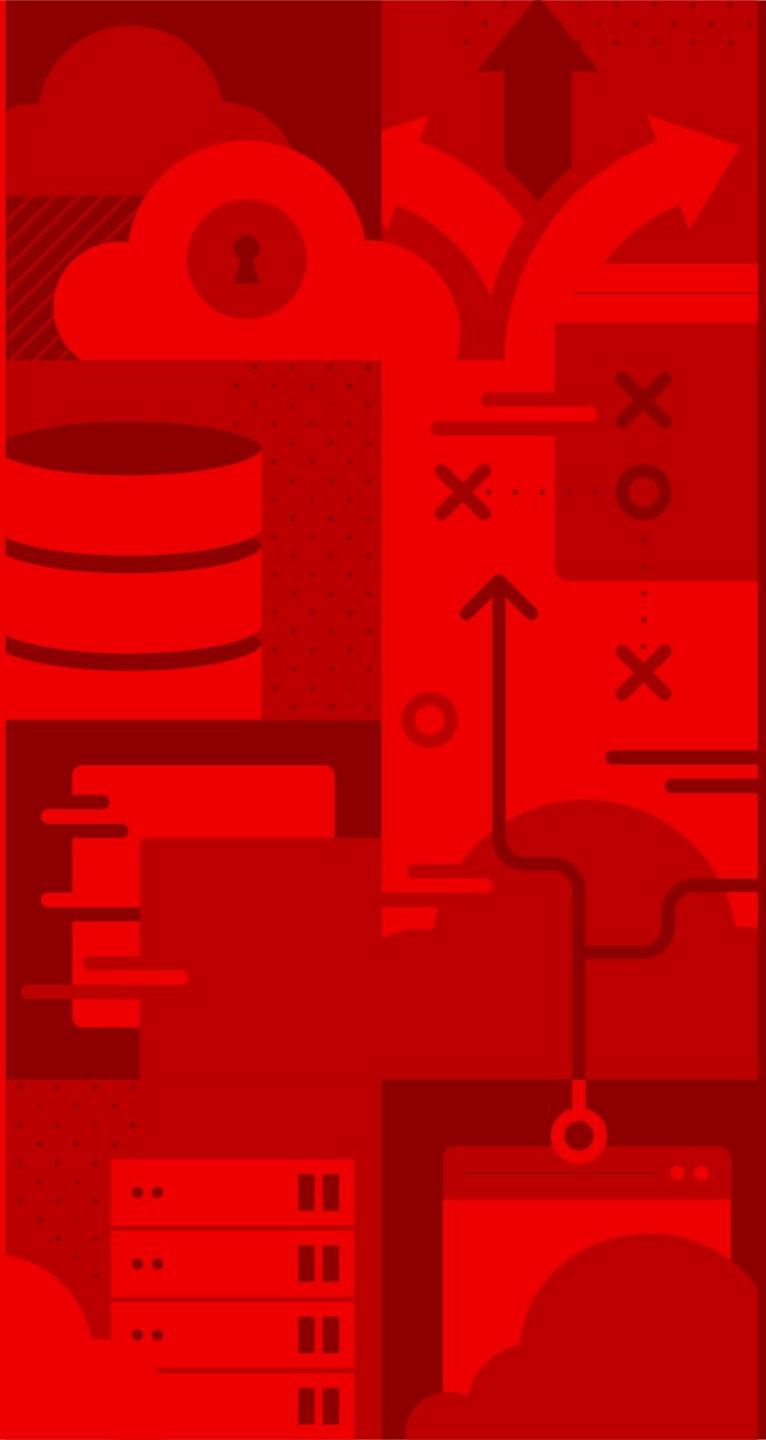
```
docker pull registry.access.redhat.com/rhel7/rhel:latest
```

Decomposition:

```
access.registry.redhat.com / rhel7 / rhel : latest
```

Generalization:

```
Registry Server / namespace / repo : tag
```



What is Kubernetes?

What is Kubernetes?

- “An open-source system for automating deployment, scaling, and management of containerized applications” (<https://kubernetes.io/>)
- Layman terms – Kubernetes is a scheduler for containers (and other things!)
- Abstracts away the details of infrastructure
- “Kubernetes is the new kernel.” – Jessie Frazelle



History of Kubernetes

- Came from Google (announced in 2014)
- Influenced by Google's own cluster manager (Borg/Omega)
- Greek word for Helmsman or Governor
 - The captain of the container ship
- Written in GOLANG
- Google and the Linux foundation formed the CNCF (Cloud Native Computing Foundation) and donated the first project (2015)



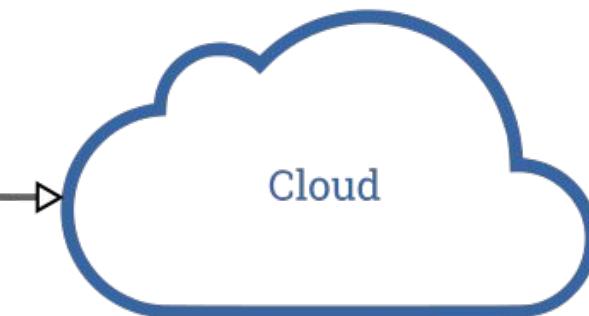
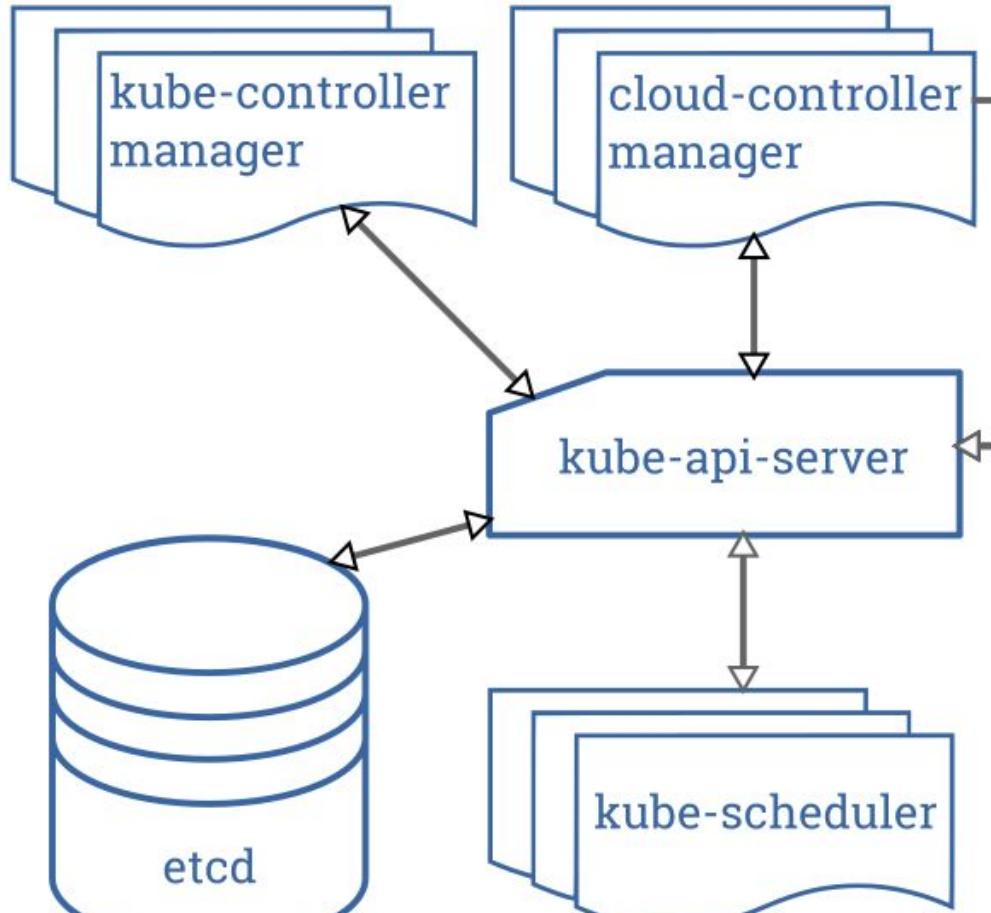
Is Kubernetes the right tool for the job?

“You should be using tools and platforms that enable you to get business logic to production as fast as safe and as secure as possible”

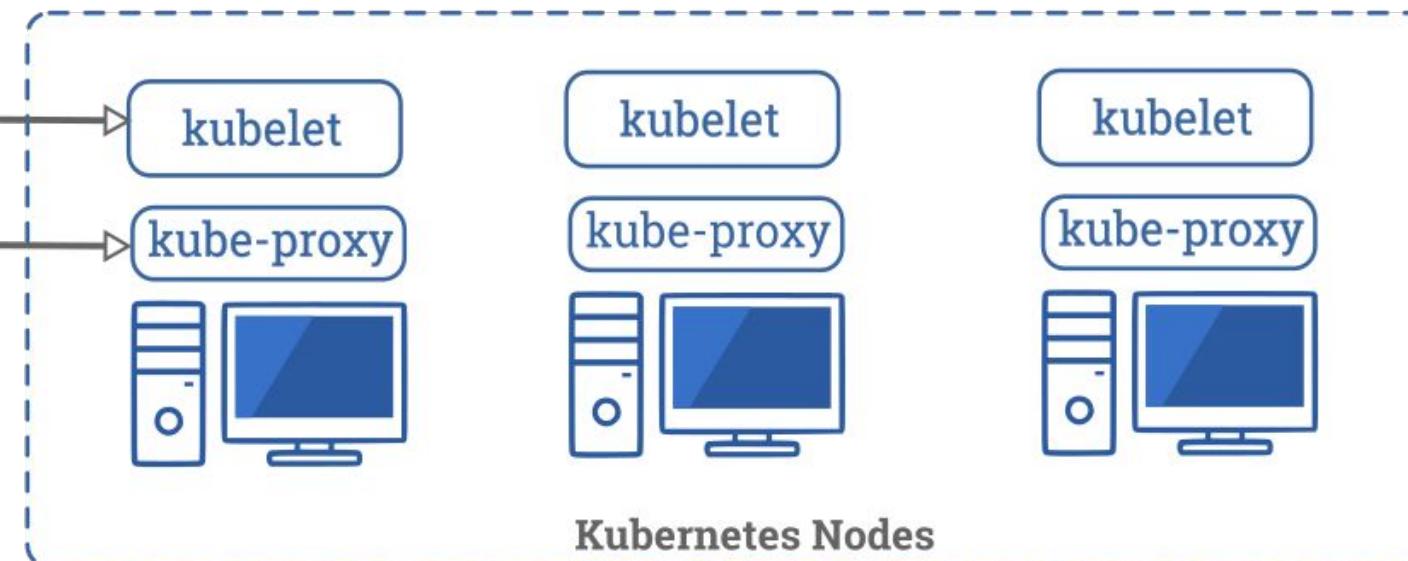
- Me

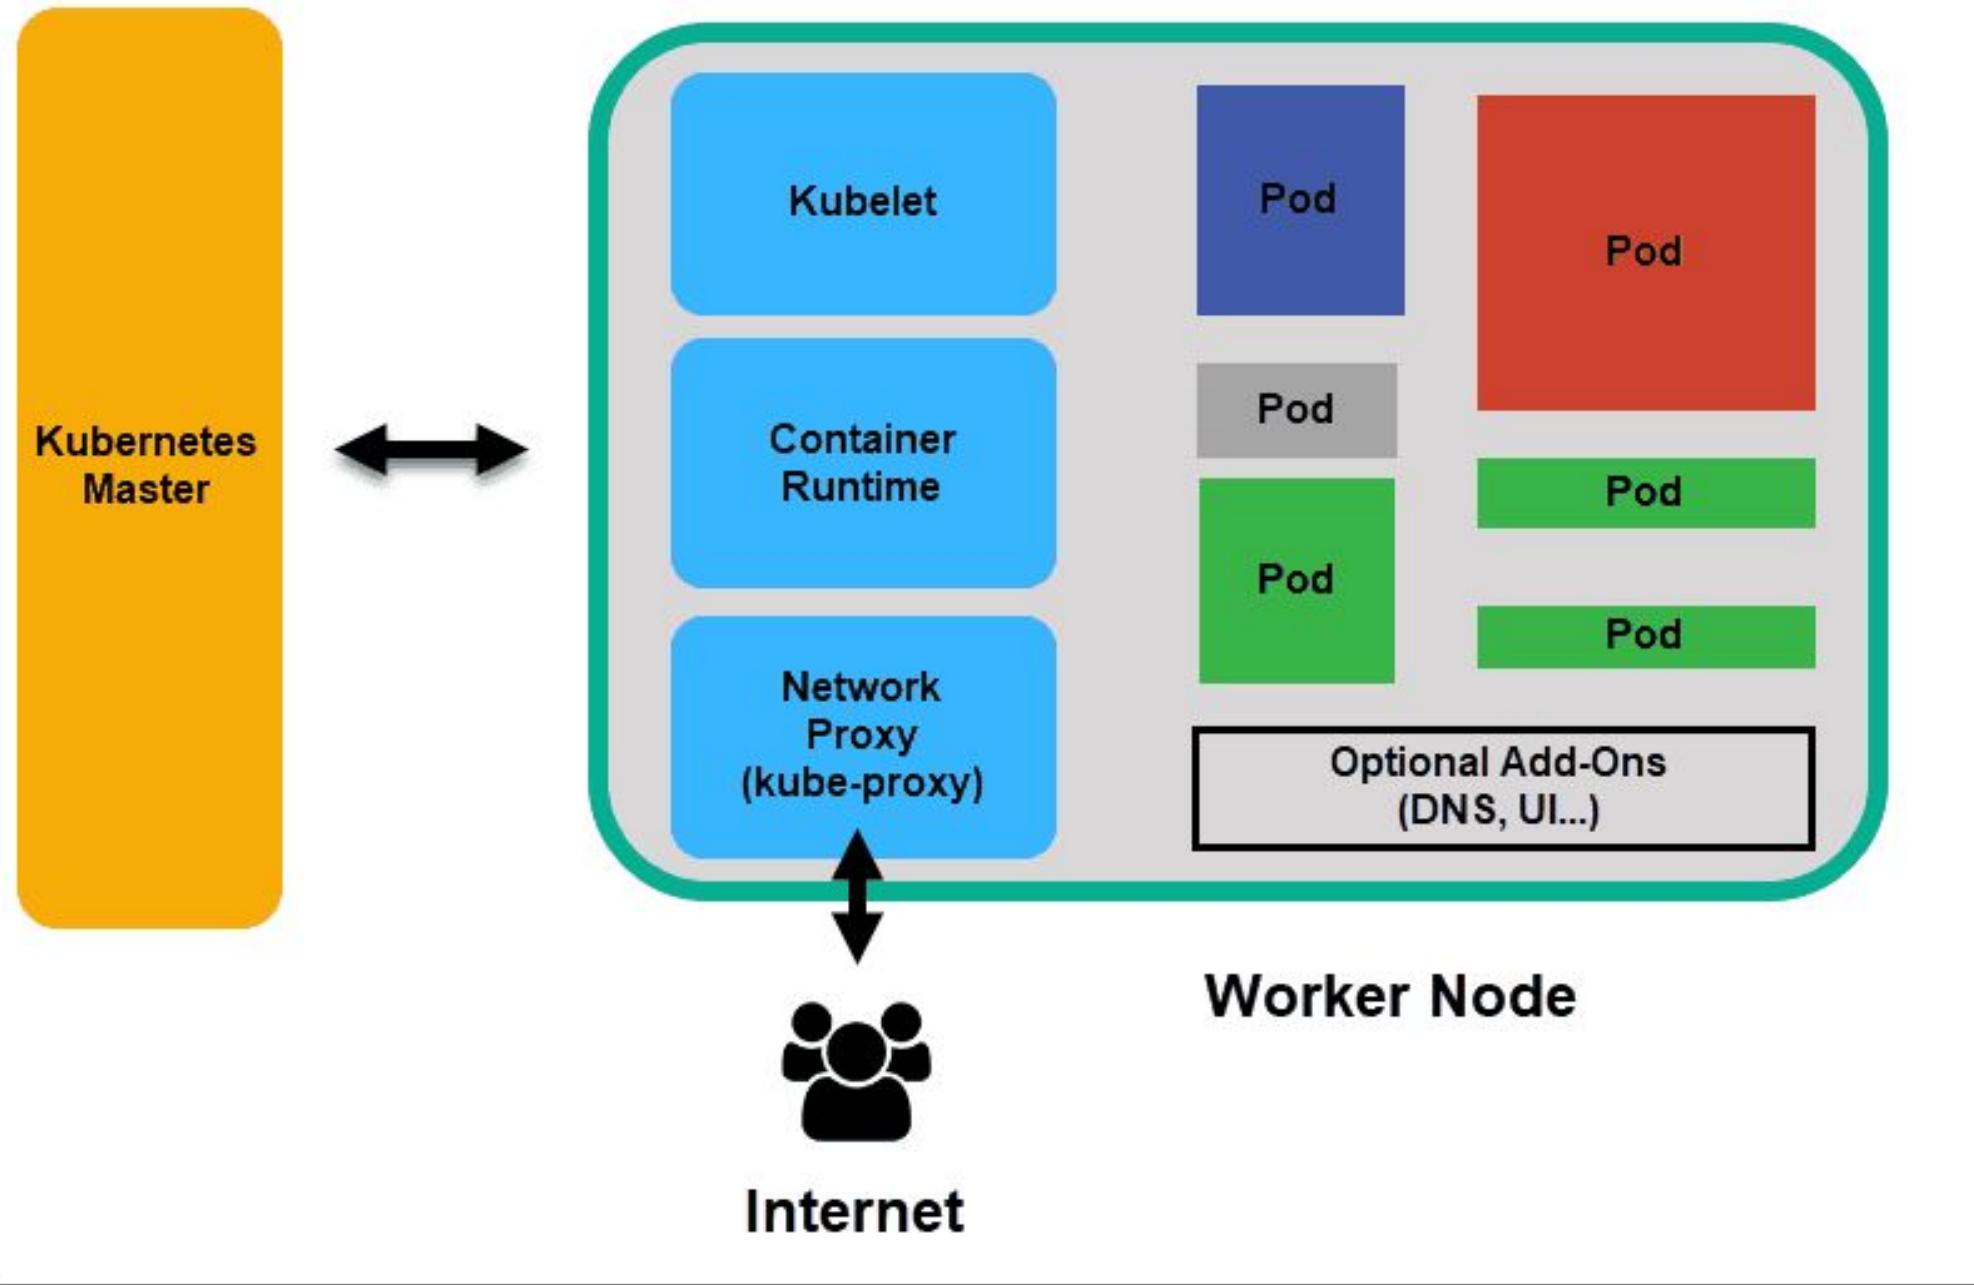
Hint: Kubernetes is a great tool for this!

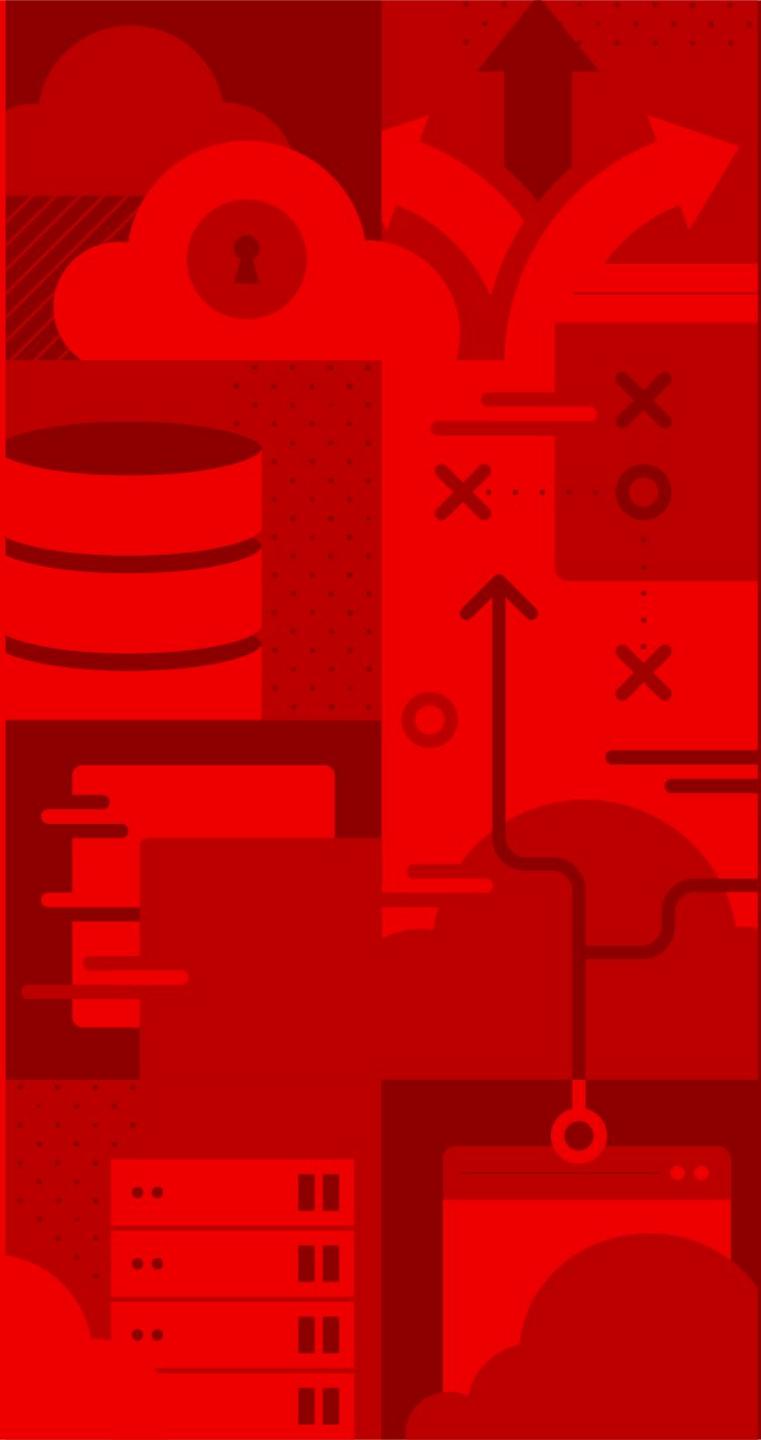
Kubernetes Control Plane



Kubernetes Nodes







Kubernetes basic building blocks

Declarative API - The unsung hero

The Control Plane's Declarative API is the most underrated and important part of Kubernetes

- Inception: Orchestrator of containers
- ~~Today~~ Yesterday: Orchestrator of containers and container adjacent things
- ~~Tomorrow~~ Today: Orchestrator of Orchestrators and Clouds
- Tomorrow: Orchestrator of WebAssembly (WASM/WASI) / Anything!

“Kubernetes is Infrastructure as Data” - Kelsey Hightower of Google

“Containers is just how it started, Kubernetes is bound for much more” - Bassam Tabbara of Upbound

Object Model

All Objects have:

- GVK (Group/Version/Kind)
- Metadata
- “spec” - This is what I want to happen
- “status” - This is what the actual state is

**Objects are described in YAML and converted to JSON when sent to the API server.*

Resource definition

```
status:  
  availableReplicas: 2  
  conditions:  
  - lastTransitionTime:  
    lastTransitionTime: "2023-01-10T12:00:00Z"  
status:  
  containerStatuses:  
  - image: nginx:1.14.22  
    imageID: ""  
    lastState: {}  
    name: nginx  
    ready: false  
    restartCount: 0  
    state:  
      waiting:  
        message: Back-off pulling image "nginx:1.14."  
        reason: ImagePullBackOff  
  replicas: 2  
  updatedReplicas: 2
```

Group: apps Version: v1beta
Kind: Deployment

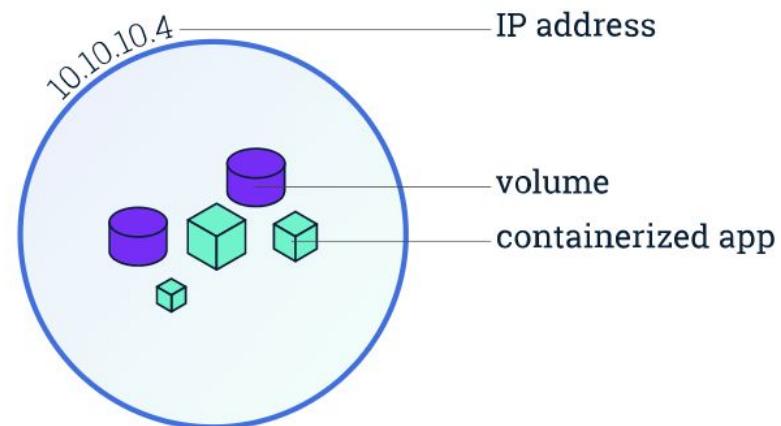


Pod

A pod is the atomic unit of an application in Kubernetes.

One (or more) containers that share:

- Networking (IP address)
- Linux namespace
- Storage
- Memory



Volumes / Configmaps / Secrets

Volumes, Configmaps, and Secrets are used to insert data into pods at runtime

- **Volumes** are essentially just a directory on disk
 - Ephemeral or Persistent (PV/PVC)
- **Configmaps** are a declarative way to store and insert configuration data in to pods as volumes
- **Secrets** are just like configmaps only they store their data in RAM and can be encrypted and obfuscated from the host.

Service

A networking construct to abstract the Pod

3 basic types of Services:

- ClusterIP – basic load balancing (internal to cluster only)
- NodePort – Translates ClusterIP:port to NodeIP:port
- LoadBalancer – integrates with 3rd party / external LB

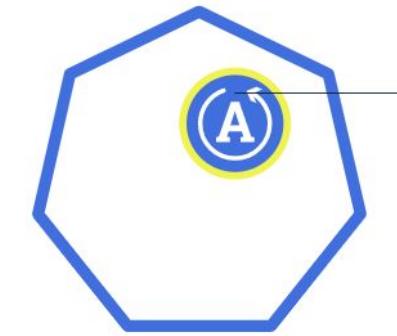


ReplicaSets / Deployments

ReplicaSets - Scales Pods...That's it...

Deployments - Declarative updater for pods and ReplicaSets

- What is running and where?
- What services and resources are available to it?
- Policies for how things should behave (IFTTT logic)
- Rollback / Scaleup / Rollout definitions



Stateful Sets / Daemon Sets / Cron Jobs

Stateful Sets

- A special deployment type for stateful workloads
- Stable and ordered constructs
- Volumes scale out with pods

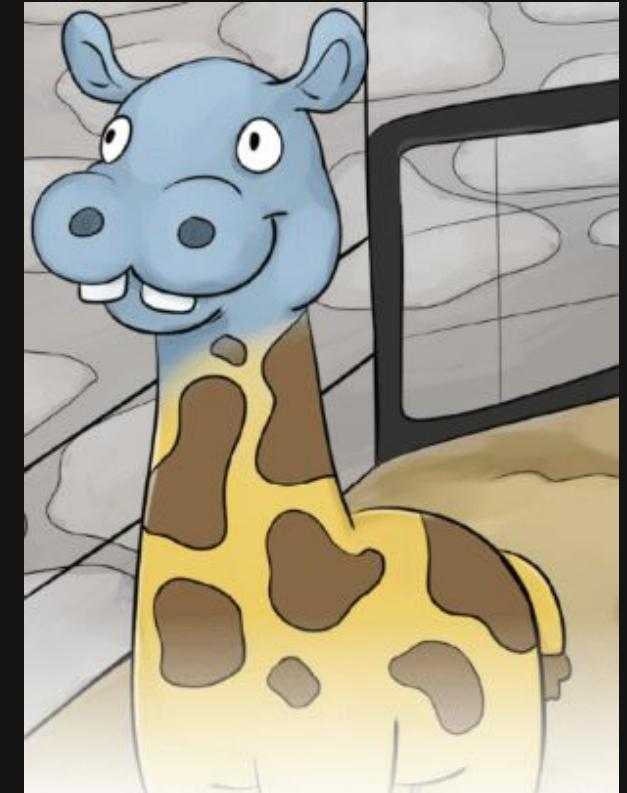
Daemon Sets

- A special deployment type that runs on a defined set of nodes
- As nodes get added that match, the pods scale
- Useful for infrastructure services

Cron Jobs

- Batch processing
- Serverless (like)
- Schedule **when** you want to run a pod

Extending Kubernetes: Custom Resource Definitions



Custom Resource Definitions

- Extend the Kubernetes API with your own GVK
- Inherit the rich high-level API abstractions and data model
- Your resources become first class citizens in the Kubernetes ecosystem
- Operator pattern pairs custom resources with custom controllers to program specific knowledge of your applications into the Kubernetes ecosystem.

Workshop



Please choose the lab code for this session, enter the activation key, and your e-mail address then click **Submit**.

Lab Code: 348f - OCP4 and Container Storage for Admins

Activation Key:

E-Mail Address:

All fields are required.

- **Unless your event organizer says otherwise, we will not e-mail you and your e-mail address will be deleted from this system after this session is over.**
Normally it is only used for tracking this session.
- You may need to refresh this page if you do not see an option for this lab session in the dropdown.
- If you are unsure which lab code to choose or what the activation key is please notify a lab assistant.

Welcome to: OCP4 and Container Storage for Admins

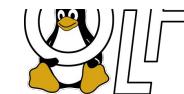
Your assigned lab GUID is **e590**

Let's get started! Please read these instructions carefully before starting to have the best lab experience:

- Save the above **GUID** as you will need it to access your lab's systems from your workstation.
- Consult the lab instructions *before* attempting to connect to the lab environment.
- Please note the following information about your lab environment:
 - Openshift Console: <https://console-openshift-console.apps.cluster-qq8pz.qq8pz.sandbox1506.opentlc.com>
 - Openshift API for command line 'oc' client: <https://api.cluster-qq8pz.qq8pz.sandbox1506.opentlc.com:6443>
 - Download oc client from <http://d3s3zqyaz8cp2d.cloudfront.net/pub/openshift-v4/clients/ocp/stable-4.7/openshift-client-linux.tar.gz>
 - Access the workshop at '<https://dashboard-lab-ocp-cns.apps.cluster-qq8pz.qq8pz.sandbox1506.opentlc.com>'
 - Login with 'kubeadmin' and '**[REDACTED]**'
 - Workshop may not be accessible until rollout finishes shortly.
 - You can access your bastion via SSH:
 - ssh lab-user@bastion.qq8pz.sandbox1506.opentlc.com
 - Make sure you use the username 'lab-user' and the password '**[REDACTED]**' when prompted.
- The following URLs and information will be used in your lab environment. Please only access these links when the lab instructions specify to do so:
 - **Username:** kubeadmin
 - **Password:** Ecu7ShakvS0KtF3Zob
 - Note: The lab instructions may specify other host names and/or URLs.

WARNING: You should only click FORGET SESSION if **requested** to do so by a lab attendant.

[FORGET SESSION](#)



Red Hat

Authorize Access

Service account dashboard-user in project lab-ocp-cns is requesting permission to access your account (kube:admin)

Requested permissions

user:info

Read-only access to your user information (including username, identities, and group membership)

You will be redirected to https://dashboard-lab-ocp-cns.apps.cluster-qq8pz.qq8pz.sandbox1506.opentlc.com/oauth_callback

Allow selected permissions

Deny

The screenshot shows a web-based workshop interface for OpenShift Container Storage. The top navigation bar includes tabs for 'Terminal' (which is active) and 'Console'. The main content area displays the 'Environment Overview' page, which provides an introduction to the OpenShift 4 cluster setup. The page mentions the cluster is running on AWS and uses Rook/Ceph for storage. It also notes the cluster has 3 master nodes, 3 worker nodes, and 1 bastion host. On the left, a sidebar lists various workshop modules: Environment Overview, Installation and Verification, Application Management Basics, Application Storage Basics, MachineSets, Machines, and Nodes, Infrastructure Nodes and Operators, Deploying and Managing OpenShift Container Storage, and OpenShift Log Aggregation. A vertical toolbar on the right offers options to reload the workshop, terminal, or console, and to open them in new tabs. The bottom of the page features footer links for GitHub, LinkedIn, and YouTube.

OpenShift and Container Storage for Admins

WORKSHOP MODULES

Environment Overview

Installation and Verification

Application Management Basics

Application Storage Basics

MachineSets, Machines, and Nodes

Infrastructure Nodes and Operators

Deploying and Managing OpenShift Container Storage

OpenShift Log Aggregation

[GitHub](#) [LinkedIn](#) [YouTube](#)

Environment Overview

Environment Overview

You will be interacting with an OpenShift 4 cluster that is running on Amazon Web Services. During the lab you will also install OpenShift Container Storage, based on Rook/Ceph.

The basics of the OpenShift 4 installation have been completed in advance. The OpenShift cluster is essentially set to all defaults and looks like the following:

- 3 master nodes
- 3 worker nodes
- 1 bastion host

Terminal Console

[~] \$

Reload Workshop
Reload Terminal
Reload Console
Open Workshop
Open Terminal
Open Console

Workshop Instructions

- **Everything is done in browser** - no local commands or installs needed on your laptop
- Works best in Chrome or Firefox
- **Turn off VPN** (we use websockets extensively), **pause AdBlock** for the lab domain (there are no ads)
- Go slow, read, re-read
- Type things (I know it's painful)
- **Ask Questions!**

Lab Request URL:
<https://red.ht/olf-kube>

Password: r3dh4t!

Exercise 1: Environment Overview

What you're working with

Environment Overview

You will be interacting with an OpenShift 4 cluster that is running on Amazon Web Services.

The basics of the OpenShift 4 installation have been completed in advance. The OpenShift cluster is essentially set to all defaults and looks like the following:

- 3 control (master) nodes
- 2 worker nodes
- 1 bastion host

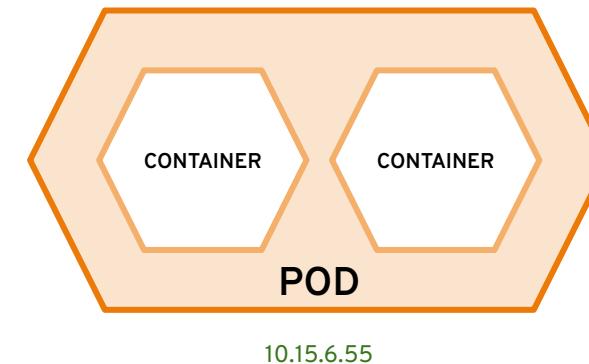
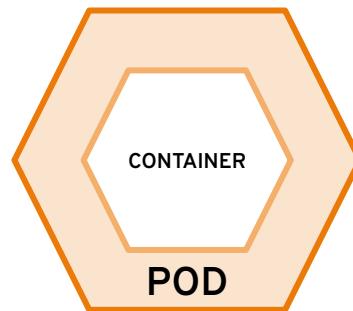
Exercise 2: Installation and Verification

Checking that everything
is OK

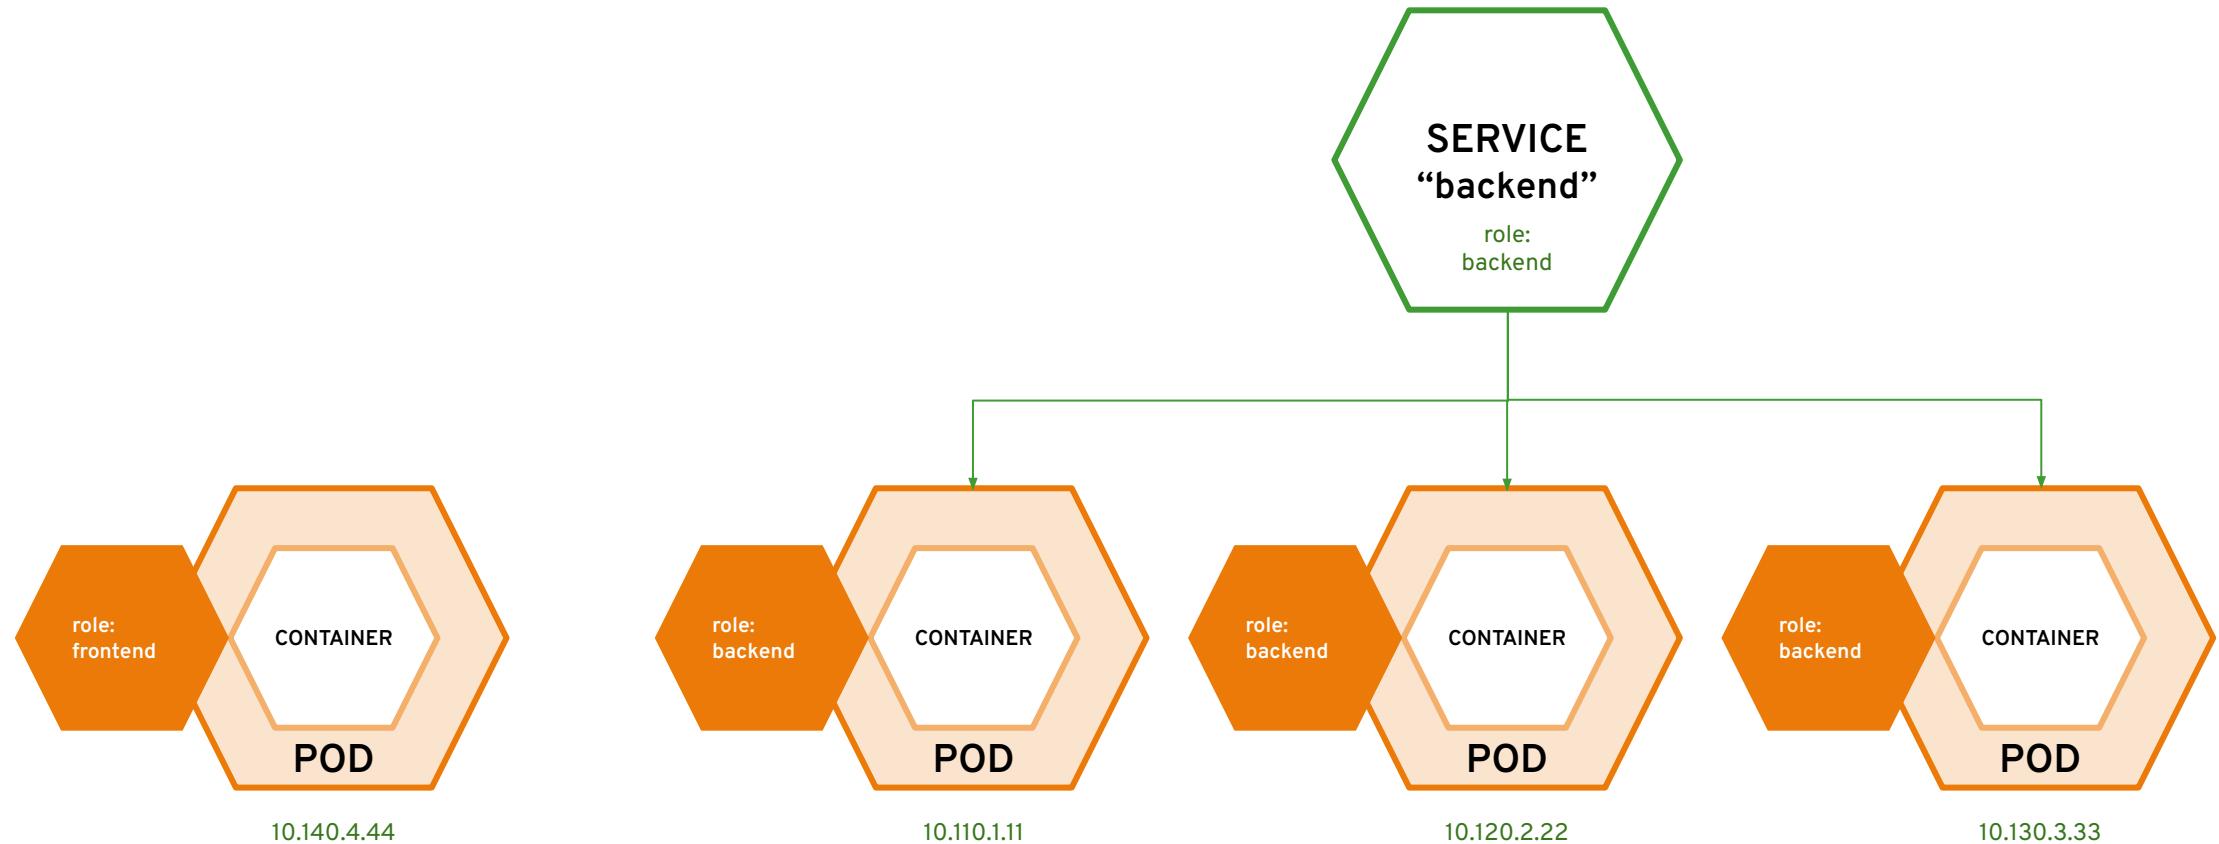
Exercise 3: Application Management Basics

Fundamental Kubernetes
concepts

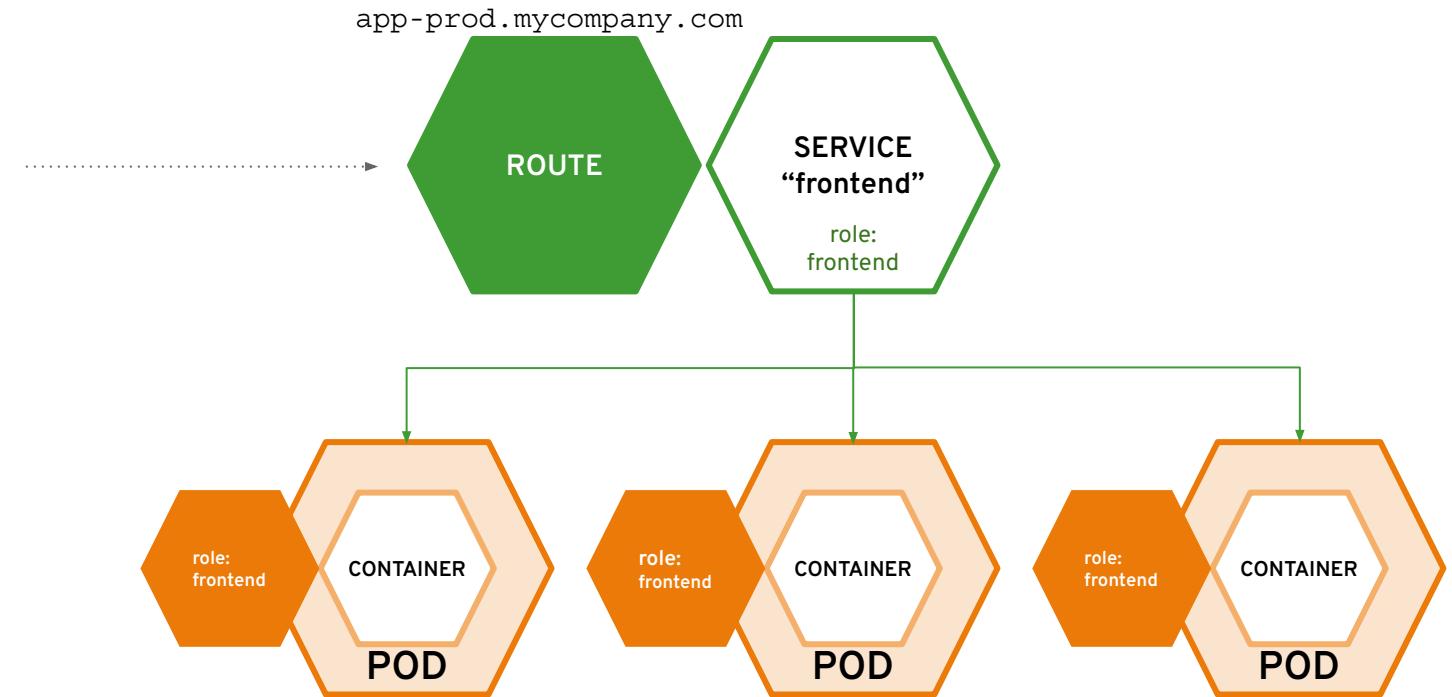
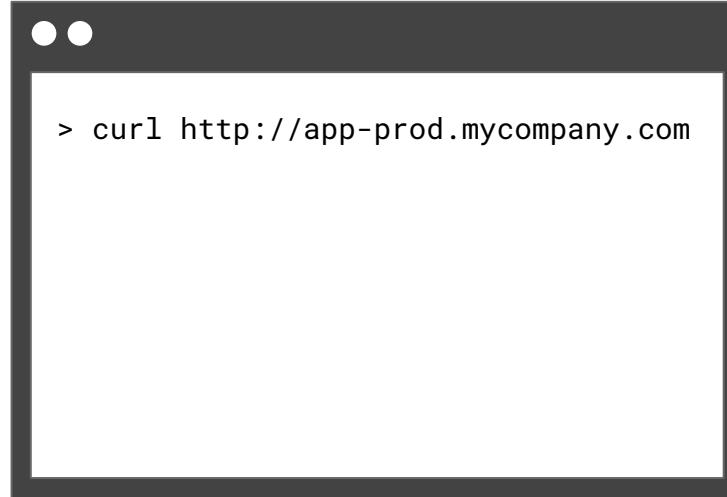
containers are wrapped in pods which are units of deployment and management



services provide internal load-balancing and service discovery across pods



routes make services accessible to clients outside the environment via real-world urls

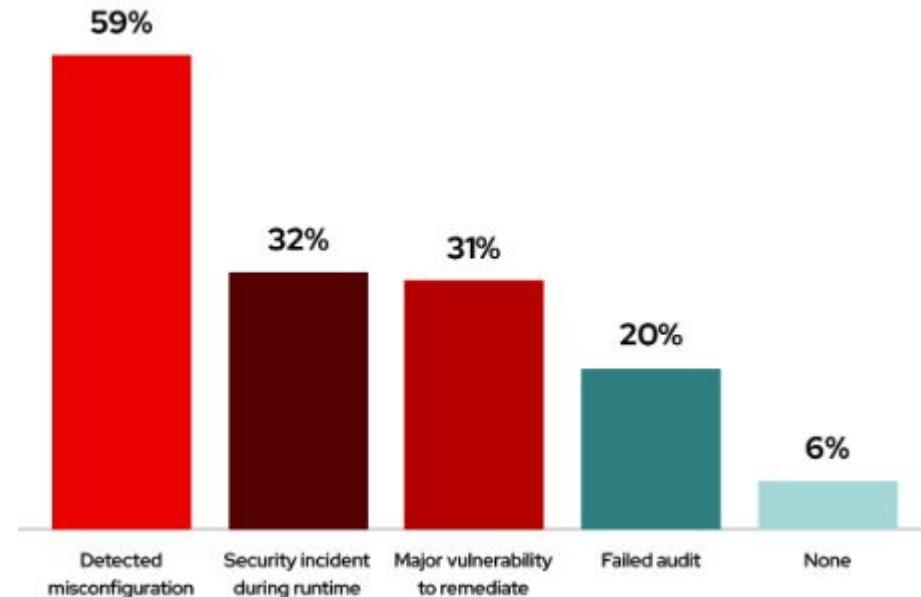




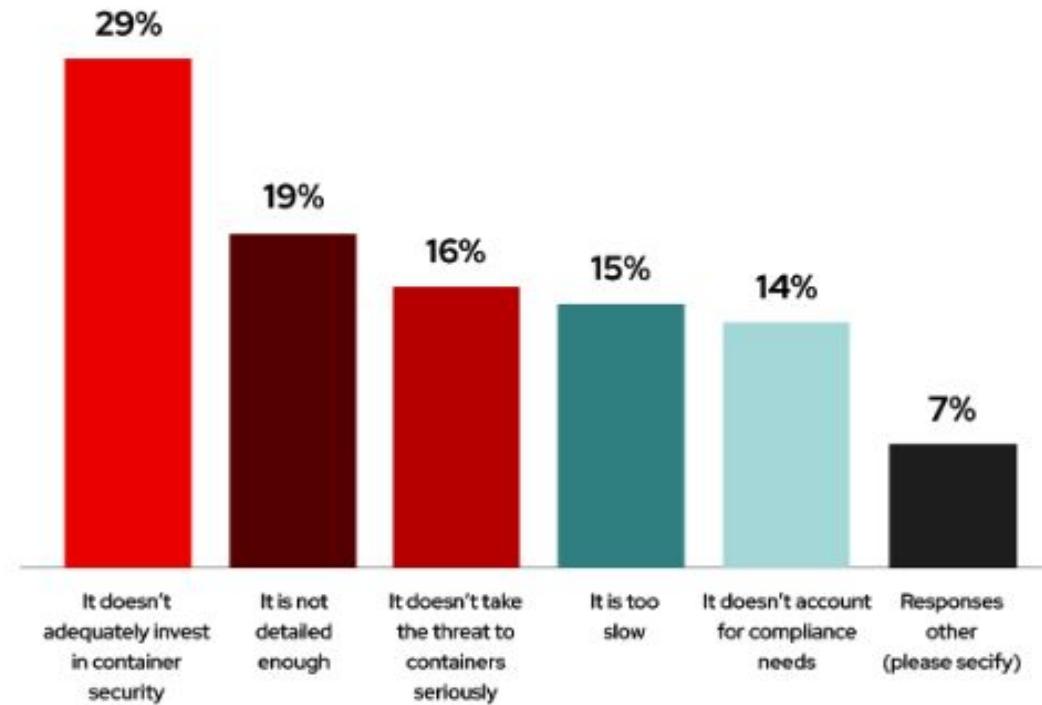
Kubernetes Security Best Practices

94% of respondents experienced at least one security incident in their Kubernetes environments in the last 12 months

Misconfiguration is the leading cause of security incidents, by a wide margin



What is your biggest concern about your company's container strategy?



Considerations for Securing Containers and Kubernetes

NIST 800-190

"Use container-specific host OSs instead of general-purpose ones to reduce attack surfaces."

CNCF Kube Security Audit

"...the underlying hosts, components, and environment of a Kubernetes cluster must be configured and managed. This management has a direct impact on the capabilities of the cluster..."

Gartner Market Guide for Cloud Workload Protection

"The best way to secure these rapidly changing and short-lived workloads is to start their protection proactively in the development phase ..."

"Replace antivirus (AV)-centric strategies with a "zero-trust execution"/default deny/application control approach to workload protection where possible."

Sources

[NIST Special Publication 800-190 Application Container Security Guide](#)

[CNCF Cloud Native Security Whitepaper](#)

[Kubernetes Security Whitepaper](#), Trail of Bits, May 31, 2019

Gartner: Market Guide for Cloud Workload Protection Platforms, ID G00356240, April 8, 2019

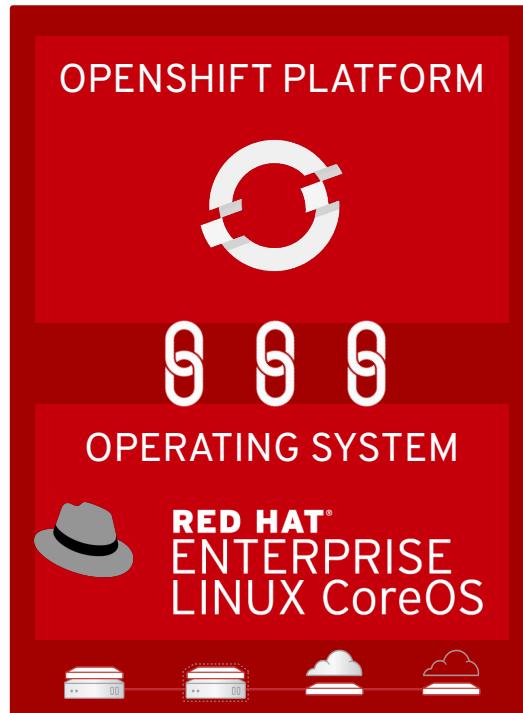


Red Hat Enterprise Linux CoreOS

Controlled Immutability / Container Optimized

Role in OpenShift Ecosystem

OPENSHIFT 4

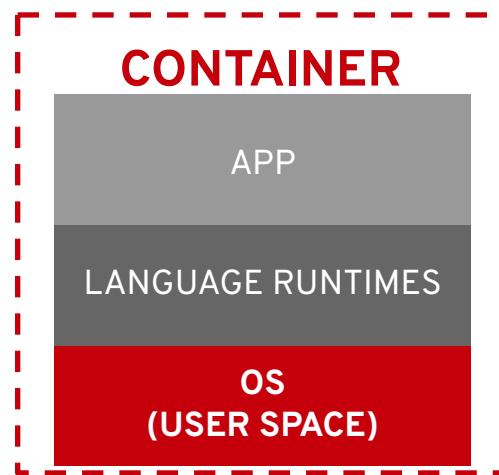


Managed by the OpenShift Cluster

- ▶ Considered a member of an OpenShift Deployment
- ▶ Configuration managed by the Machine Config Operator
 - Container runtime
 - Kubelet configuration
 - Authorized container registries
 - SSH Configuration
 - Multiple machine pools can be created
- ▶ Continuously monitoring for configuration drift
- ▶ Deploy the [File Integrity operator](#) to monitor for changes to files

The Red Hat Universal Base Image

The base image for all your needs -- enterprise architecture, security and performance



The Red Hat Universal Base Image is based on RHEL and made available at no charge by a new end user license agreement.

Development

- ▶ Minimal footprint (~90 to ~200MB)
- ▶ Programming languages (Modularity & AppStreams)
- ▶ Enables a single CI/CD chain

Production

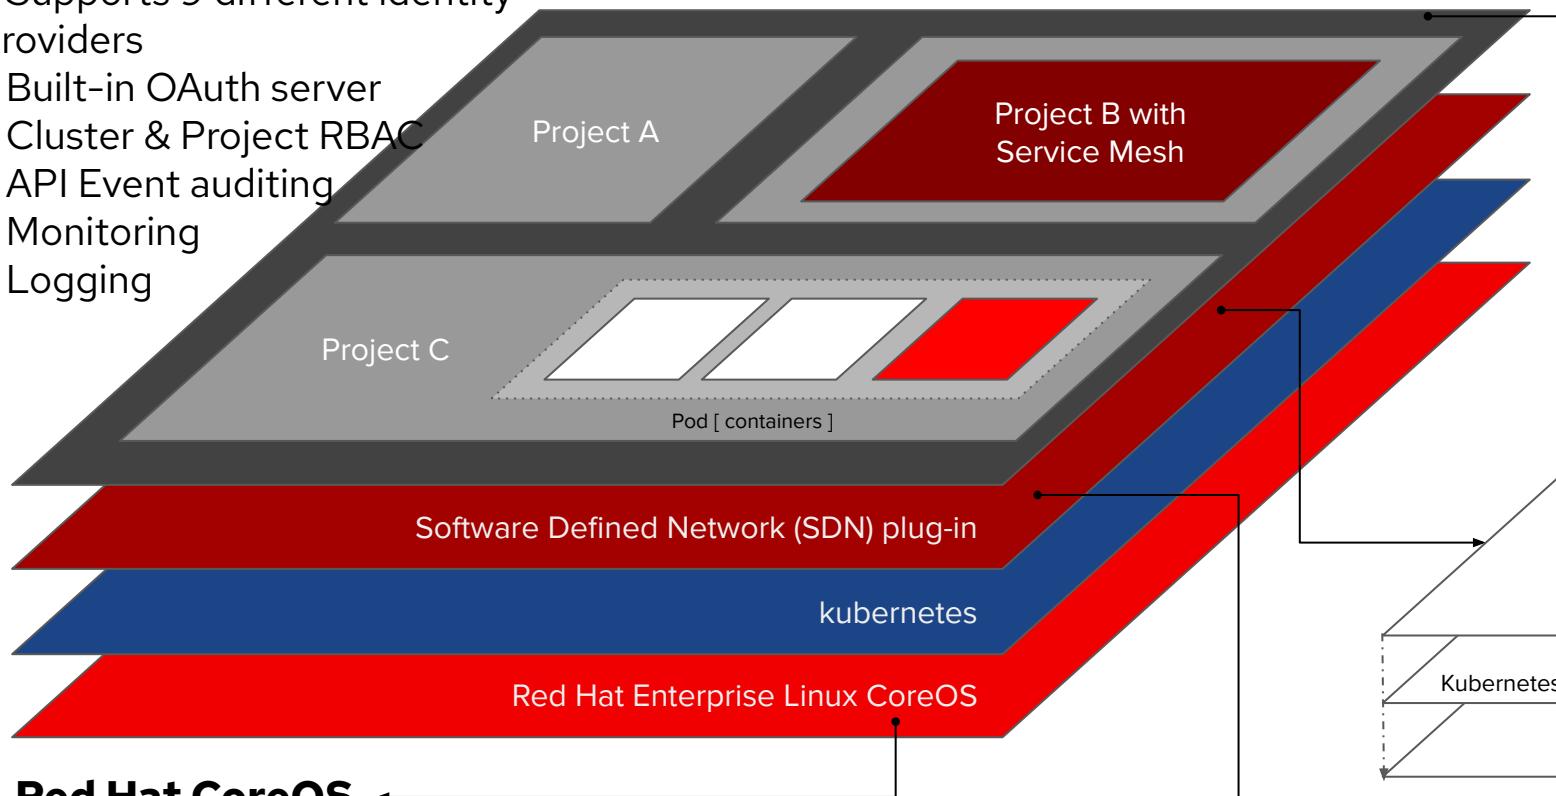
- ▶ Supported as RHEL when running on RHEL
- ▶ Same Performance, Security & Life cycle as RHEL
- ▶ Can attach RHEL support subscriptions as RHEL

Defense in Depth

CONFIDENTIAL designator

OpenShift Container Platform

- Supports 9 different identity providers
- Built-in OAuth server
- Cluster & Project RBAC
- API Event auditing
- Monitoring
- Logging

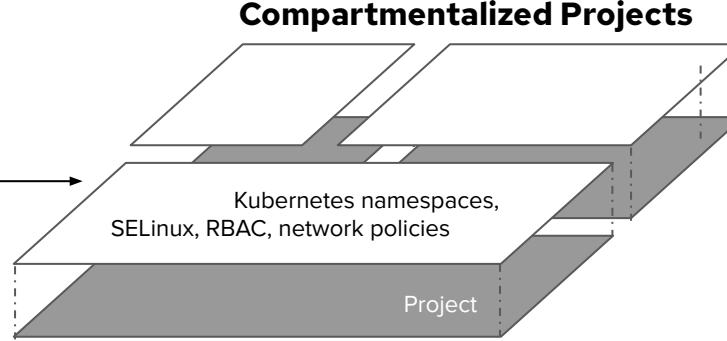


Red Hat CoreOS

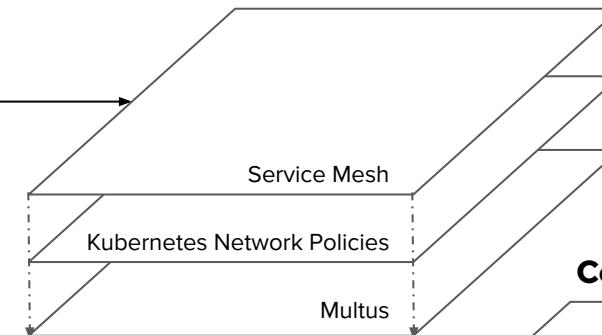
- Minimized attack surface
- Controlled Immutability
- SELinux on by default
- Kernel namespaces and Cgroups
- CRI-O container runtime, Kubelet
- Auditd for host-level audit
- FIPS enablement
- RHCOS volume encryption

Security Context Constraints (Kubernetes Admission Controller)

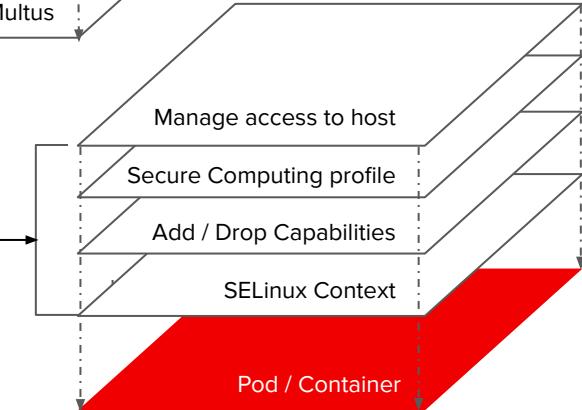
Red Hat Universal Base Image
containers can leverage RHEL
FIPS capabilities



Network Security



Container Security



Comprehensive, Layered Approach to Security

Control

Fleet

Application Lifecycle and Locality

Workload

Vulnerability analysis

App config analysis

APIs for CI/CD integrations

Trusted content

Container registry

Build management

CI/CD pipeline

BUILD



Protect

Fleet Management

Policy admission controller

Compliance assessments

Risk profiling

Kubernetes platform lifecycle

Identity and access management

Platform data

Deployment policies

DEPLOY



Detect & Respond

Fleet Observability & Alerts

Runtime behavioral analysis

Auto-suggest network policies

Threat detection / incident response

Container isolation

Network isolation

Application access and data

Observability

RUN

DevSecOps

Build Security into the App

- Trusted Content
- Enterprise Registry
- Control and Automate builds
- DevSecOps - Must be in the pipeline

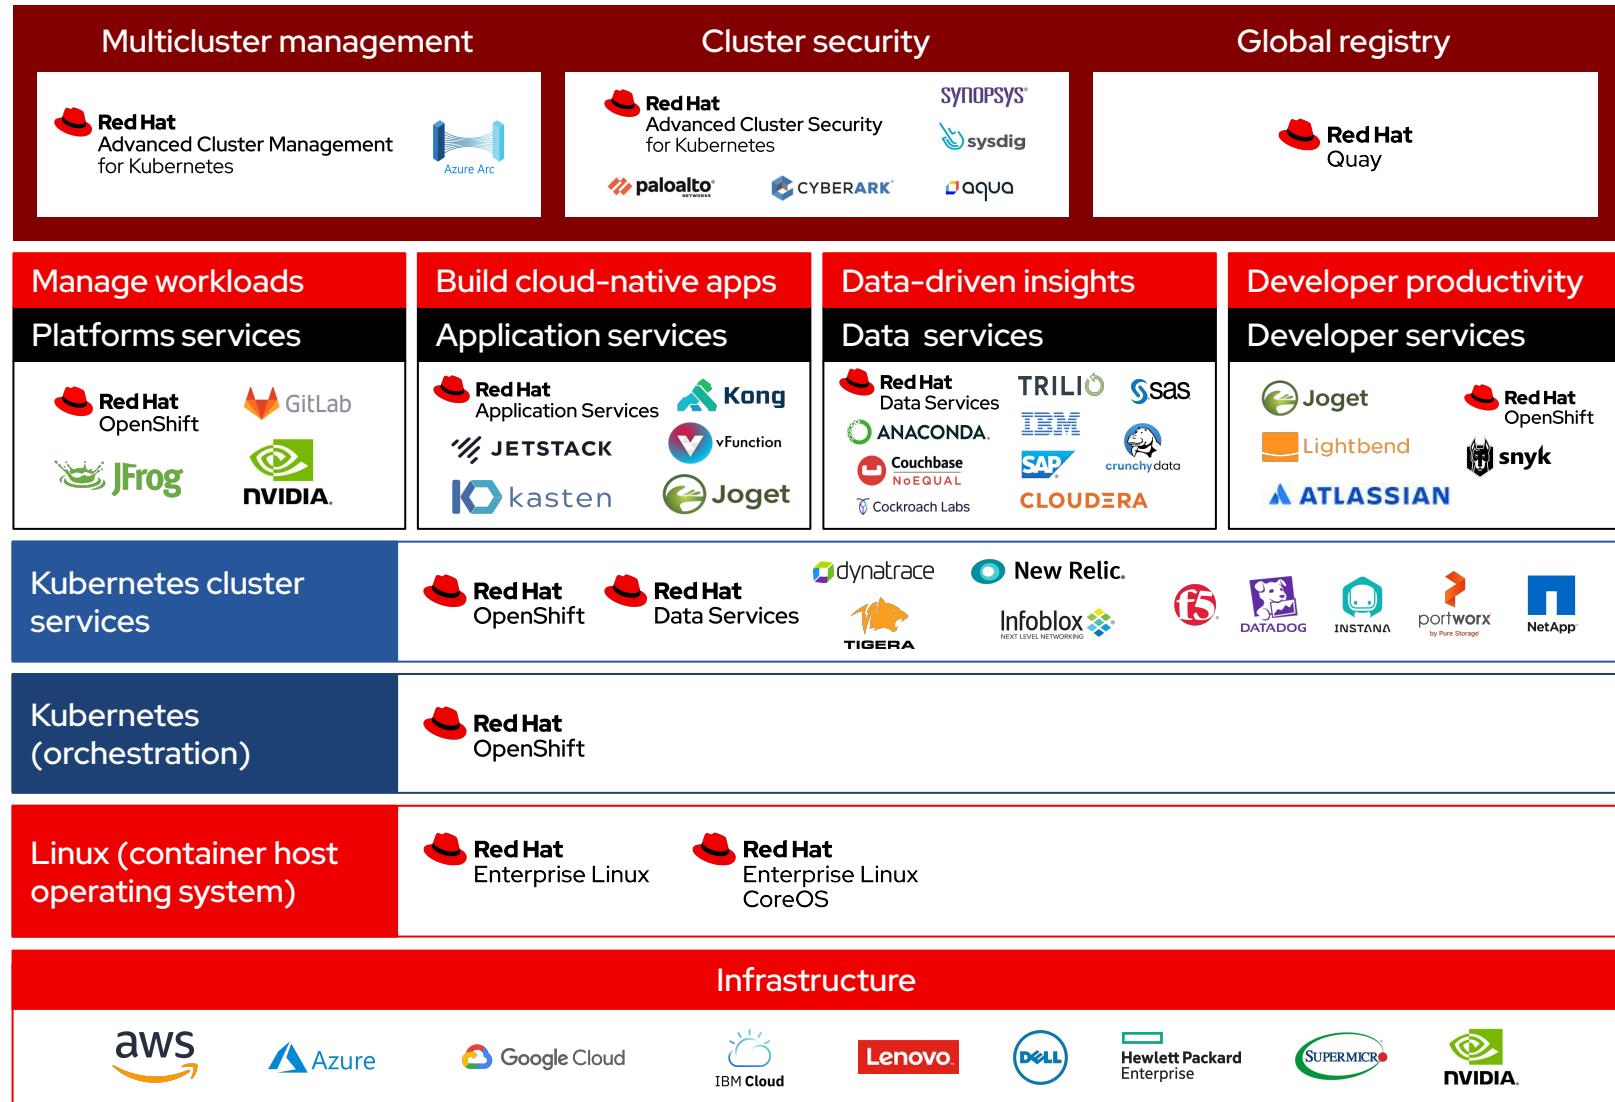
Manage deployment config, security, compliance

- Platform configuration and life-cycle management
- Identity and access management
- Security for platform data and attached storage
- Deployment Policies

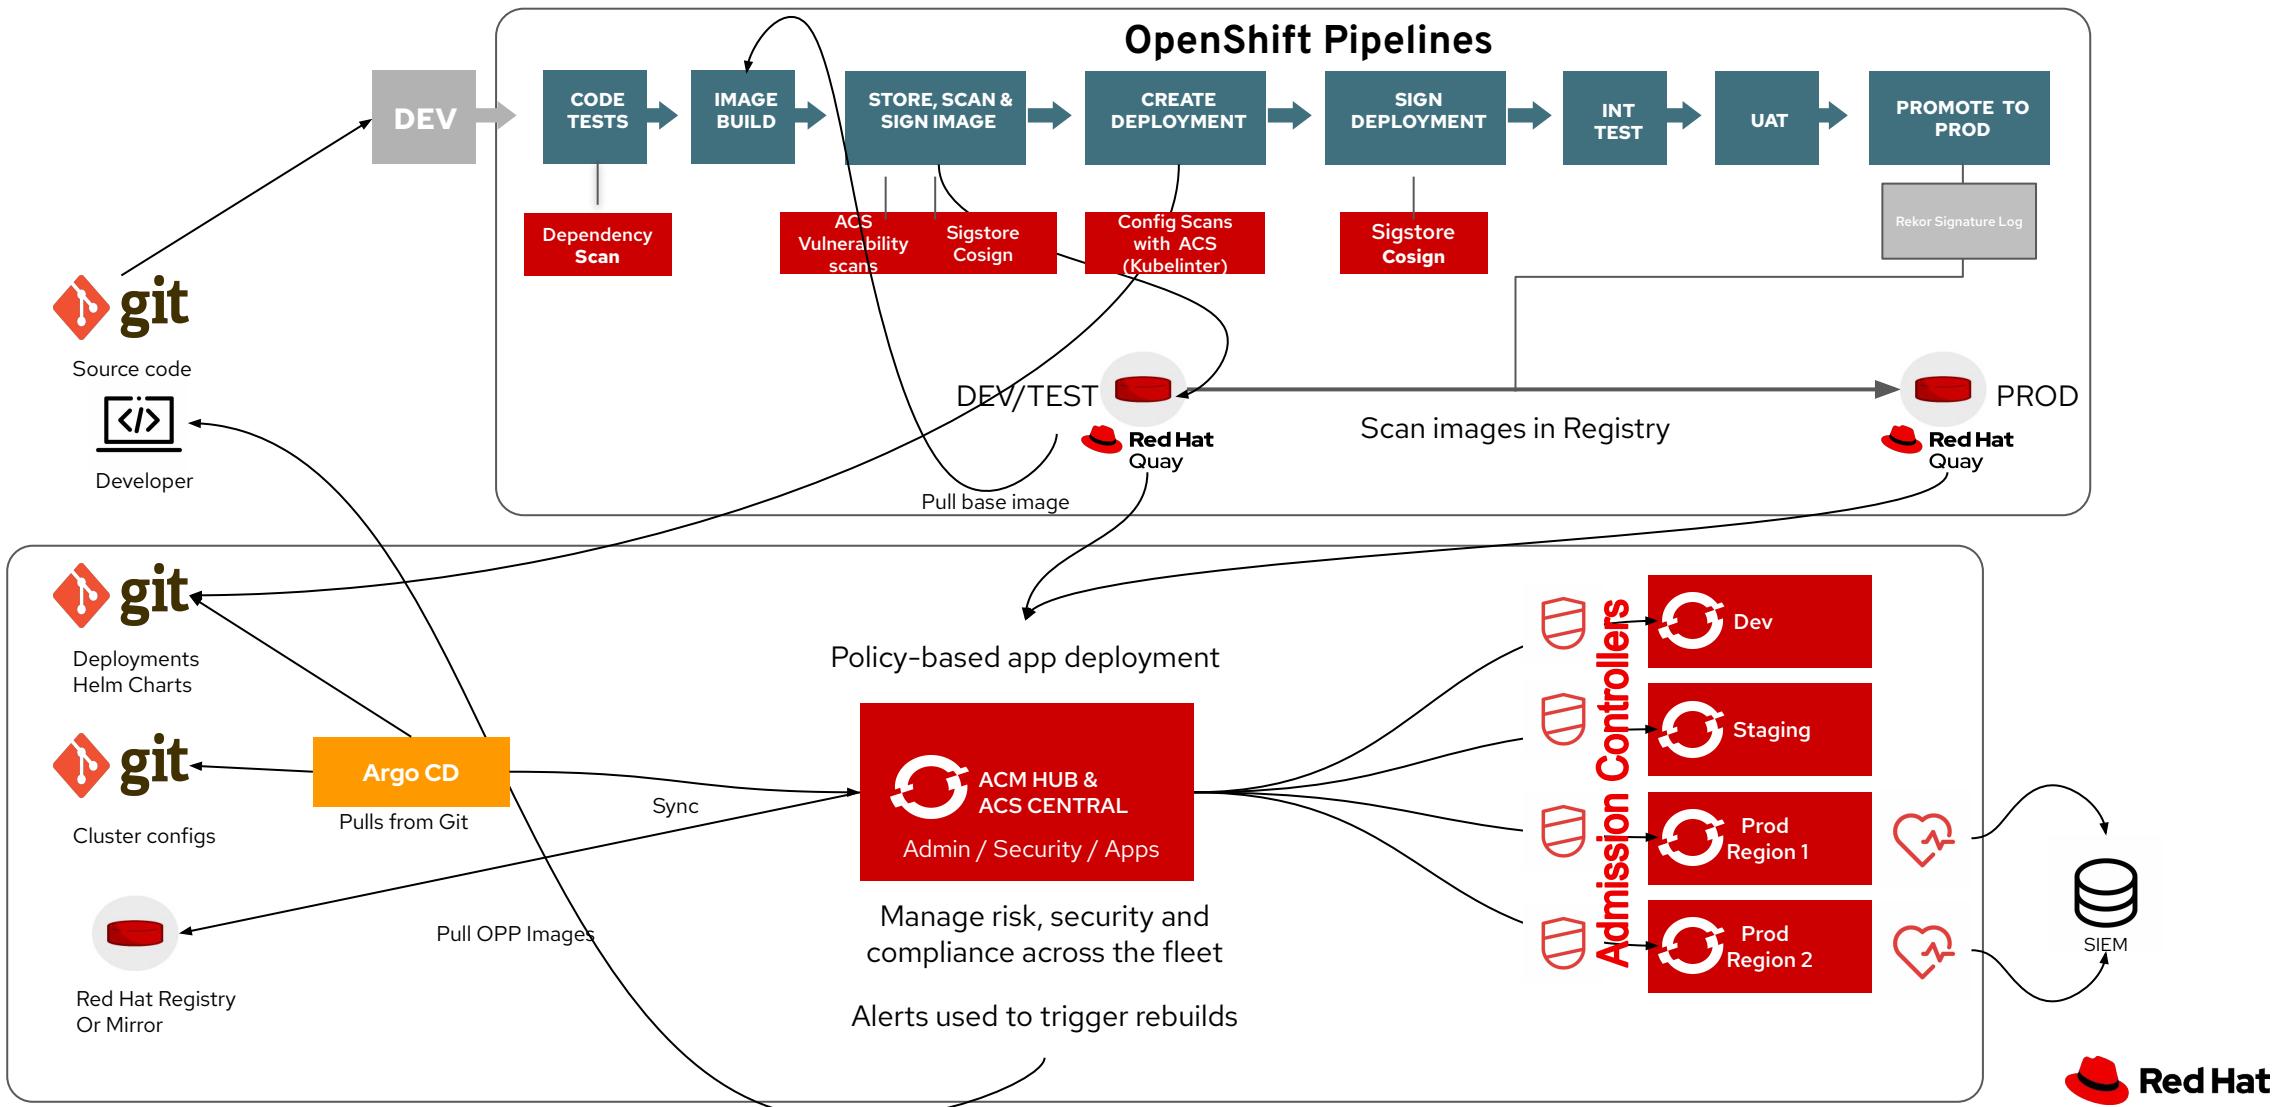
Runtime Protection

- Container Isolation
- Application and Network Isolation
- Secure Access to Apps
- Observability

Extend security with best in class tooling



Policy Based Multi-cluster DevSecOps



Resources

Container security maturity model -

<https://www.redhat.com/en/resources/container-security-maturity-model-whitepaper>

Managing NIST 800-53 Controls in a Multicloud OpenShift Environment -

<https://cloud.redhat.com/blog/managing-nist-800-53-controls-in-a-multicloud-openshift-environment-part-1>

Container and Kubernetes security: An evaluation guide -

<https://www.redhat.com/en/resources/container-kubernetes-security-evaluation-detail>

State of Kubernetes security report -

<https://www.redhat.com/en/resources/state-kubernetes-security-report>

Resources

How OpenShift enables container security -

<https://www.redhat.com/en/technologies/cloud-computing/openshift/security>

A layered approach to container and Kubernetes security -

<https://www.redhat.com/en/resources/layered-approach-container-kubernetes-security-whitepaper>

OpenShift Security Guide (300 page in depth guide!) -

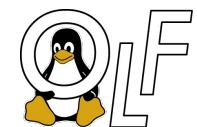
<https://www.redhat.com/rhdc/managed-files/cl-openshift-security-guide-ebook-us287757-202103.pdf>

Exercise 4: Application Storage Basics

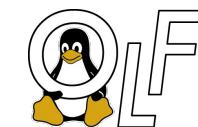
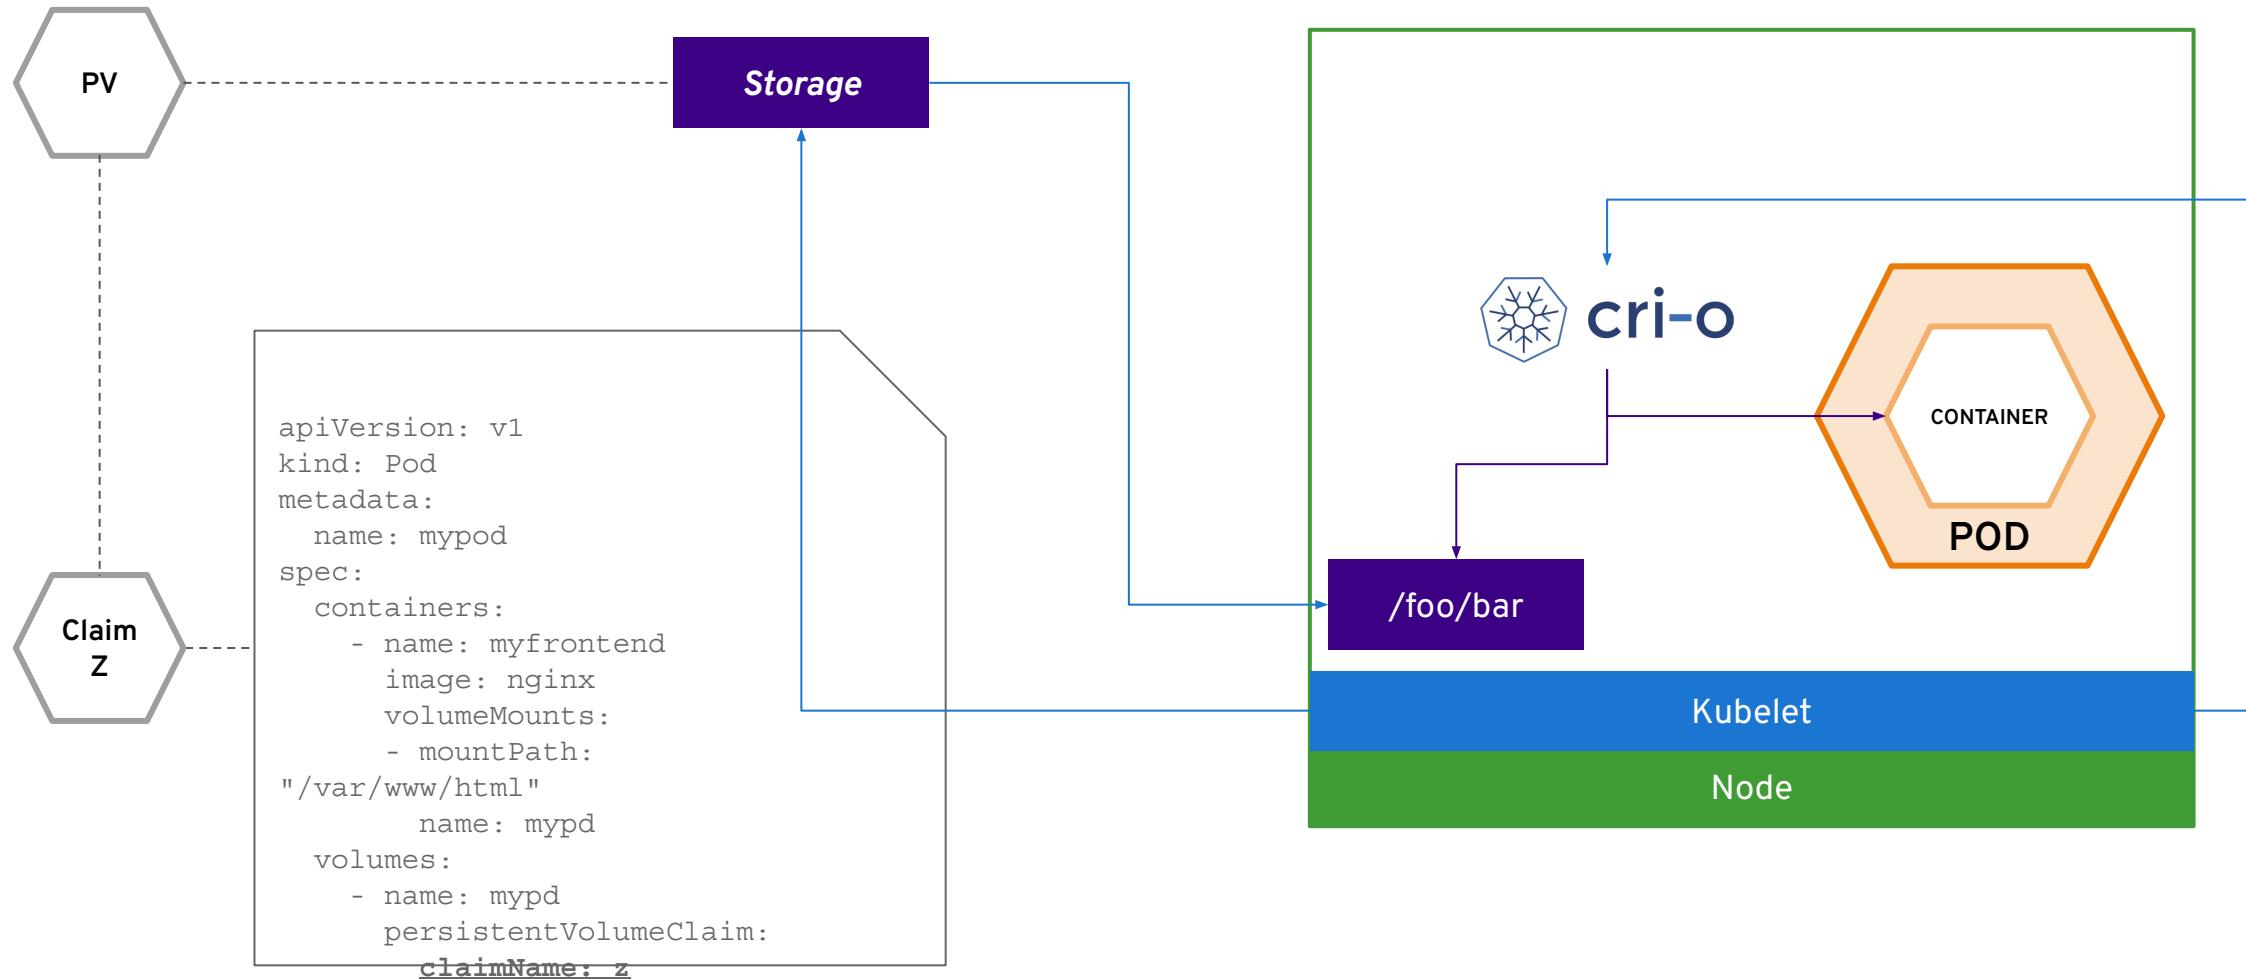
Persistent Volumes and
Persistent Volume Claims

A broad spectrum of static and dynamic storage endpoints

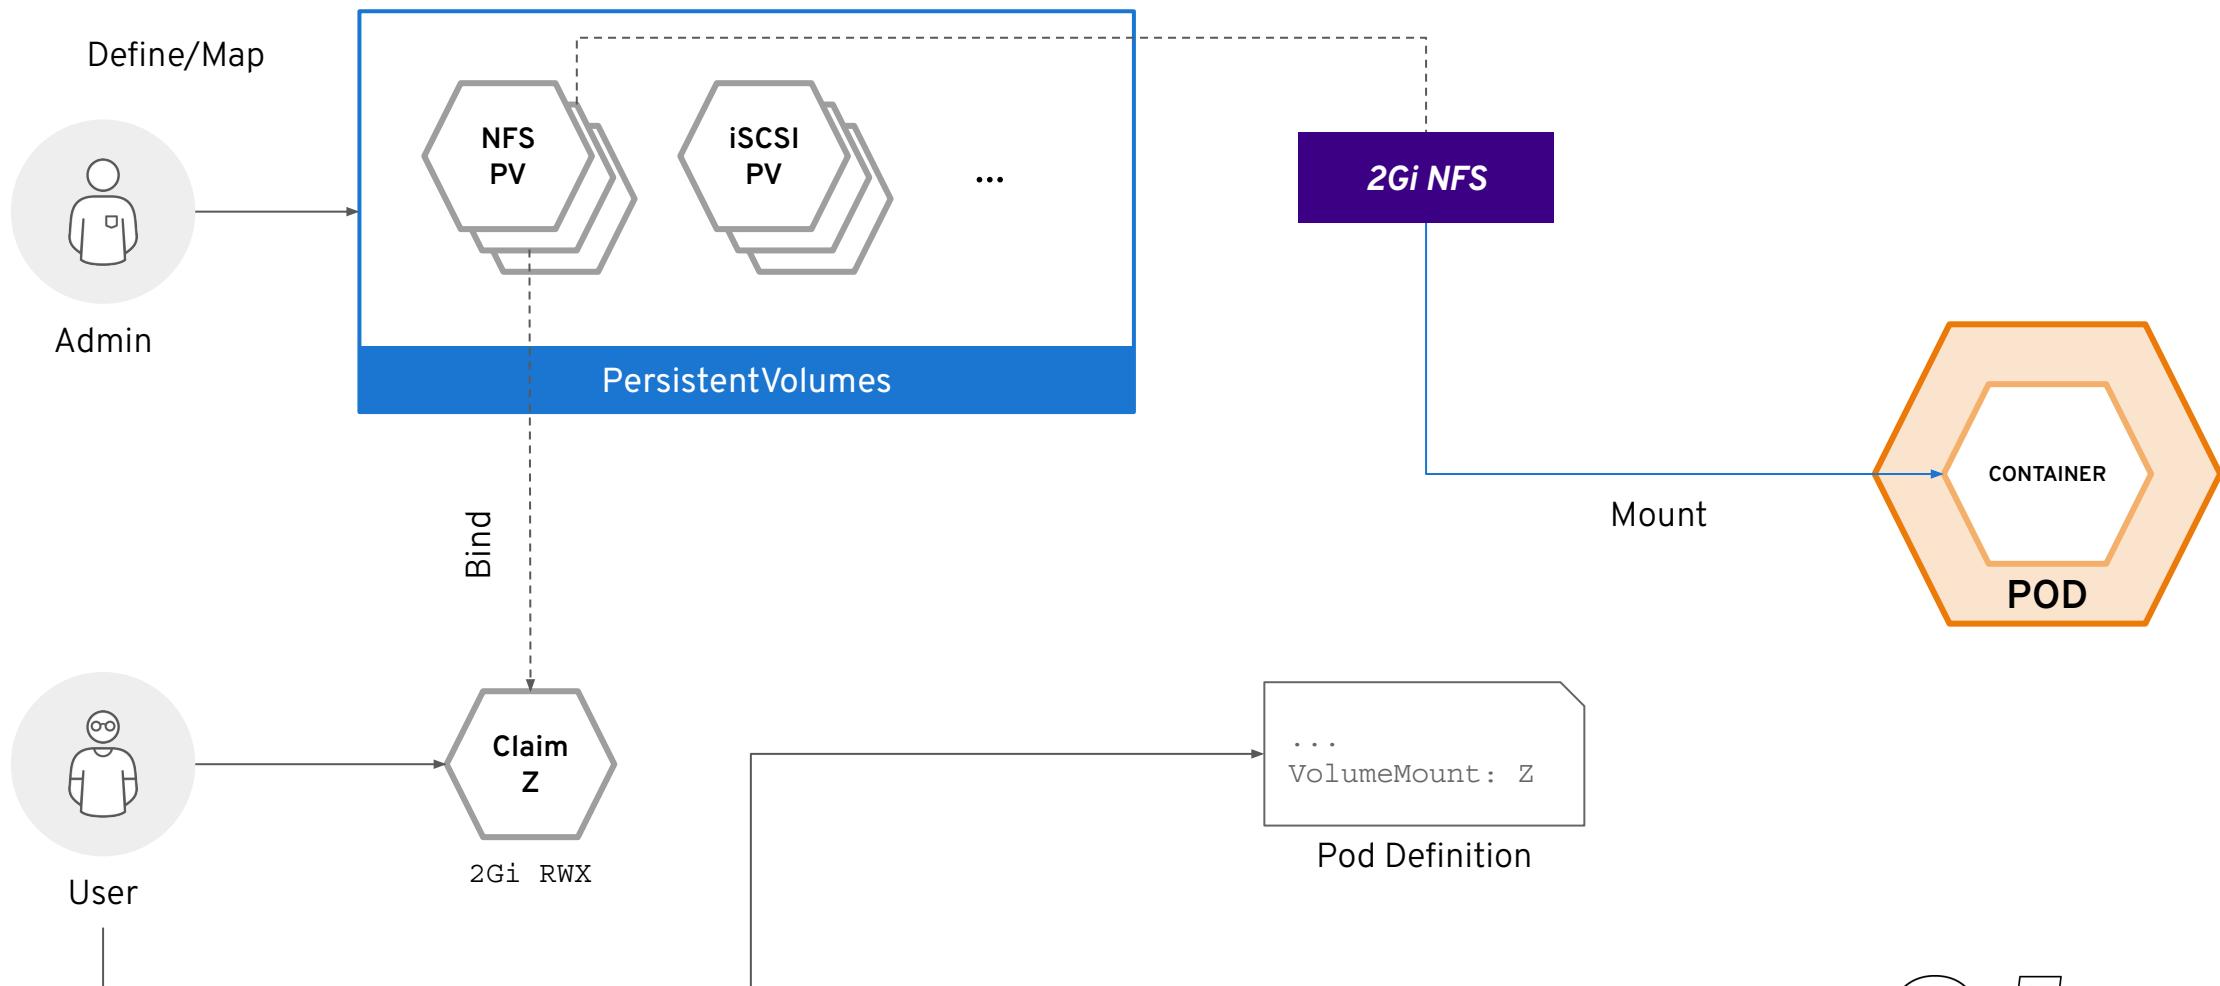
NFS	OpenStack Cinder	iSCSI	Azure Disk	AWS EBS	FlexVolume
GlusterFS	Ceph RBD	Fiber Channel	Azure File	GCE Persistent Disk	VMWare vSphere VMDK
NetApp Trident*		Container Storage Interface (CSI)**			



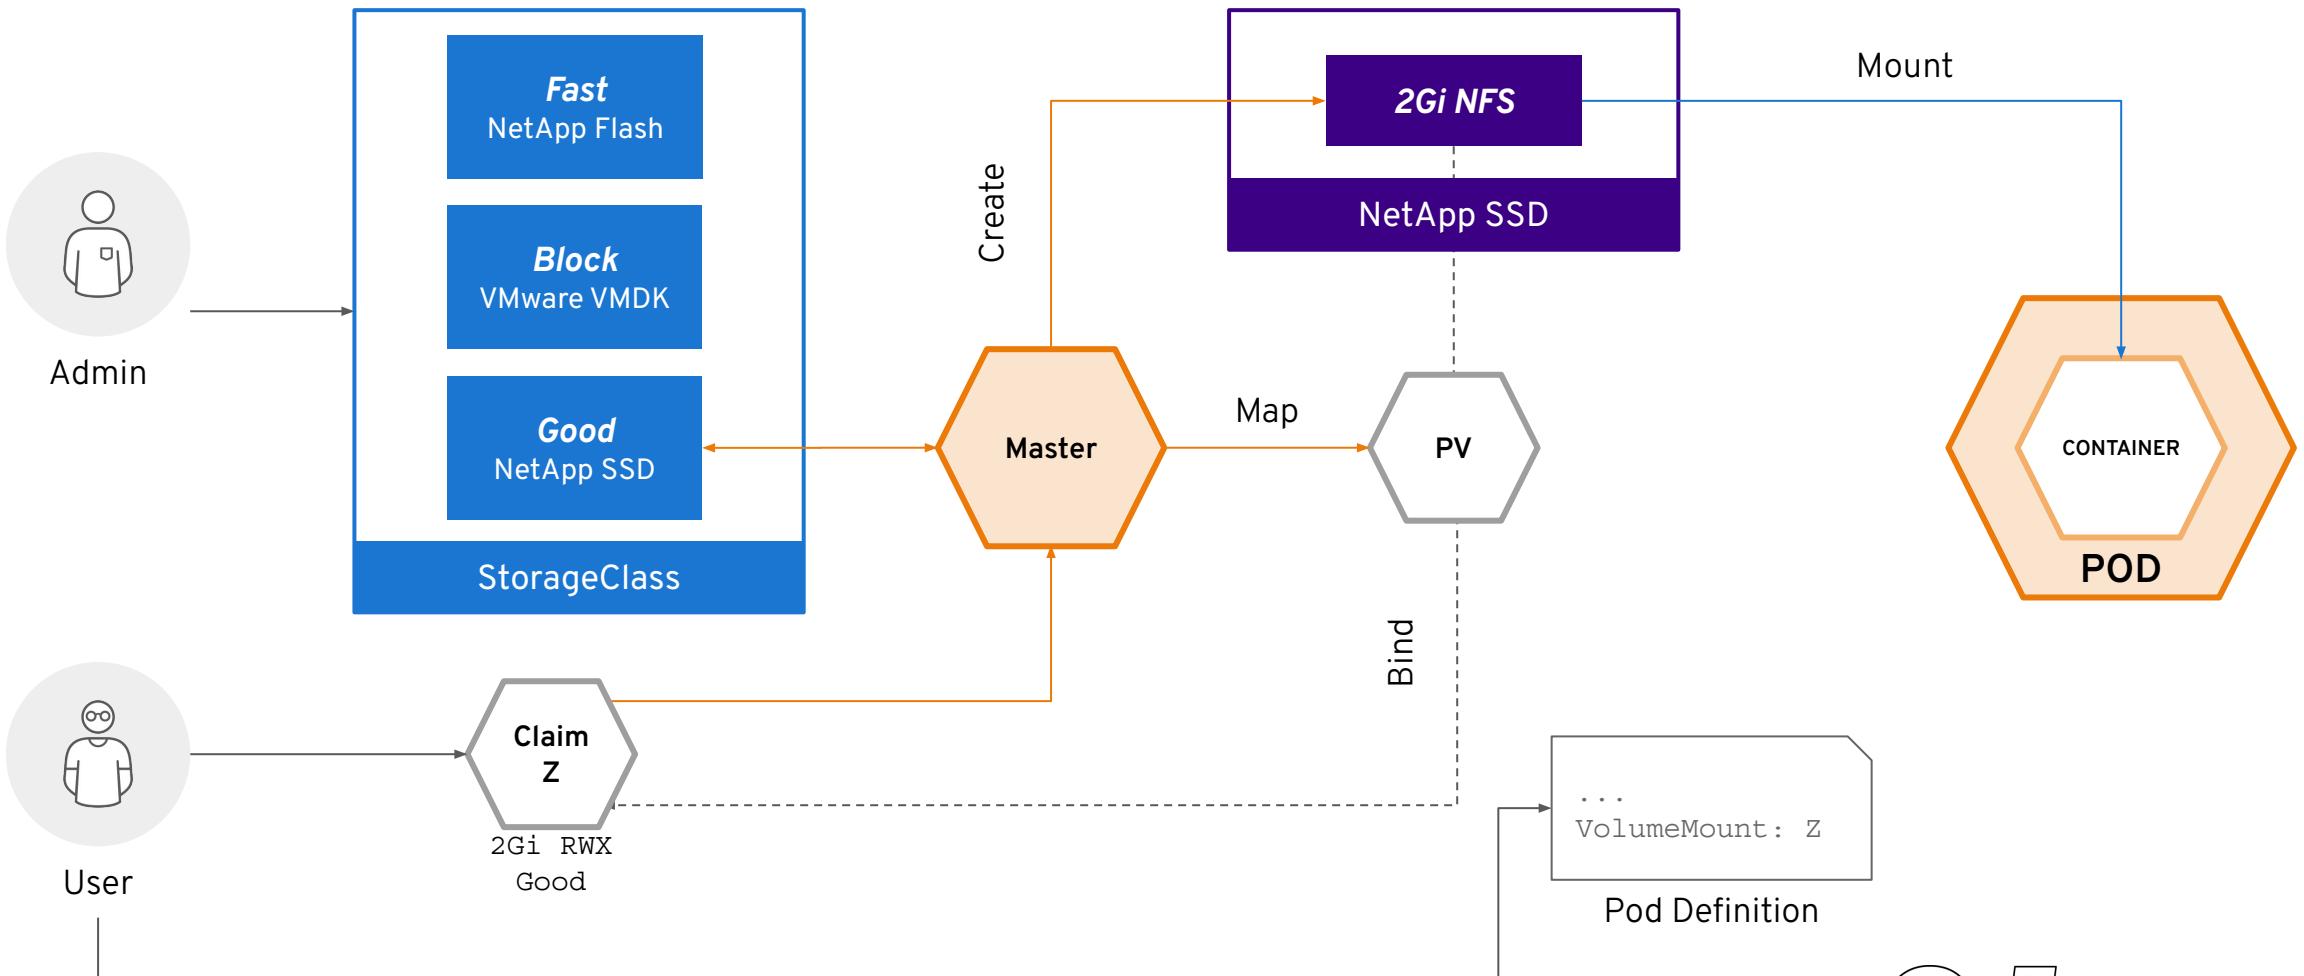
PV Consumption



Static Storage Provisioning



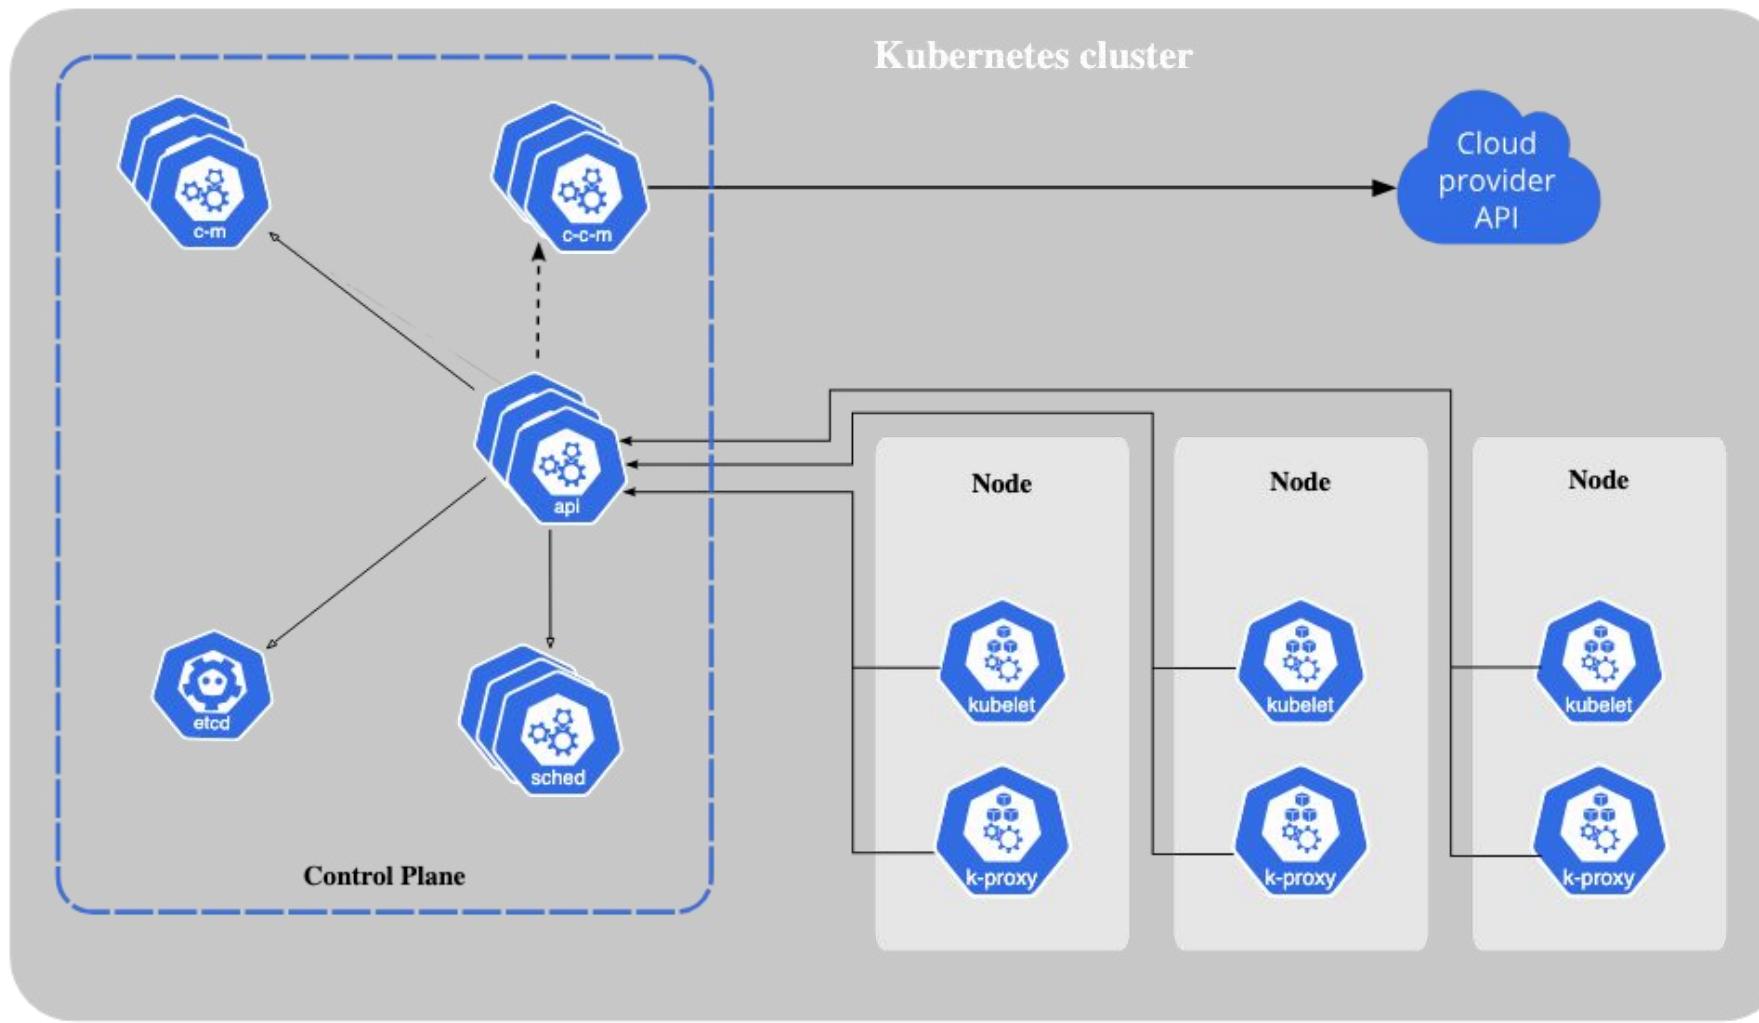
Dynamic Storage Provisioning



Exercise 5: Machinesets, Machines and Nodes

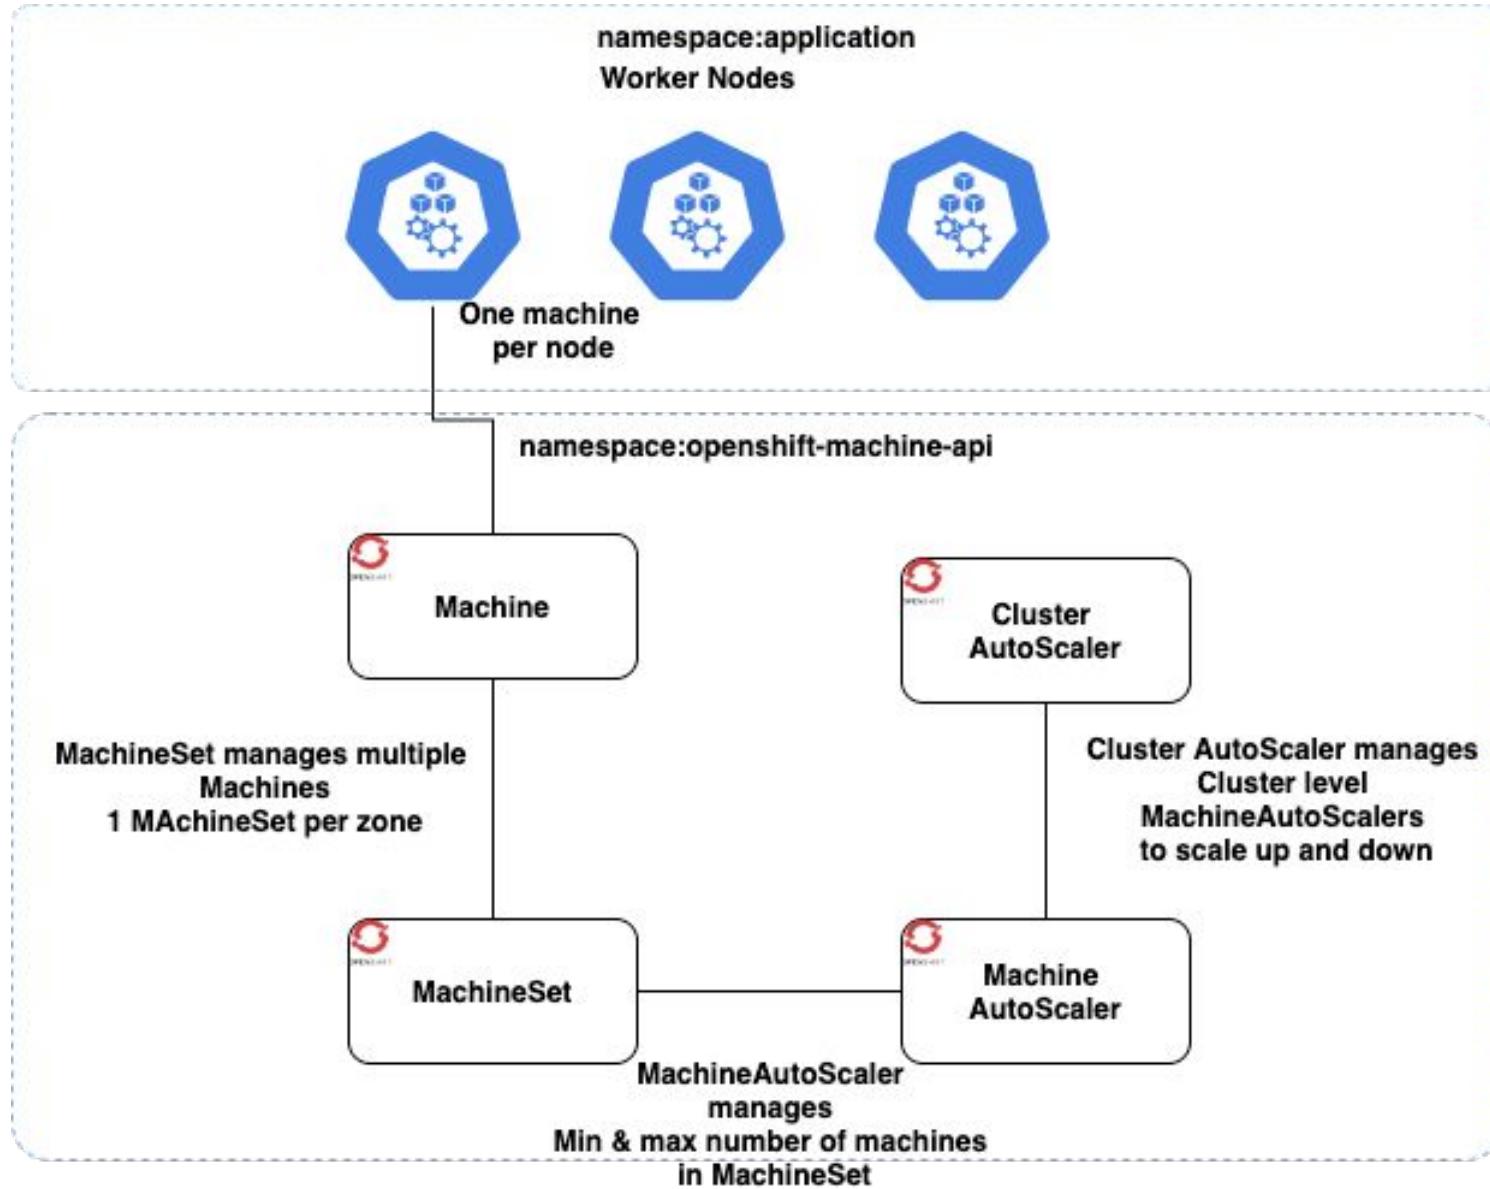
Expanding the cluster
capacity quickly and easily

Cloud Controller Manager



Cluster API

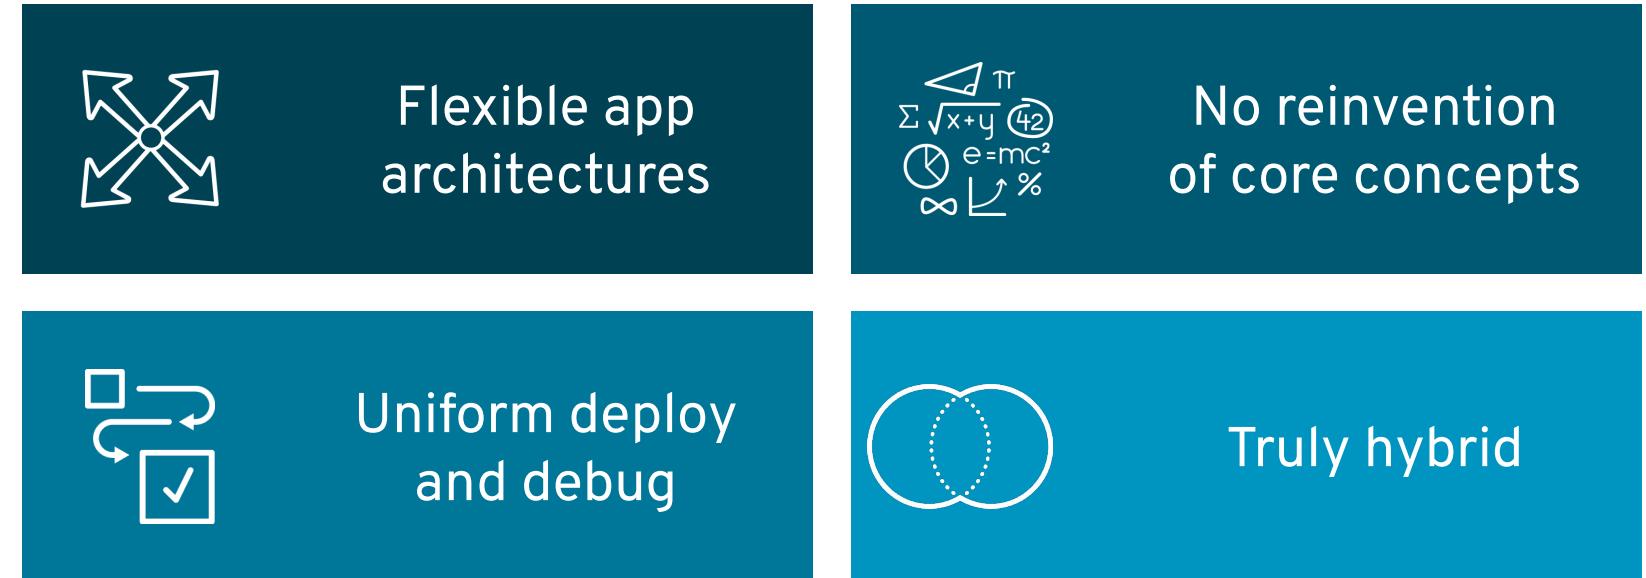
CONFIDENTIAL designator



Exercise 6: Infrastructure Nodes and Operators

Common Day2
configuration tasks

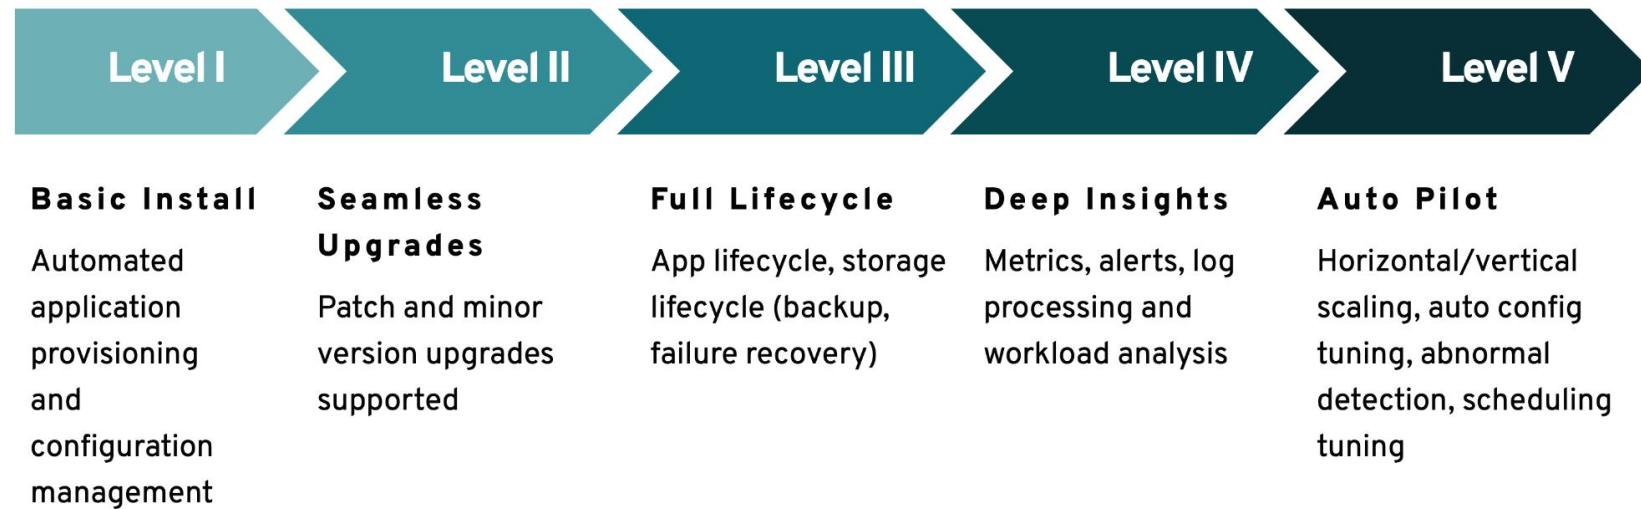
Kubernetes-native day 2 management



Operators codify operational knowledge and workflows to automate life-cycle management of containerized applications with Kubernetes

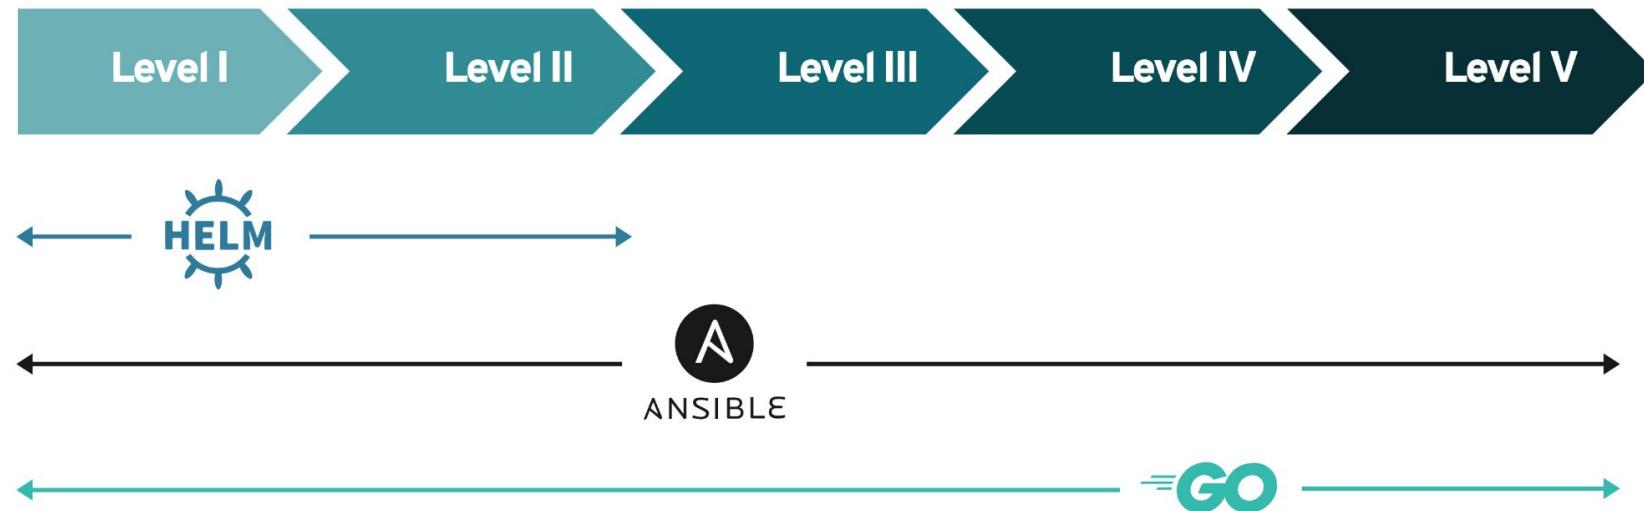
OPERATOR CAPABILITY LEVELS

Operators come in different maturity levels in regards to their lifecycle management capabilities for the application or workload they deliver. The capability models aims to provide guidance in terminology to express what features users can expect from an Operator.



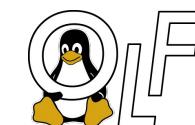
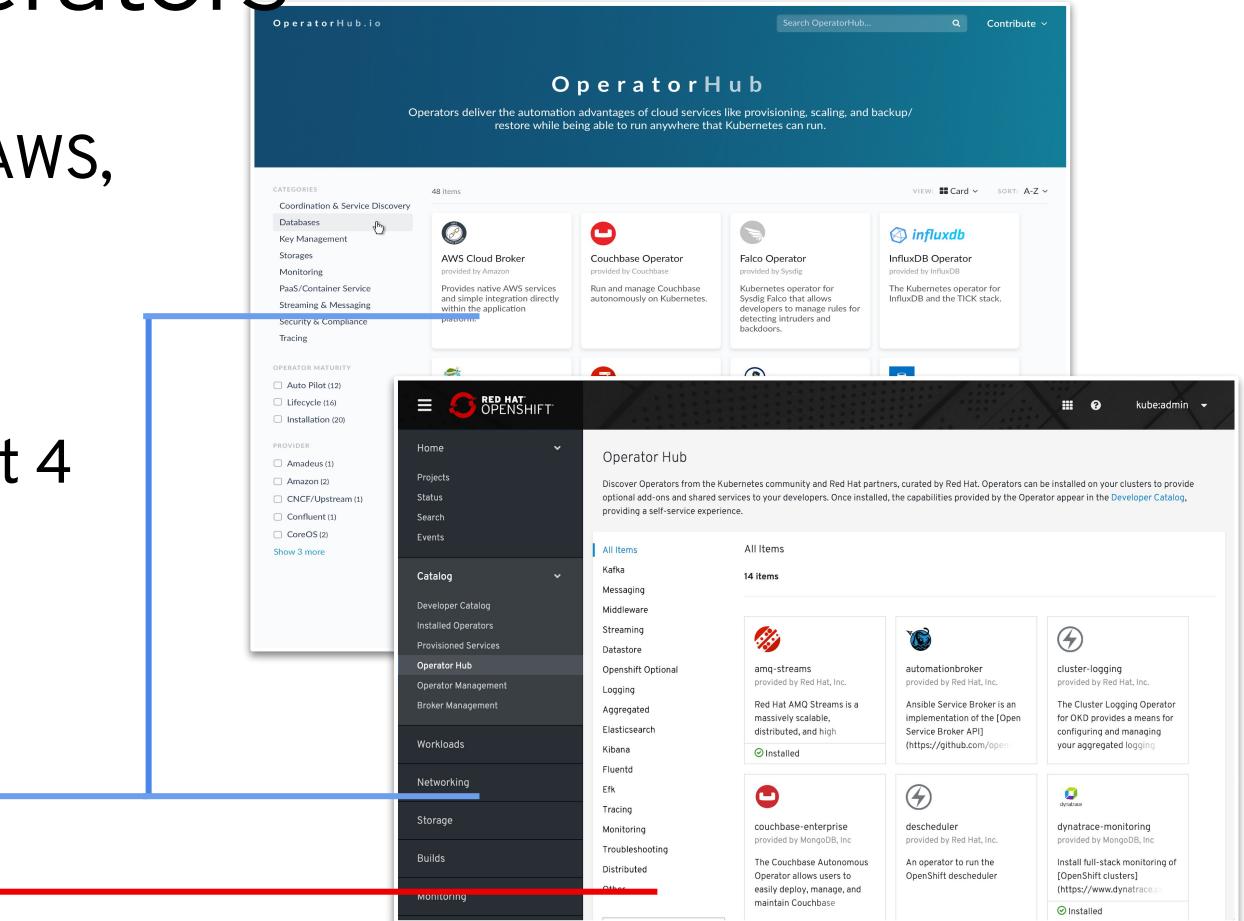
OPERATOR CAPABILITY LEVELS

Operators come in different maturity levels in regards to their lifecycle management capabilities for the application or workload they deliver. The capability models aims to provide guidance in terminology to express what features users can expect from an Operator.

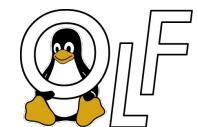


Operator Hub and certified Operators

- OperatorHub.io launched by Red Hat, AWS, Microsoft and Google
- OpenShift Operator Certification
- OperatorHub integrated into OpenShift 4



Operator Hub and certified Operators





Crunchy Postgres for Kubernetes

Production Postgres Made Easy

PGO is developed with many years of production experience in automating Postgres management on Kubernetes, providing a seamless cloud native Postgres solution to keep your data always available.

- **PostgreSQL Cluster Provisioning:** [Create, Scale, & Delete PostgreSQL clusters with ease](#), while fully customizing your Pods and PostgreSQL configuration!
- **High-Availability:** Safe, automated failover backed by a [distributed consensus based high-availability solution](#). Uses [Pod Anti-Affinity](#) to help resiliency; you can configure how aggressive this can be! Failed primaries automatically heal, allowing for faster recovery time. You can even create regularly scheduled backups as well and set your backup retention policy
- **Disaster Recovery:** [Backups](#) and [restores](#) leverage the open source [pgBackRest](#) utility and [includes support for full, incremental, and differential backups as well as efficient delta restores](#). Set how long you want your backups retained for. Works great with very large databases!
- **Monitoring:** [Track the health of your PostgreSQL clusters](#) using the open source [pgMonitor](#) library.
- **Clone:** [Create new clusters from your existing clusters or backups](#) with efficient data cloning.
- **TLS:** All connections are over [TLS](#). You can also [bring your own TLS infrastructure](#) if you do not want to use the provided defaults.
- **Connection Pooling:** Advanced [connection pooling](#) support using [pgBouncer](#).
- **Affinity and Tolerations:** Have your PostgreSQL clusters deployed to [Kubernetes Nodes](#) of your preference. Set your [pod anti-affinity](#), node affinity, Pod tolerations and more rules to customize your deployment topology!
- **Full Customizability:** Crunchy PostgreSQL for Kubernetes makes it easy to get your own PostgreSQL-as-a-Service up and running and fully customize your deployments, including:
 - Choose the resources for your Postgres cluster: [container resources and storage size. Resize at any time](#) with minimal disruption.
 - Use your own container image repository, including support [imagePullSecrets](#) and private repositories
 - [Customize your PostgreSQL configuration](#)

and much more!



Exercise 8: OpenShift Logging

A solution to visualize and corroborate log messages across pods

Vector & Loki – Logging Subsystem

The logging subsystem for Red Hat OpenShift collects container logs and node logs. These are categorized into types:

- ▶ **application** - Container logs generated by non-infrastructure containers.
- ▶ **infrastructure** - Container logs from namespaces `kube-*` and `openshift-*`, and node logs from journald.
- ▶ **audit** - Logs from auditd, kube-apiserver, openshift-apiserver, and ovn if enabled.

Vector & Loki – Logging Subsystem

Components

- **Loki:** a search and analytics engine to store logs
- **Vector:** gathers logs and sends to Elasticsearch.
- **Kibana:** A web UI for Elasticsearch.

Access control

- Cluster administrators can view all logs
- Users can only view logs for their projects

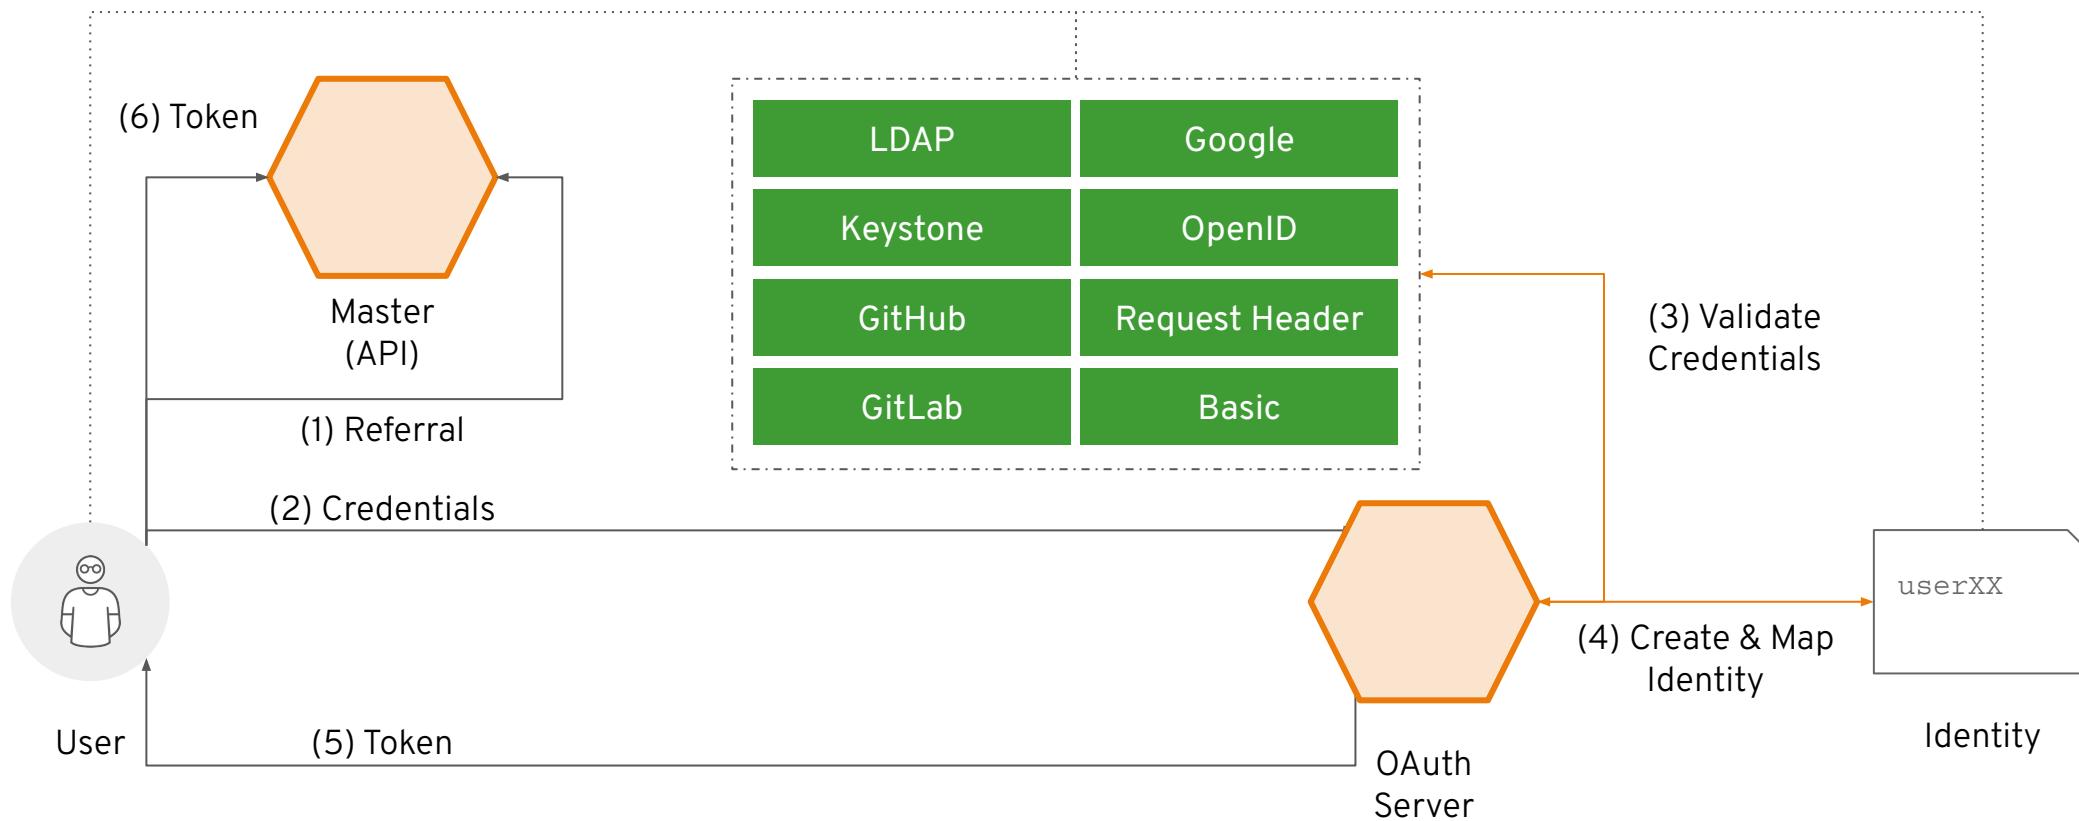
Ability to forward logs elsewhere

- External syslog, elasticsearch, Splunk, Cloudwatch, Kafka

Exercise 9: Authentication Providers, Users, and Groups

Configuring RBAC
alongside LDAP

Identity and Access Management



Fine-Grained RBAC

- Project scope & cluster scope available
- Matches request attributes (verb,object,etc)
- If no roles match, request is denied (deny by default)
- Operator- and user-level roles are defined by default
- Custom roles are supported

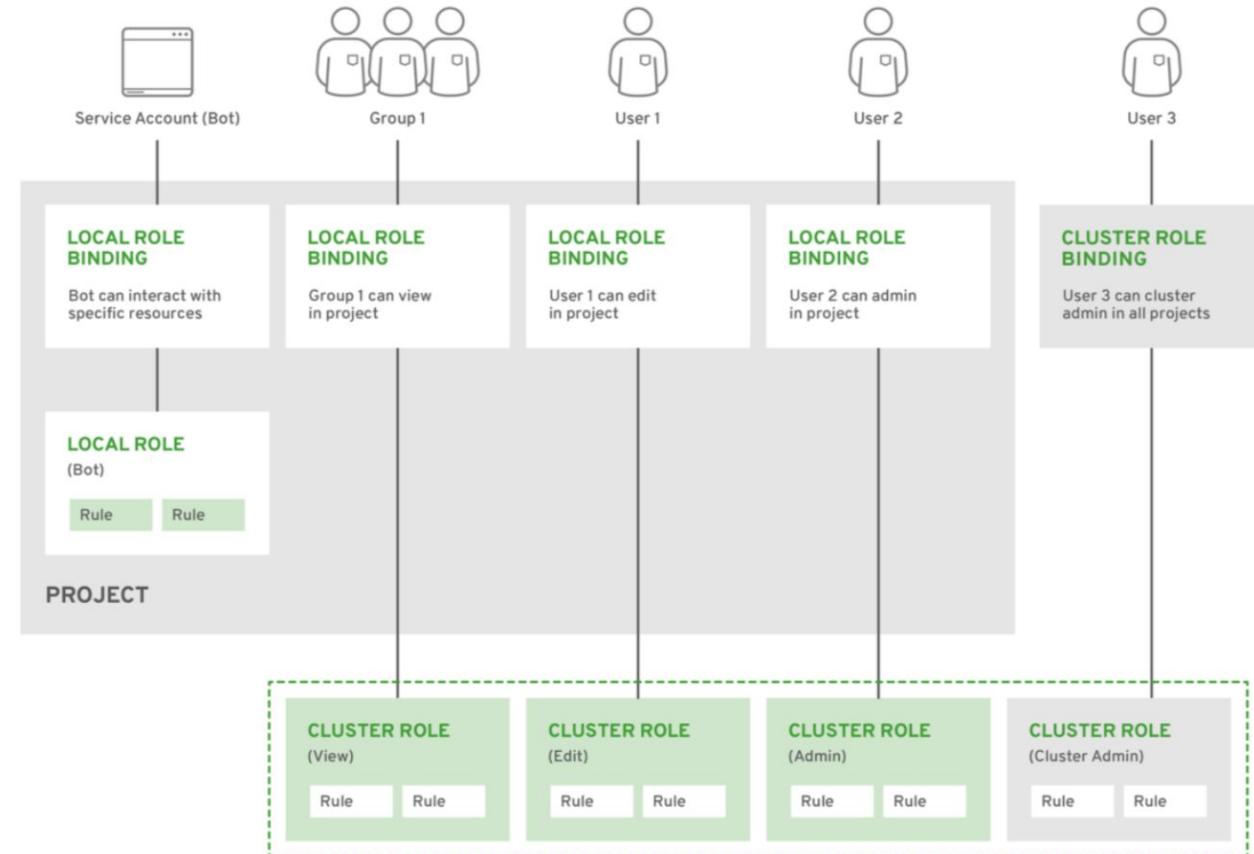


Figure 12 - Authorization Relationships

Exercise 10: OpenShift Cluster Monitoring

A no-touch built-in
monitoring solution
implemented with Red Hat
SRE best practices

OpenShift Cluster Monitoring



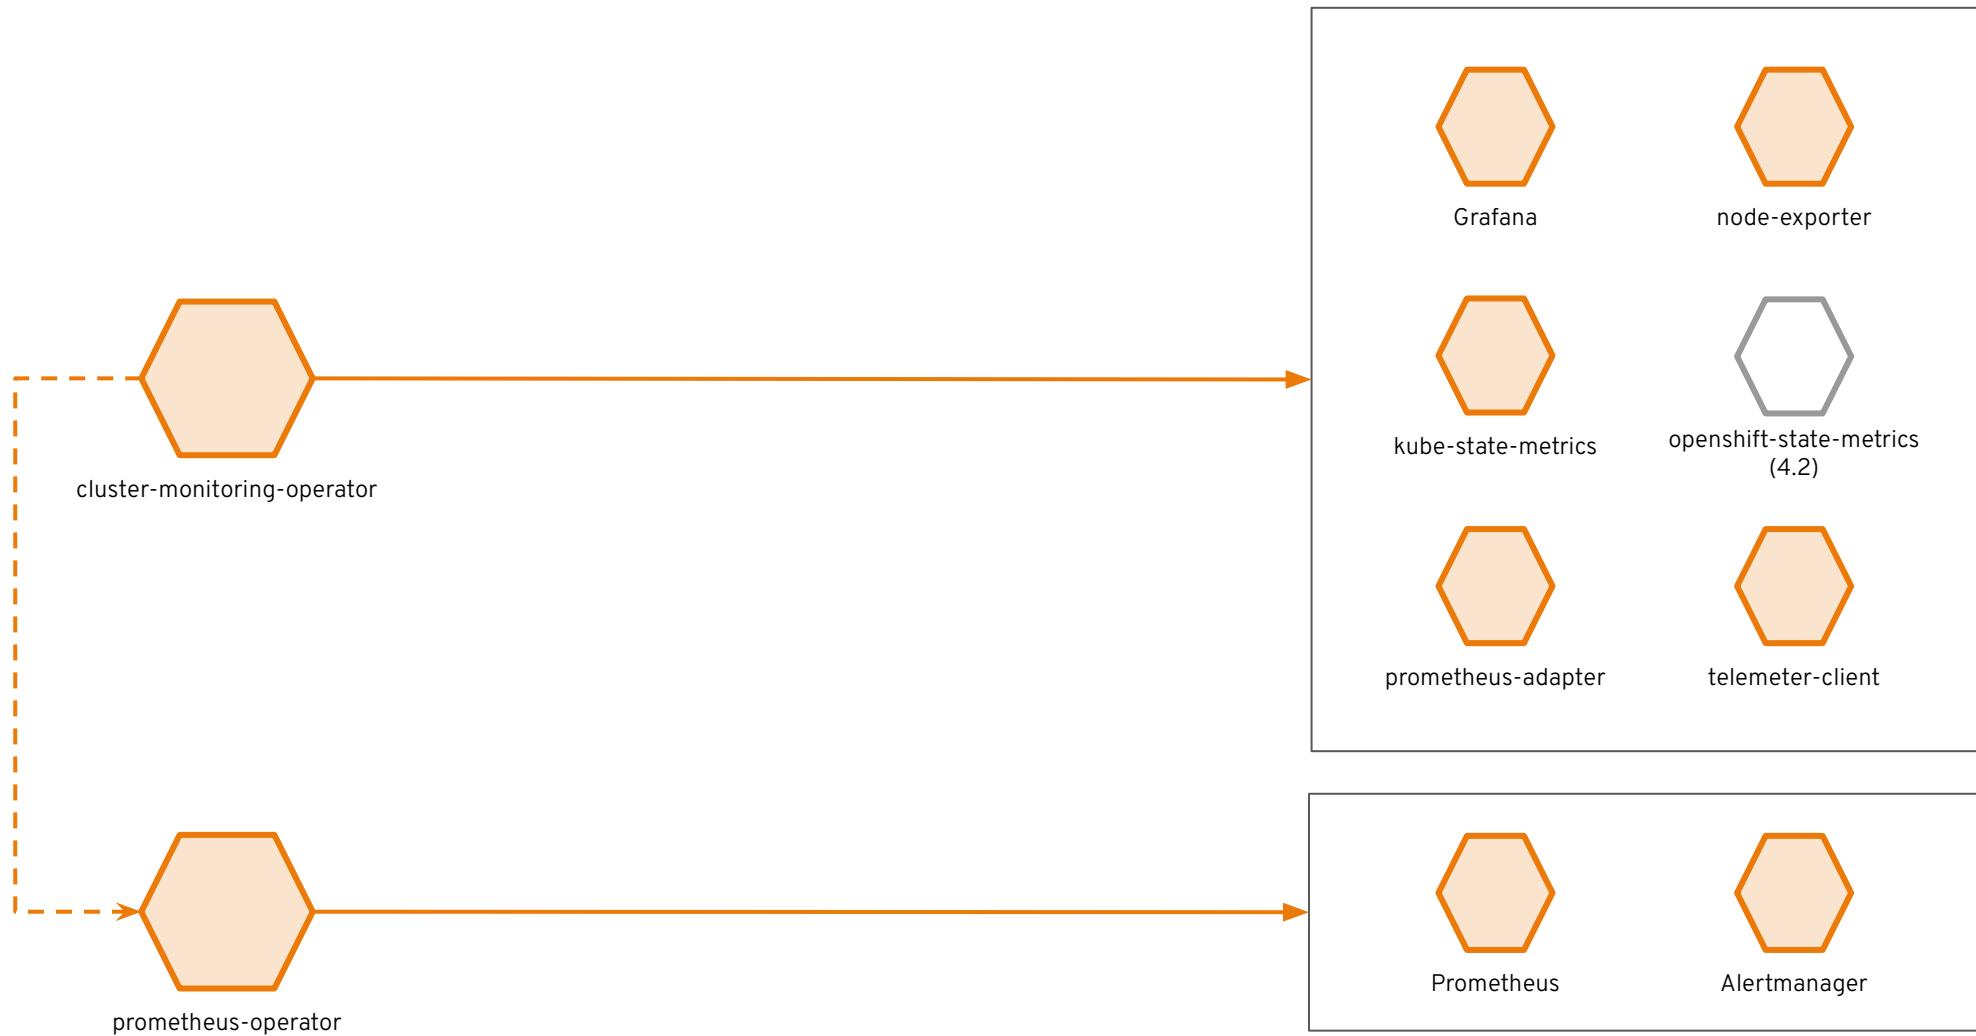
Metrics collection and storage via Prometheus, an open-source monitoring system time series database.

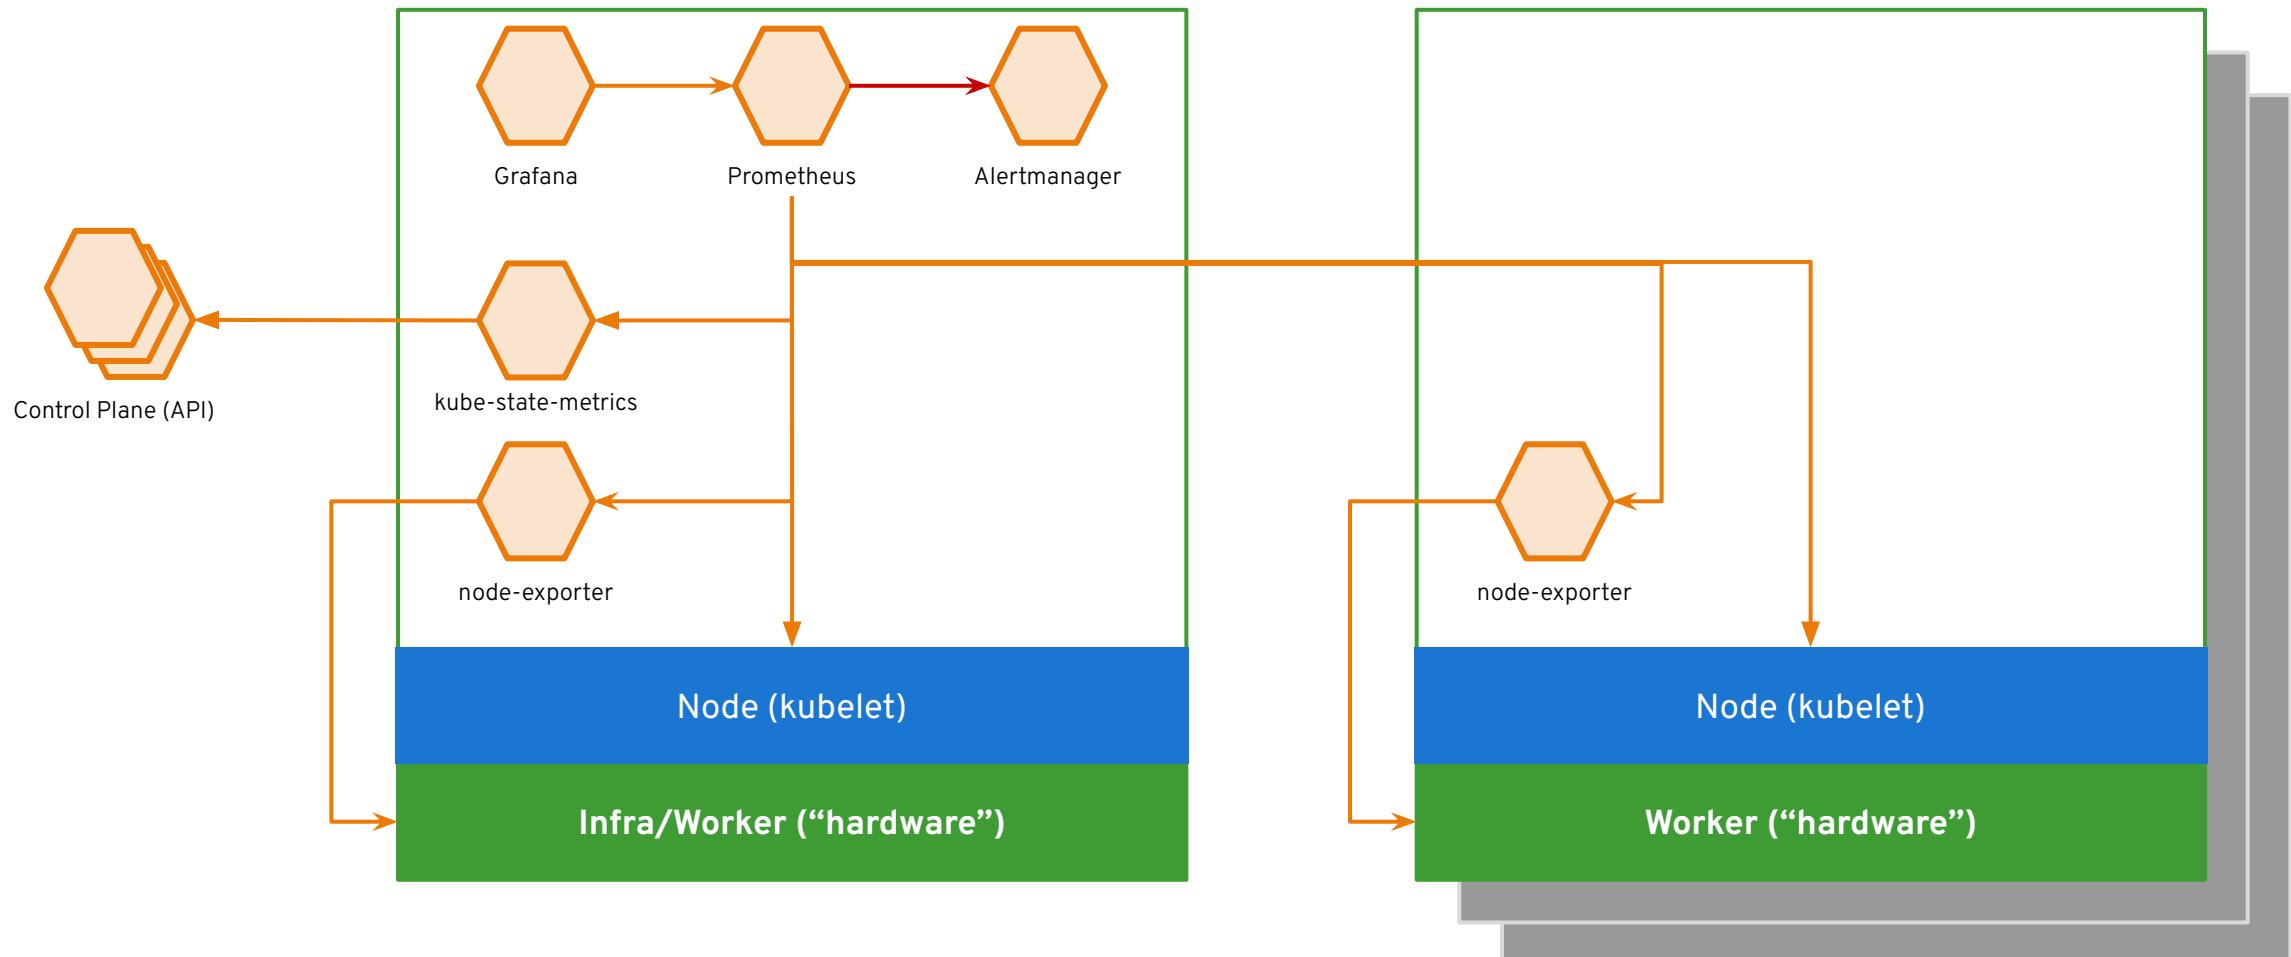


Alerting/notification via Prometheus' Alertmanager, an open-source tool that handles alerts send by Prometheus.



Metrics visualization via Grafana, the leading metrics visualization technology.





Exercise 11: Project Template, Quota, and Limits

Forcing users to be
well-behaved with respect
to resource consumption

Types of resources

Compressible: there is always more, you just have to wait for the next unit of time

- Cpu, network bandwidth, iops ...
- Compressible resources are always expressed as qyt/unit of time for example cpu/millisec

Incompressible: once used up, there is no more.

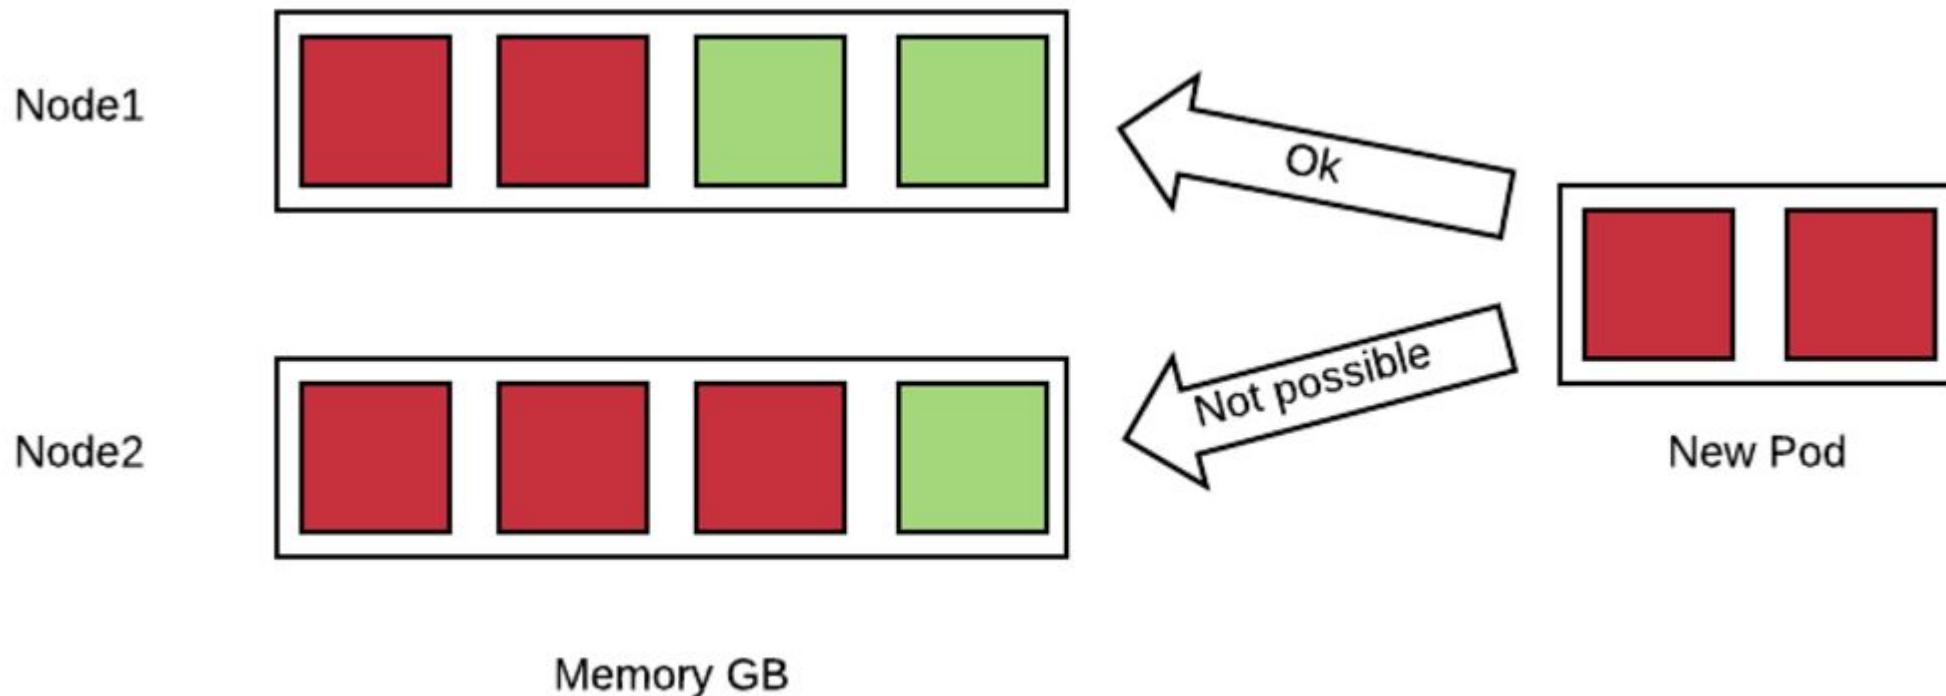
- Memory, diskspace ...

Quotas, requests and limits

Constraint	When it's Evaluated	Purpose
Quotas	Request Time	To limit how many resources a single tenant can request so that they cannot take over the cluster.
Requests	Scheduling Time	Requests are used to choose which node is the best fit for the given workload (together with other criteria).
Limits	Run Time	Limits are used to limit the maximum amount of resources that a container can use at runtime.

Scheduling a pod

(how memory requests affects scheduling a pod)



Red squares on the nodes represent the sum of the requests of previously scheduled pods.

Compare the output of “oc describe node” and “oc adm top”.

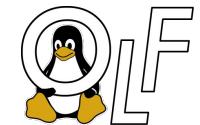
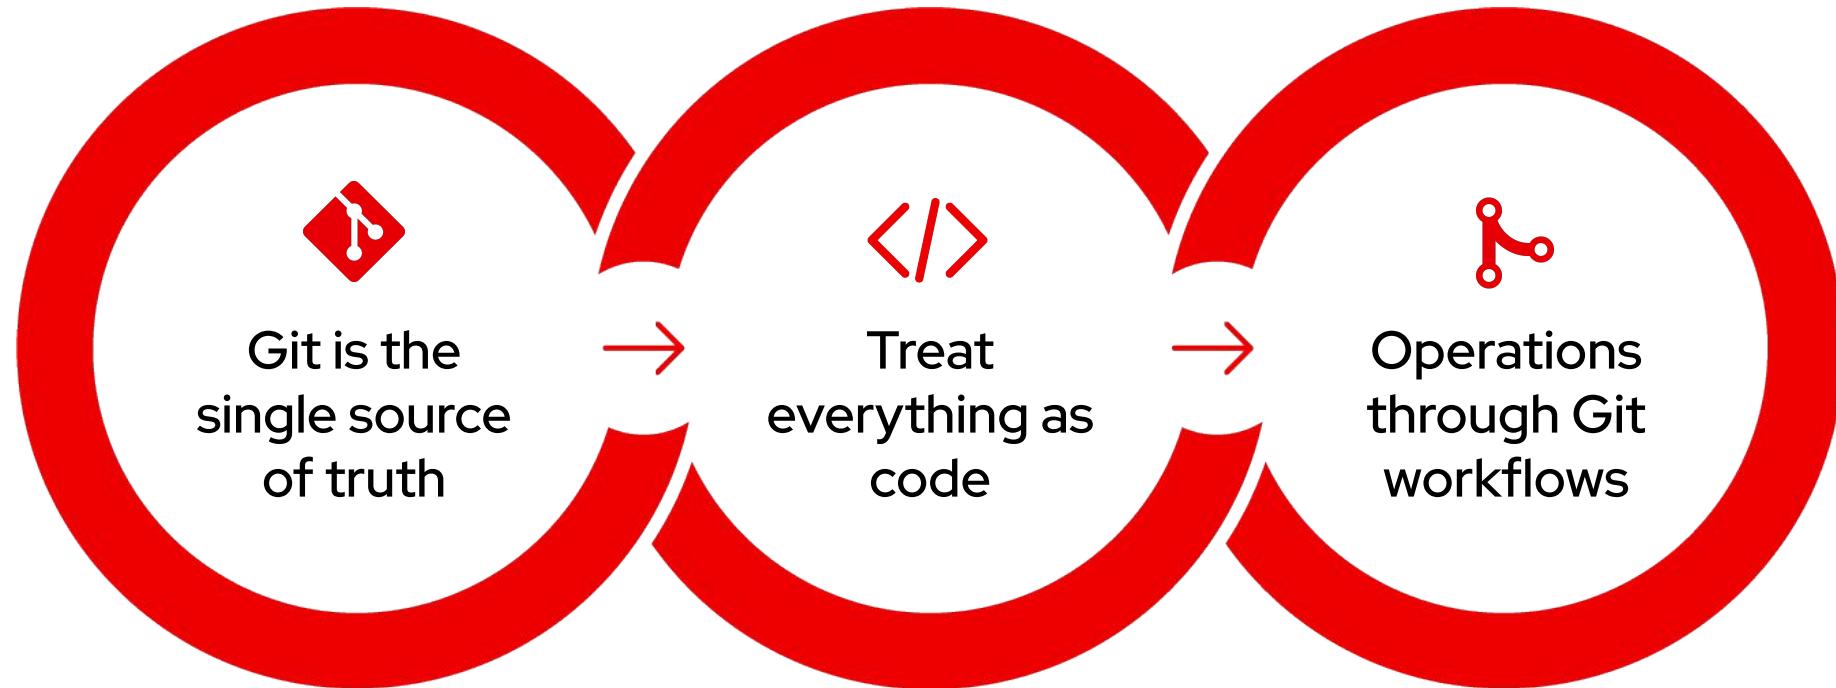
Exercise 12: Networking and NetworkPolicy

Fine-grained control over
network connectivity
between and even inside
tenants

OpenShift GitOps

What is GitOps?

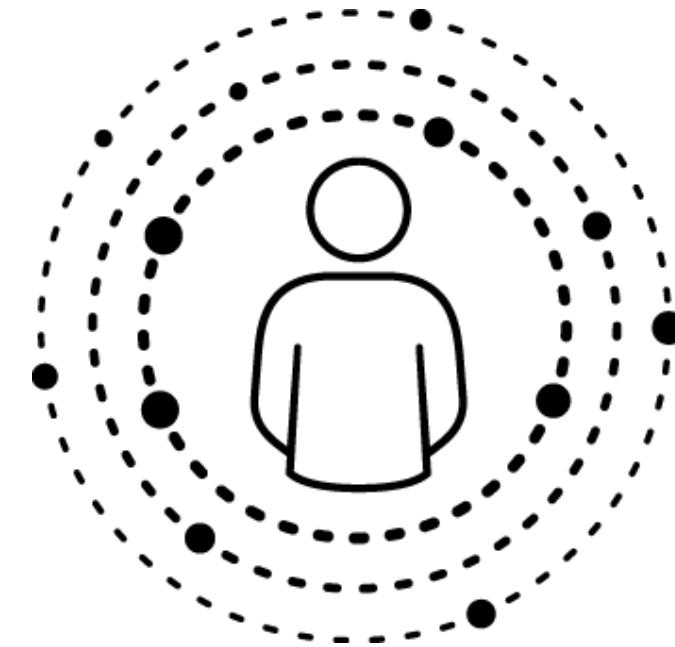
An developer-centric approach to Continuous Delivery and infrastructure operation



GitOps is for Everyone



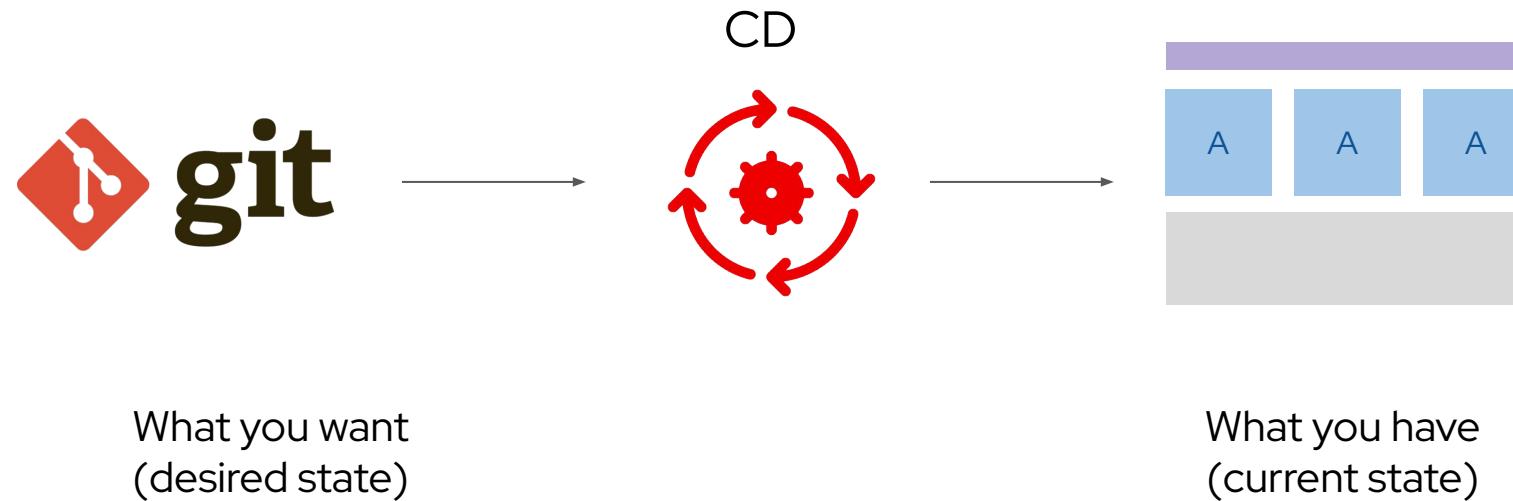
Developers



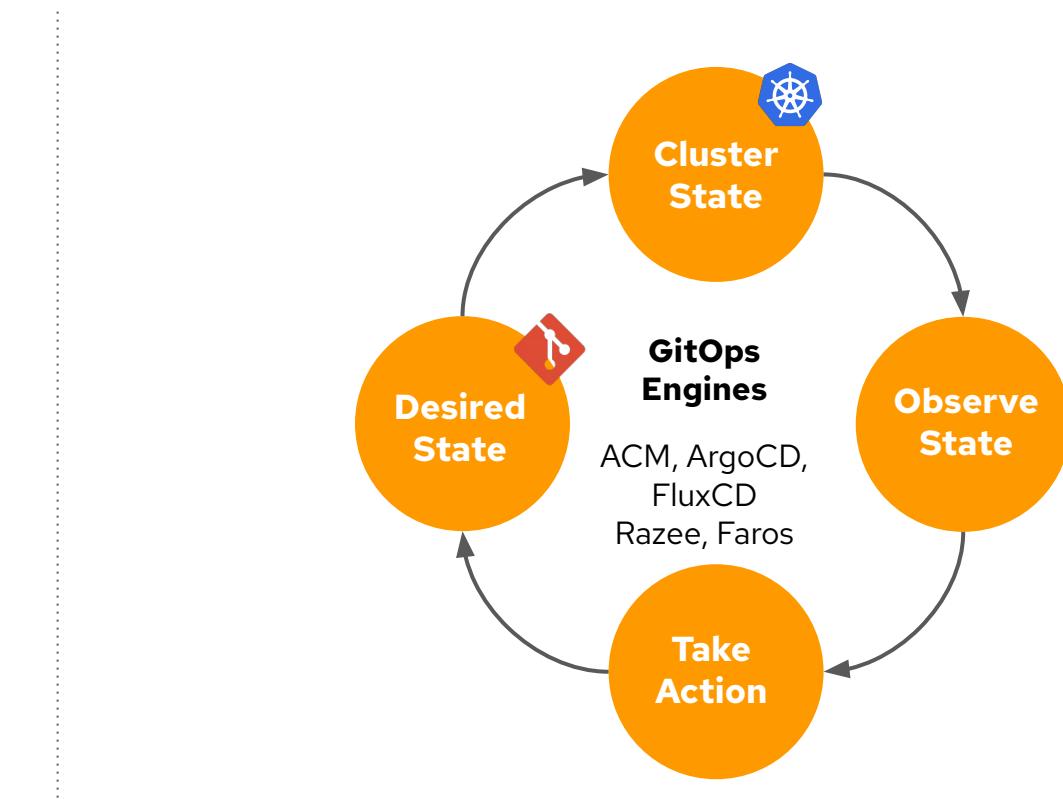
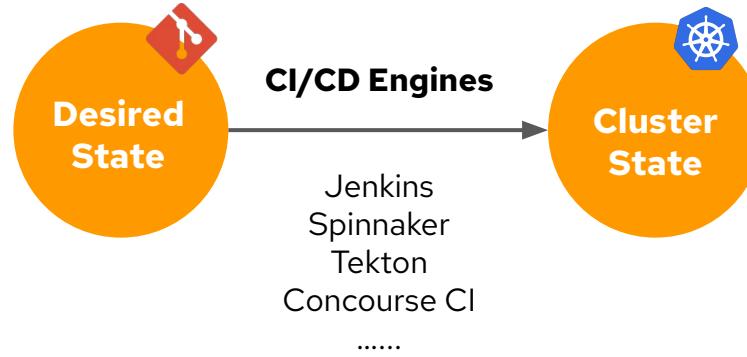
Operations

GitOps Workflow

a declarative approach to application delivery



GitOps versus CI/CD



Why GitOps?

Standard Workflow

Familiar tools and Git workflows from application development teams

Enhanced Security

Review changes beforehand, detect configuration drifts, and take action

Visibility and Audit

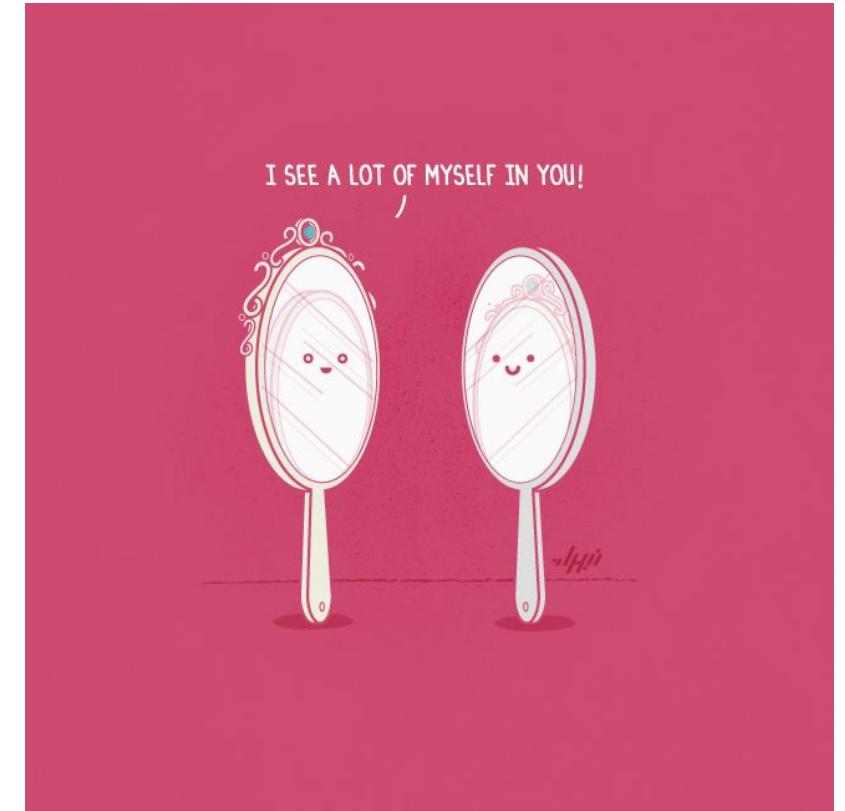
Capturing and tracing any change to clusters through Git history

Multi-cluster consistency

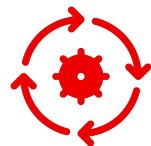
Reliably and consistently configure multiple Kubernetes clusters and deployment

Kubernetes and GitOps - A Perfect Match

- Kubernetes is a declarative environment
 - Application deployments are declared and Kubernetes scheduler makes it happen
 - OpenShift goes further in that cluster configuration is declared and Operators make it happen
- GitOps in traditional environments requires automation/scripting, declarative environment minimizes or eliminates this need
- Declarations are yaml files which are easily stored and managed in git



OpenShift GitOps



Multi-cluster config management

Declaratively manage cluster and application configurations across multi-cluster OpenShift and Kubernetes infrastructure with Argo CD



Automated Argo CD install and upgrade

Automated install, configurations and upgrade of Argo CD through OperatorHub



Opinionated GitOps bootstrapping

Bootstrap end-to-end GitOps workflows for application delivery using Argo CD and Tekton with GitOps Application Manager CLI



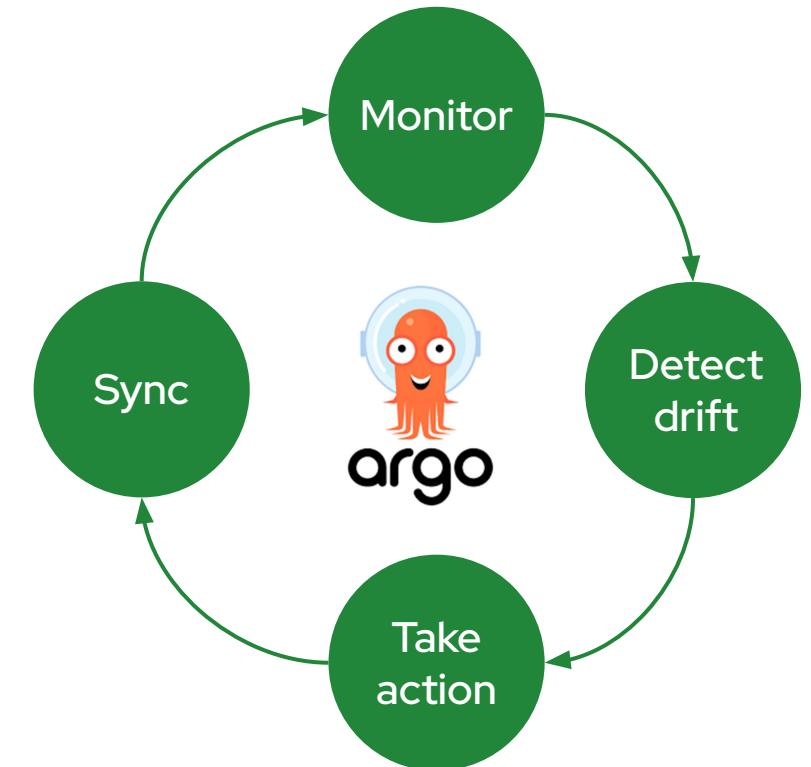
Deployments and environments insights

Visibility into application deployments across environments and the history of deployments in the OpenShift Console

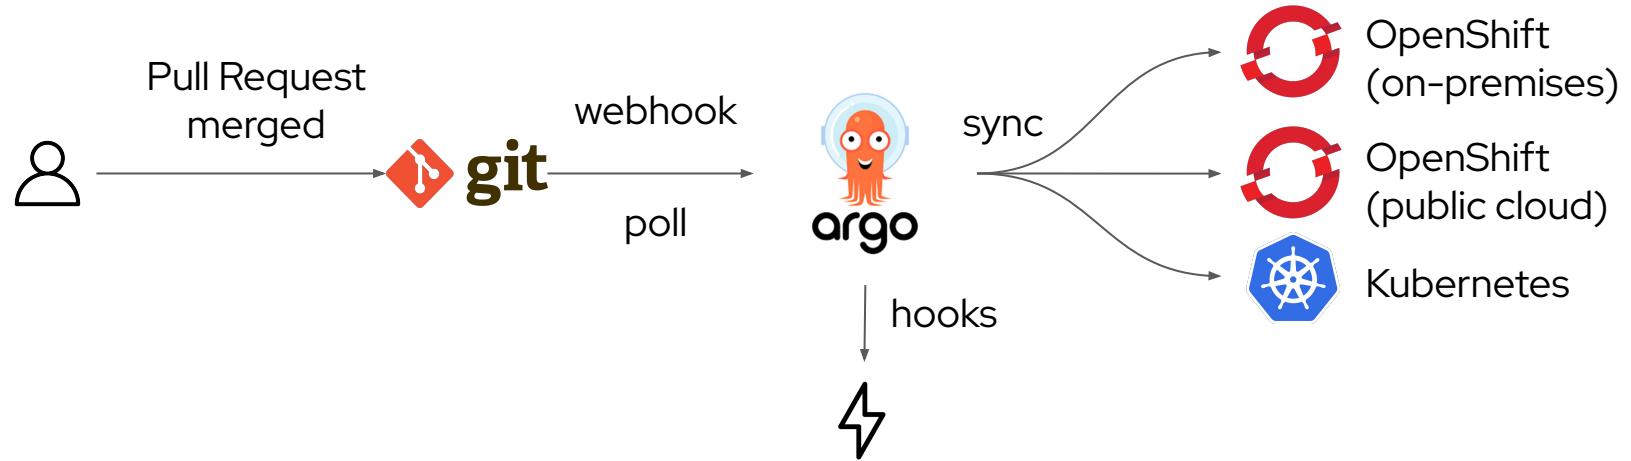
Argo CD

- Cluster and application configuration versioned in Git
- Automatically syncs configuration from Git to clusters
- Drift detection, visualization and correction
- Granular control over sync order for complex rollouts
- Rollback and rollforward to any Git commit
- Manifest templating support (Helm, Kustomize, etc)
- Visual insight into sync status and history

127



Argo CD



A screenshot of the Argo CD web interface for the "realtime" application. The top navigation bar includes "APPLICATION DETAILS", "APP DIFF", "SYNC", "SYNC STATUS", "HISTORY AND ROLLBACK", "DELETE", and "REFRESH". The main area shows the application status as "Healthy" and "Synced". A "synchronization" section indicates a successful sync to revision 49147e1. The interface displays a complex dependency graph between multiple "redisapp" and "realtimeapp" components across two clusters, with deployment and replicaset stages.

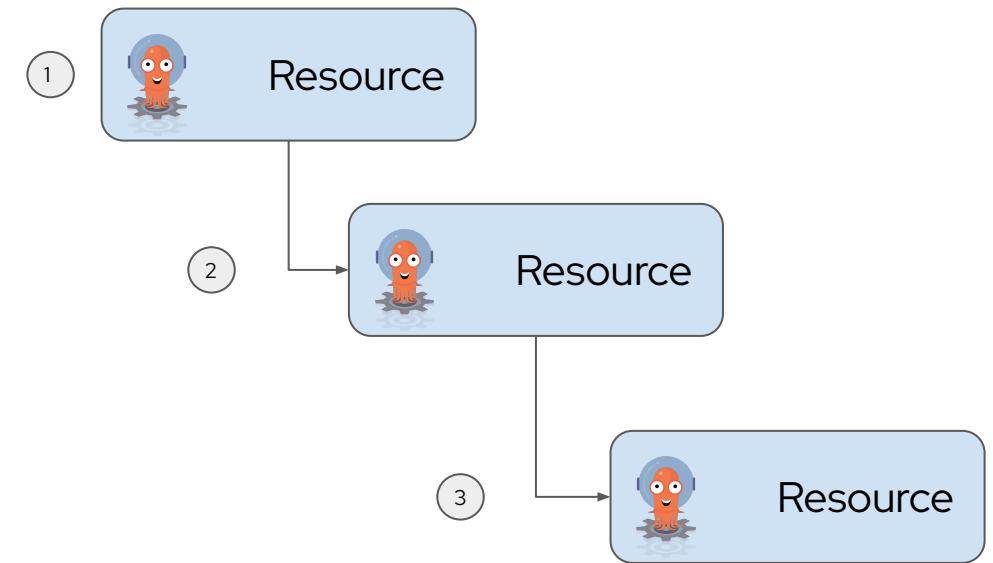
What is an Argo CD Application?

- Argo CD Application is a Custom Resource (CR) that defines the app in a declarative manner
- Application definition includes:
 - Name
 - Cluster
 - Git repository
 - Synchronization Policy
- Applications can be deployed from Argo CD GUI or CLI (argocd or kubectl or oc)

```
apiVersion: argoproj.io/v1alpha1
kind: Application
metadata: [REDACTED]
name: product-catalog-dev
namespace: argocd
spec: [REDACTED]
destination:
  namespace: argocd
  server: https://kubernetes.default.svc
project: product-catalog
source:
  path: manifests/app/overlays/dev-quay
  repoURL: https://github.com/gnunn-gitops/product-catalog.git
  targetRevision: master
syncPolicy:
  automated:
    prune: false
    selfHeal: false
```

Argo CD Sync Waves

- Sometimes you need to deploy resources in a specific order due to dependencies, pre-reqs, etc
- Argo CD enables you to order the deployment using Sync Waves. Next resource will not deploy until the first resource is “healthy”
- Very useful when resources have hard dependencies versus eventual consistency



GitOps - Avoiding Duplication

GitOps enables deployment across multiple environments and clusters, awesome!

Wait, how do we manage configuration without copying and pasting yaml everywhere?





Helm is a package manager for Kubernetes
applications that manages manifests via templating
define, install and update applications



Kustomize is a tool for customizing Kubernetes manifests via patching, it is built natively into the “kubectl/oc” cli with a standalone “kustomize” cli for advanced features.





What does it look like?



Kustomization

A file that declares any resources, and any customization to apply to them, e.g. add a common label or namespace



Overlay

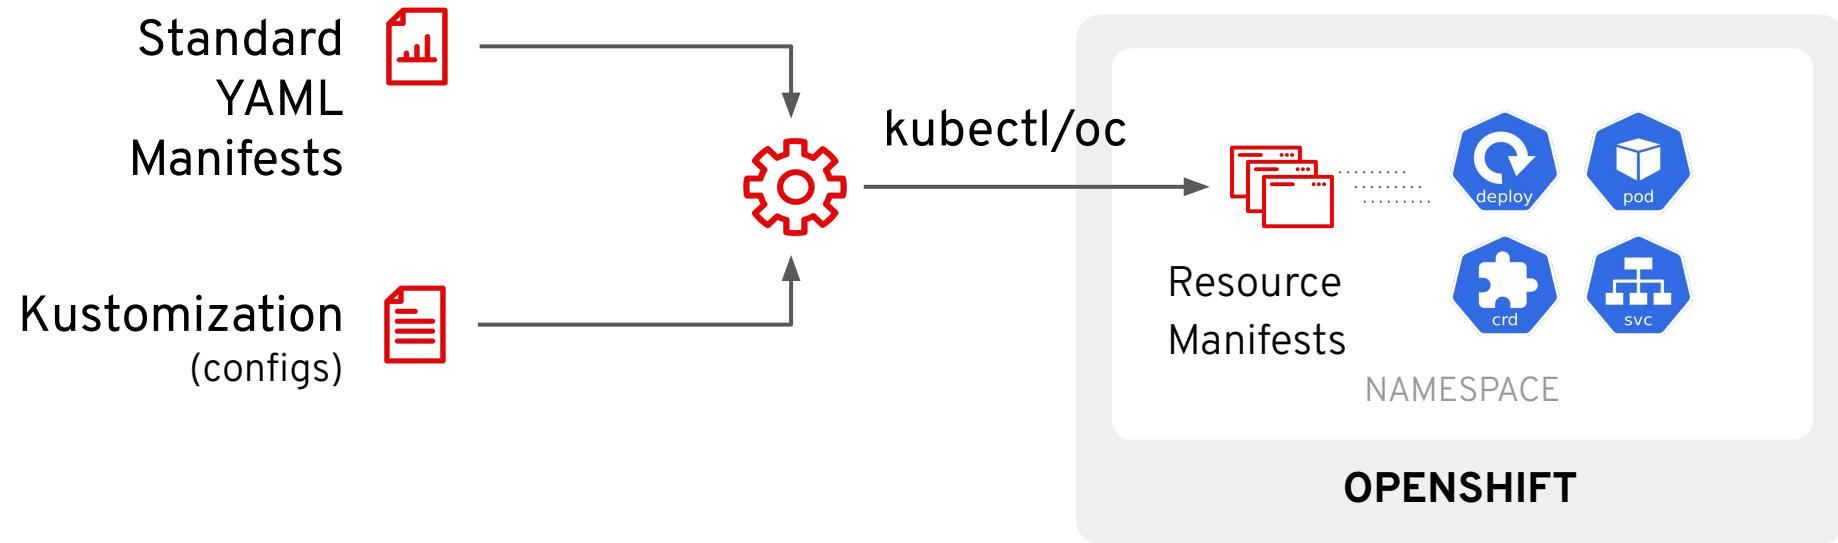
Manage variants of a configuration using overlays that modify a common base.

Example folder structure

```
~/someApp
  └── base
      ├── deployment.yaml
      ├── kustomization.yaml
      └── service.yaml
  └── overlays
      ├── development
          ├── cpu_count.yaml
          ├── kustomization.yaml
          └── replica_count.yaml
      └── production
          ├── cpu_count.yaml
          ├── kustomization.yaml
          └── replica_count.yaml
```



How does it work?



Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.

 linkedin.com/company/red-hat

 youtube.com/user/RedHatVideos

 facebook.com/redhatinc

 twitter.com/RedHat

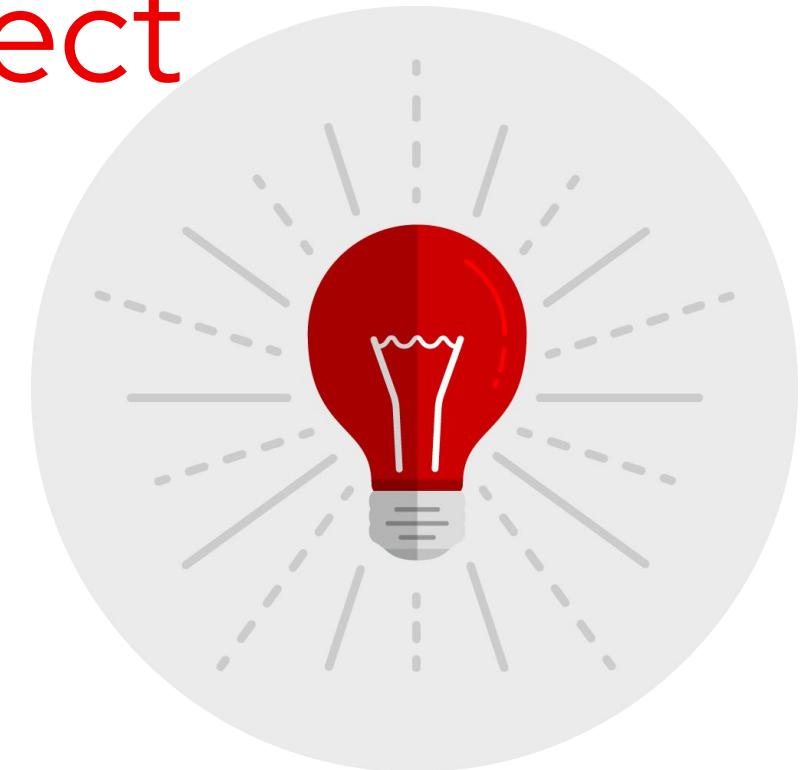
Red Hat Summit Connect

Cincinnati - Sept 14th

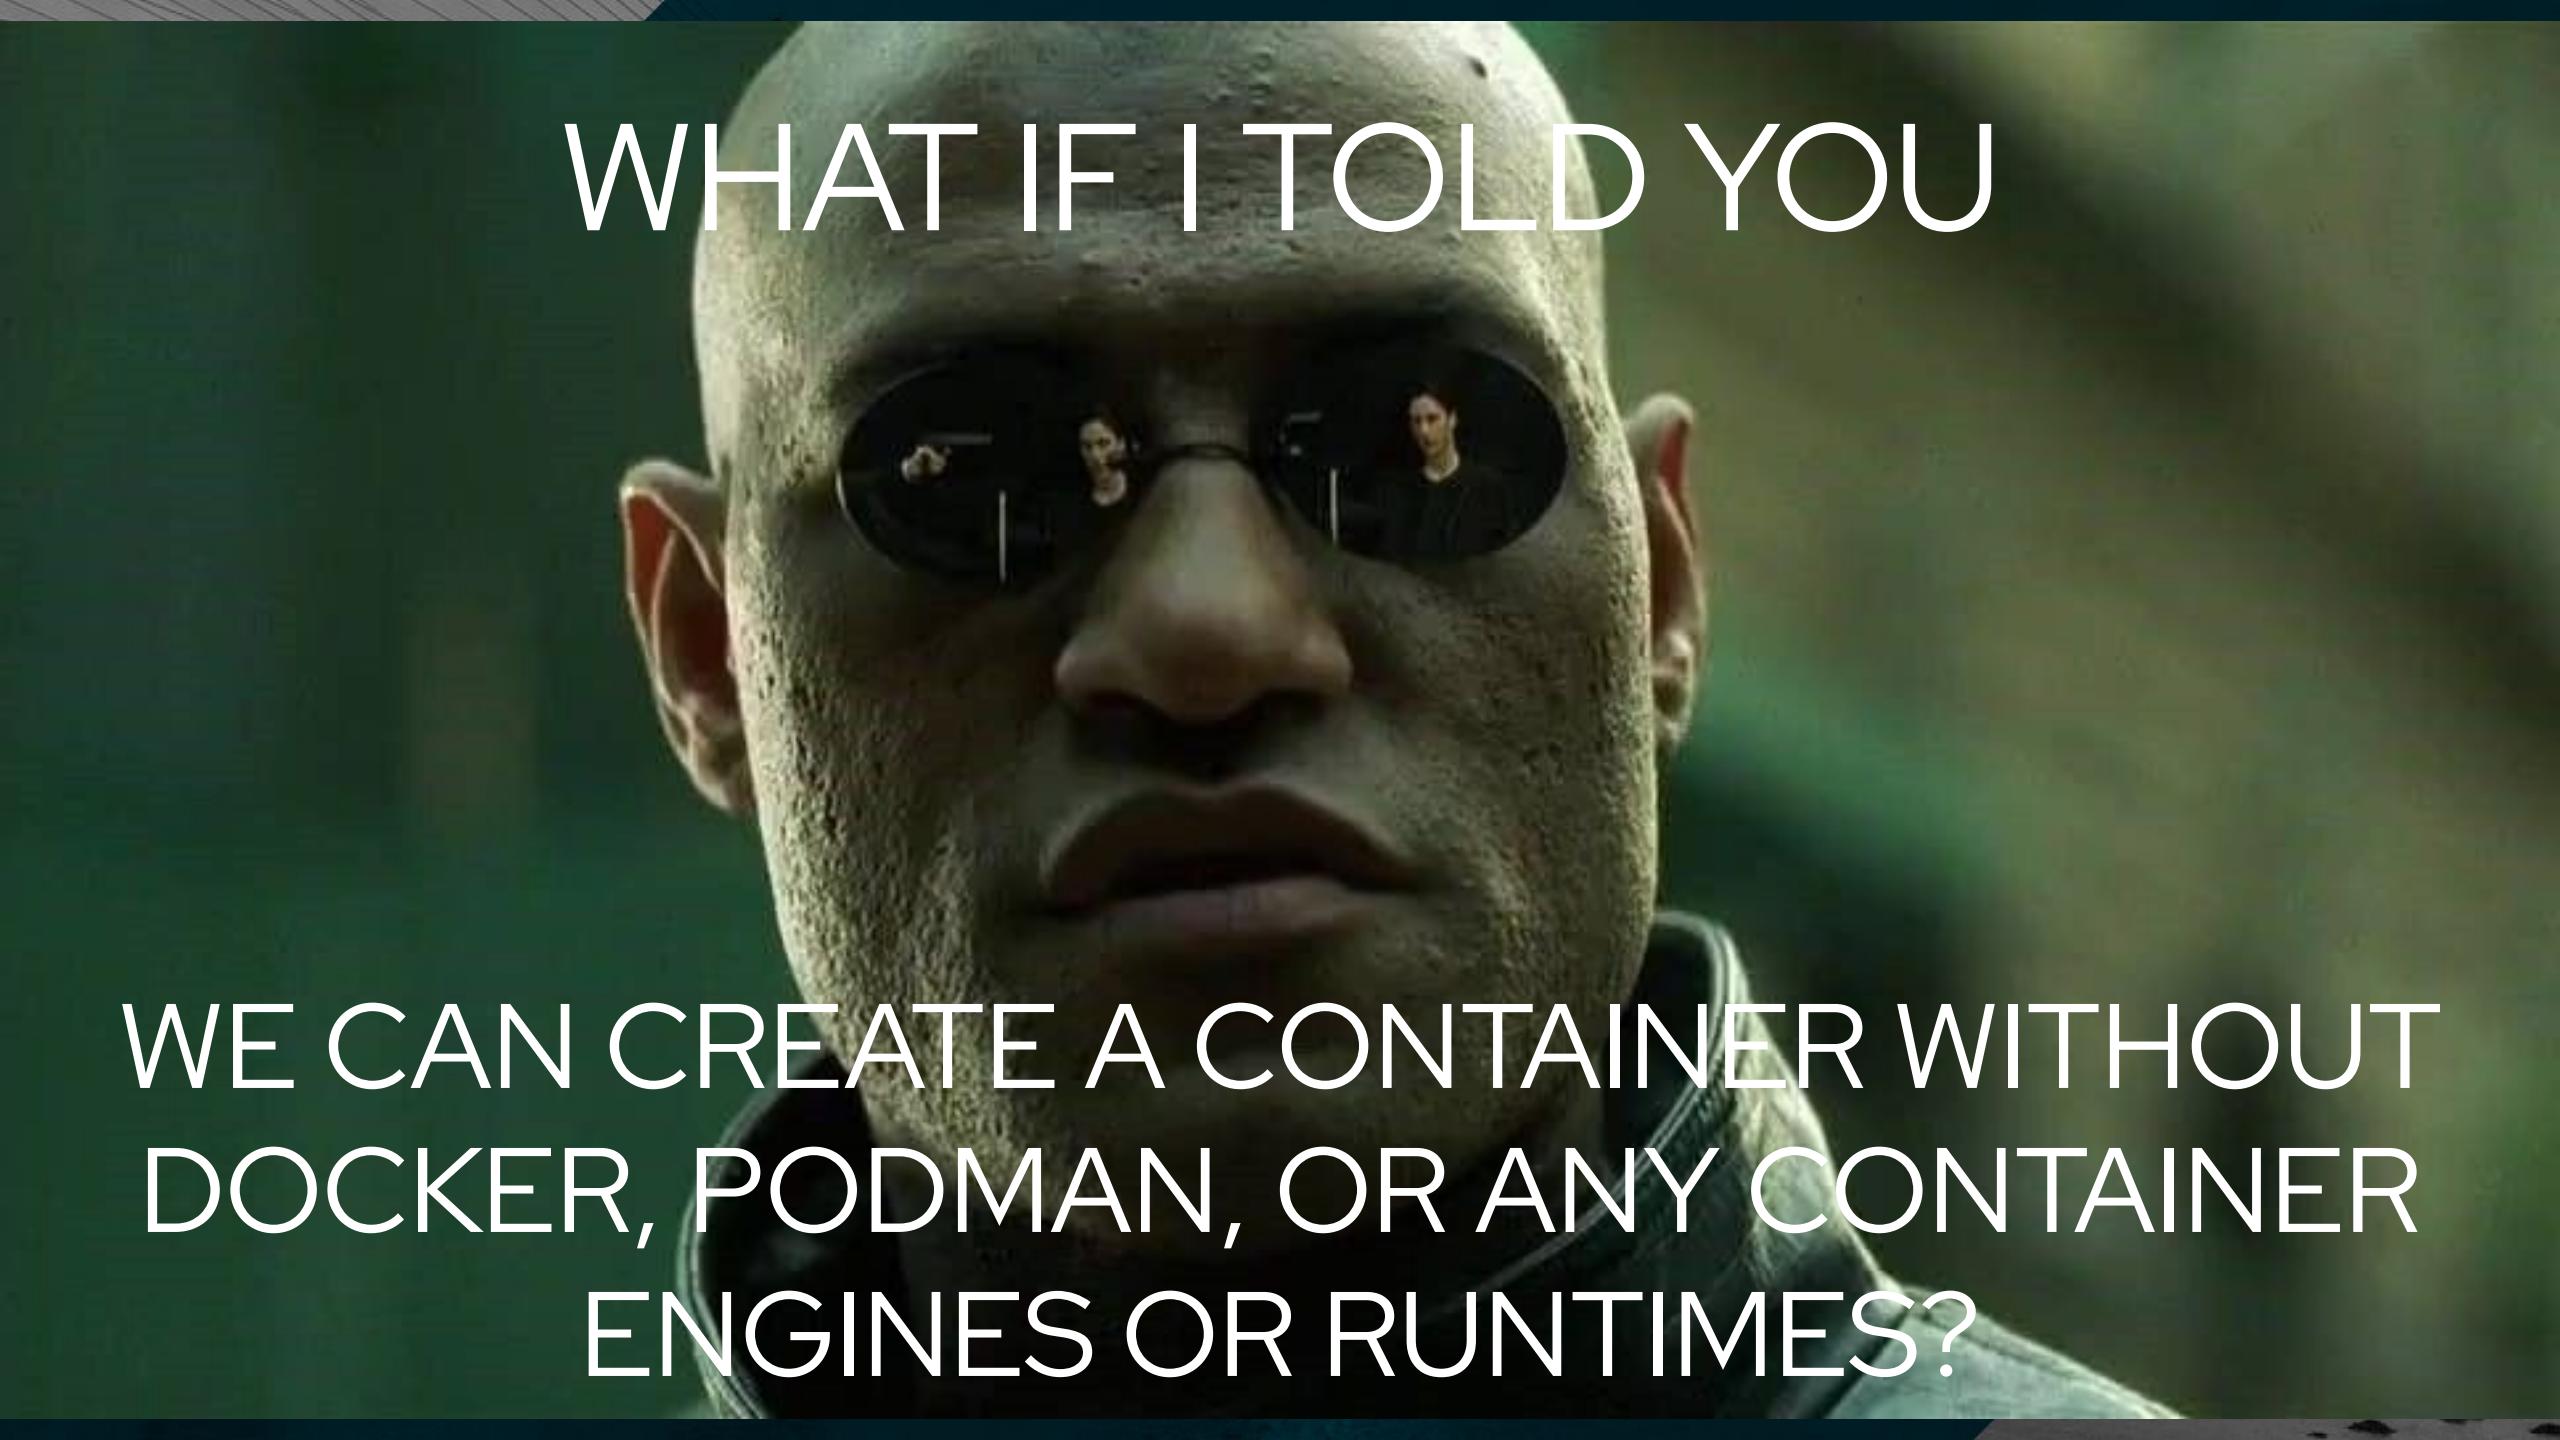
Other Cities too!

Hands On Labs

- Openshift Virtualization
- Event Driven Ansible



CONTAINERS FROM SCRATCH



WHAT IF I TOLD YOU

WE CAN CREATE A CONTAINER WITHOUT
DOCKER, PODMAN, OR ANY CONTAINER
ENGINES OR RUNTIMES?

How do we make one?

- It all starts with a process
- Then we wrap some Linux/Unix boundaries around it that....
 - Partition kernel resources (filesystem, IPC, network, PID, users, more)
 - Limit resource usage
 - Prioritize or de-prioritize CPU share and Disk I/O
 - Control freezing, stopping, starting and checkpointing
 - Account for usage of resources managed above
 - Limit what/who they can talk to
 - Limit what they can say

Namespaces

- It all starts with a process
- Then we wrap some boundaries around it that....
 - Partition kernel resources (filesystem, IPC, network, PID, users, more)
 - Limit resource usage
 - Prioritize or de-prioritize CPU share and Disk I/O
 - Control freezing, stopping, starting and checkpointing
 - Account for usage of resources managed above
 - Limit what/who they can talk to (MAC)
 - Limit what they can say

CGroups

- It all starts with a process
- Then we wrap some boundaries around it that....
 - Partition kernel resources (filesystem, IPC, network, PID, users, more)
 - Limit resource usage
 - Prioritize or de-prioritize CPU share and Disk I/O
 - Control freezing, stopping, starting and checkpointing
 - Account for usage of resources managed above
- Limit what/who they can talk to (MAC)
- Limit what they can say

SELinux / AppArmor

- It all starts with a process
- Then we wrap some boundaries around it that....
 - Partition kernel resources (filesystem, IPC, network, PID, users, more)
 - Limit resource usage
 - Prioritize or de-prioritize CPU share and Disk I/O
 - Control freezing, stopping, starting and checkpointing
 - Account for usage of resources managed above
 - Limit what/who they can talk to (MAC) (highlighted)
 - Limit what they can say

Stop Disabling SELINUX!

<https://stopdisablingselinux.com/>

#MakeSELINUXenforcingAgain
#SETENFORCE1

SECCOMP

- It all starts with a process
- Then we wrap some boundaries around it that....
 - Partition kernel resources (filesystem, IPC, network, PID, users, more)
 - Limit resource usage
 - Prioritize or de-prioritize CPU share and Disk I/O
 - Control freezing, stopping, starting and checkpointing
 - Account for usage of resources managed above
 - Limit what/who they can talk to (MAC)
 - Limit what they can say

LAB

CONTAINERS FROM SCRATCH (~20 min)

<https://red.ht/olf-container>

